



Machine perception

Image formation & Image processing 1

Matej Kristan



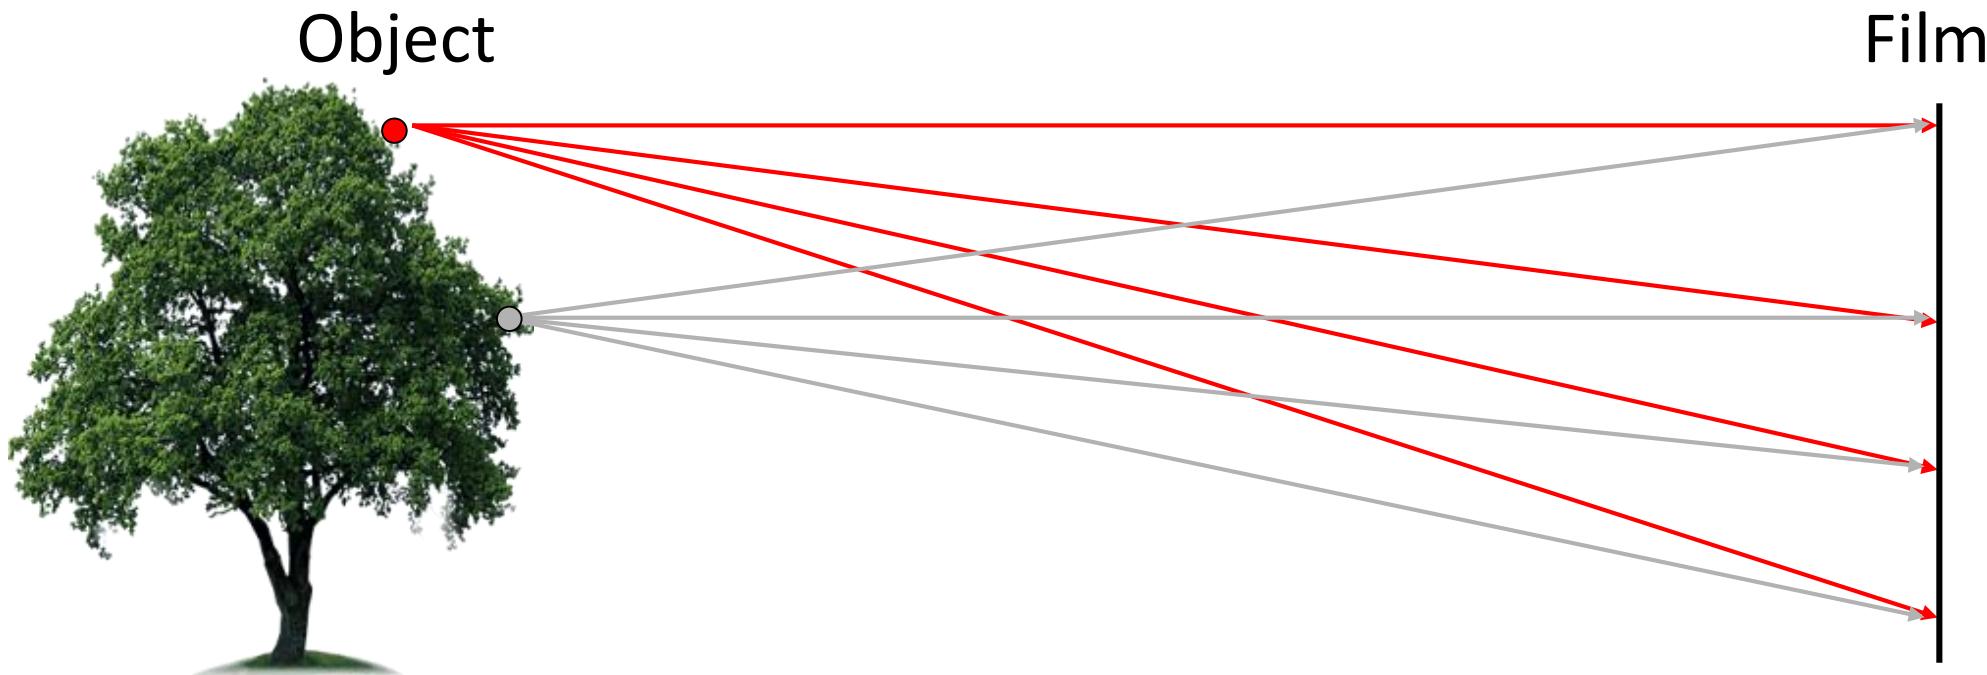
Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani



Machine perception

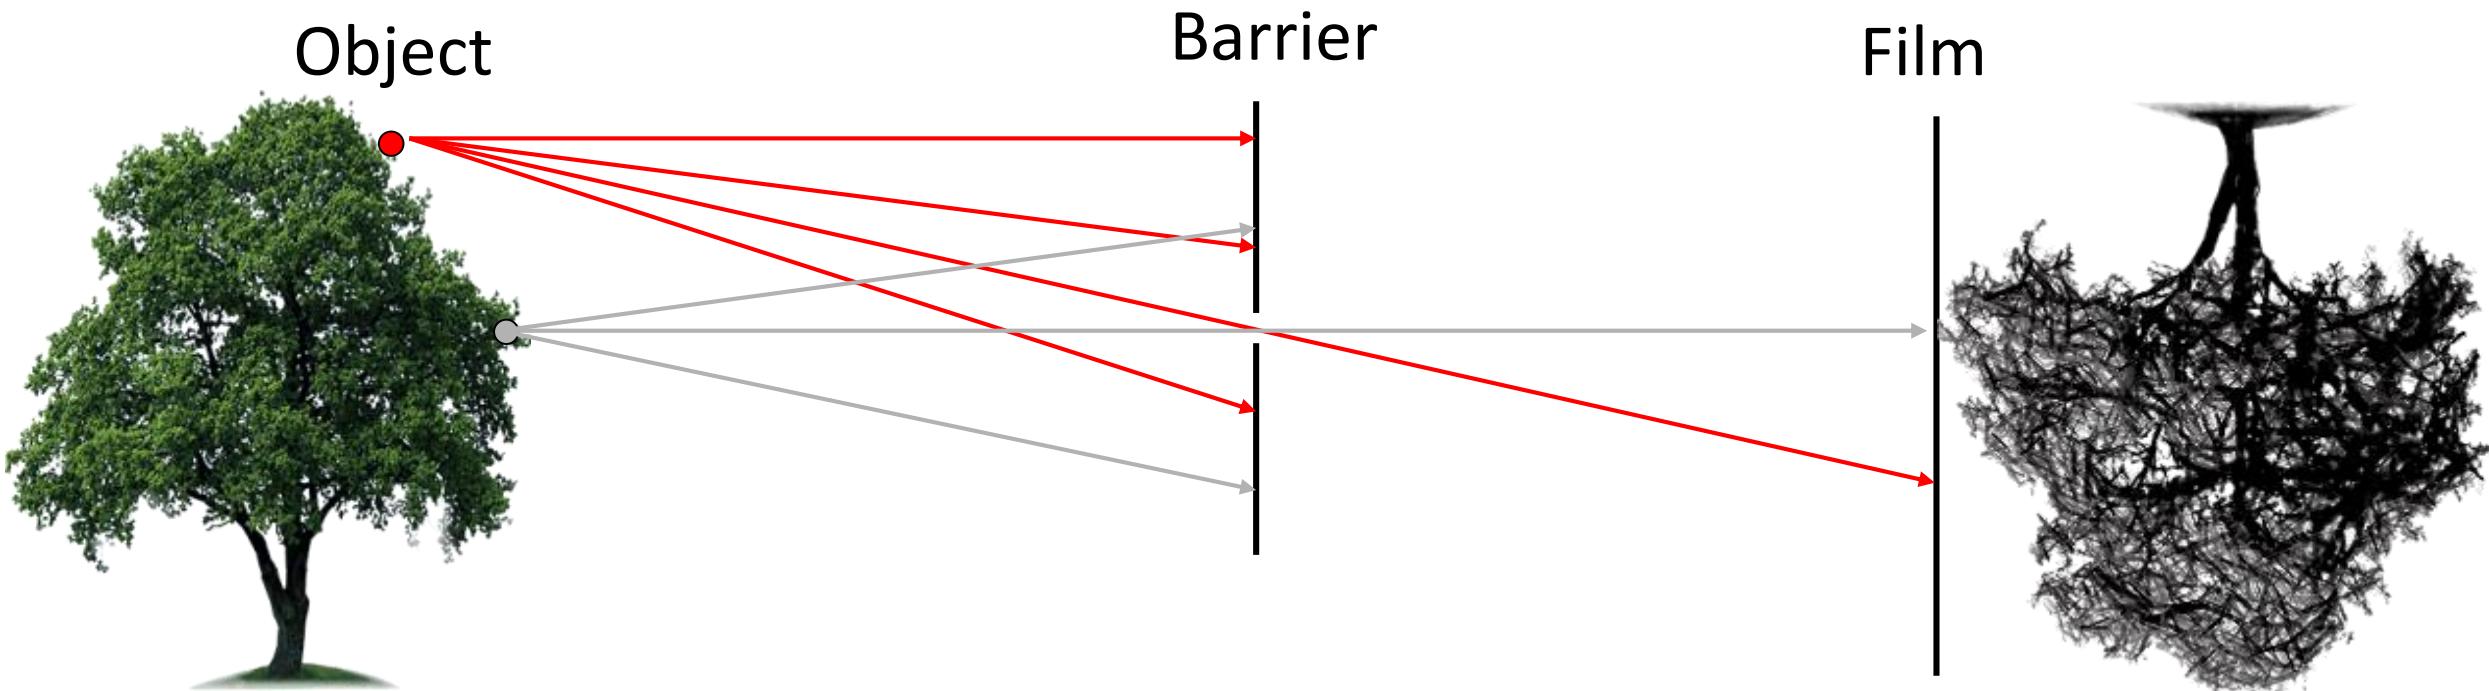
IMAGE FORMATION

Let's design a camera!



- Idea 1: put an object in front of a film...
- Do we get a good image of the object?

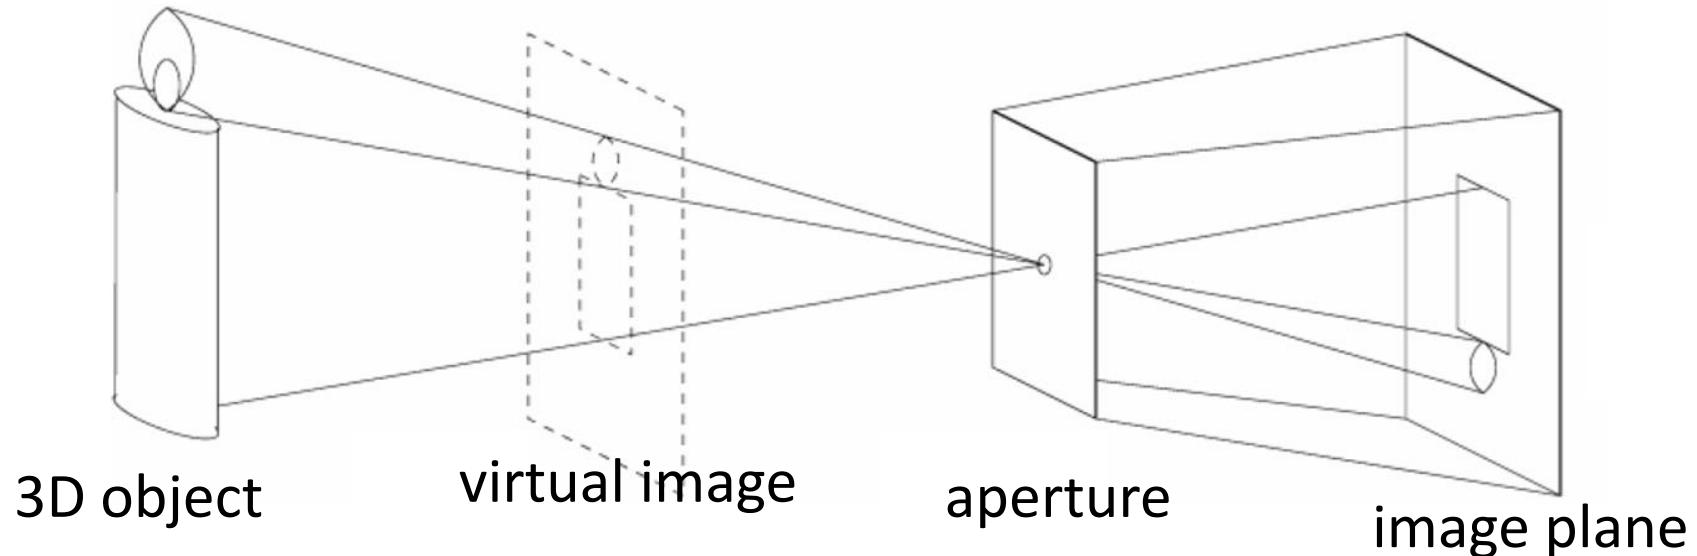
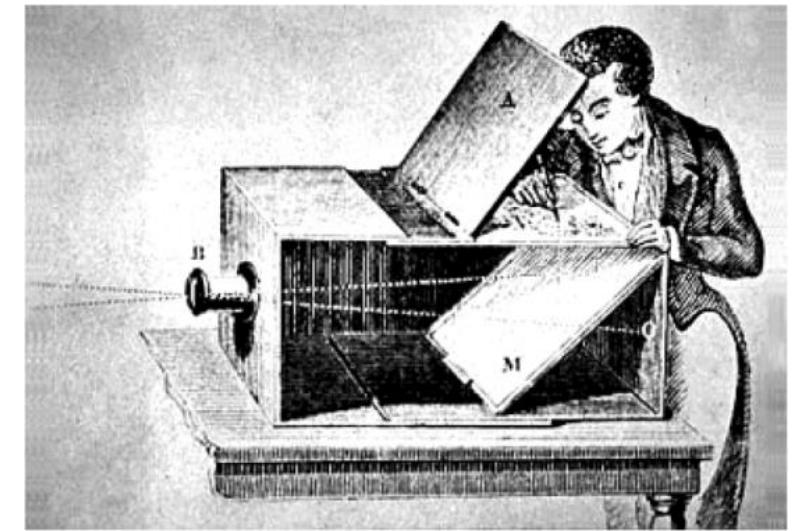
Let's design a camera!



- Add a barrier that blocks most of the rays
 - Significantly reduces blurring
 - The „hole“ is known as aperture

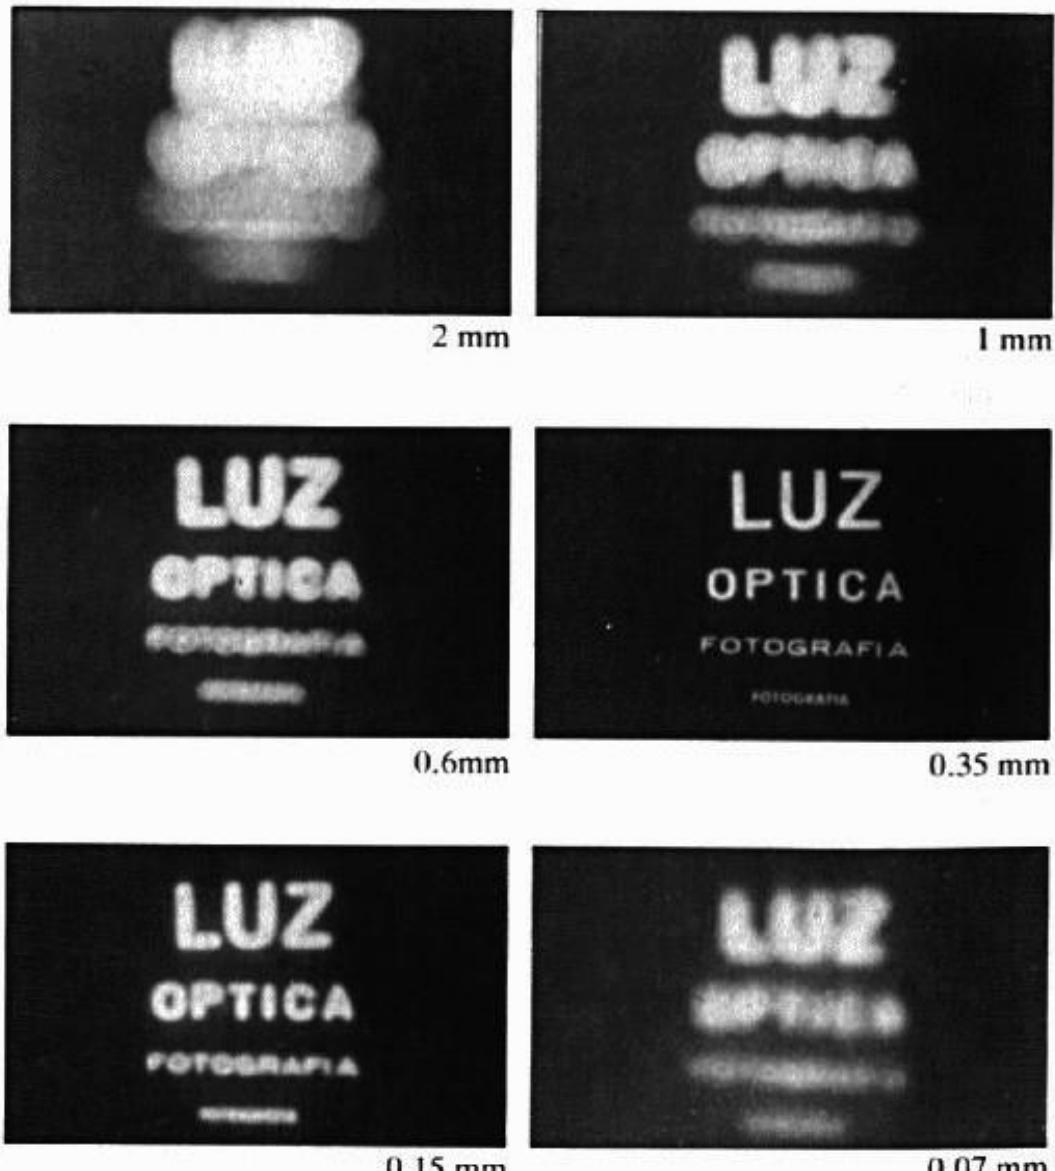
A pinhole camera

- 500 BC Mo Ti...
- ...1485 Leonardo da Vinci
- A simple standard camera model
 - A box with a small aperture
 - Works in practice

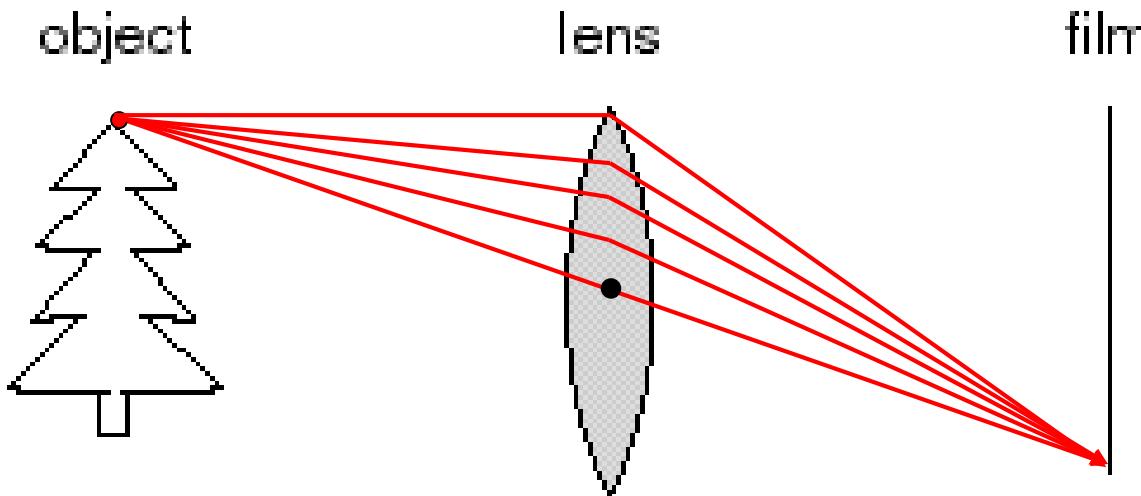


Effects of the aperture size

- Too large – multiple directions averaging, resulting in a blurred image.
- Too small – light starts diffracting, causing blurred image.
- In general small number of rays hit the film, which results in a dark image.
- How do we deal with this?

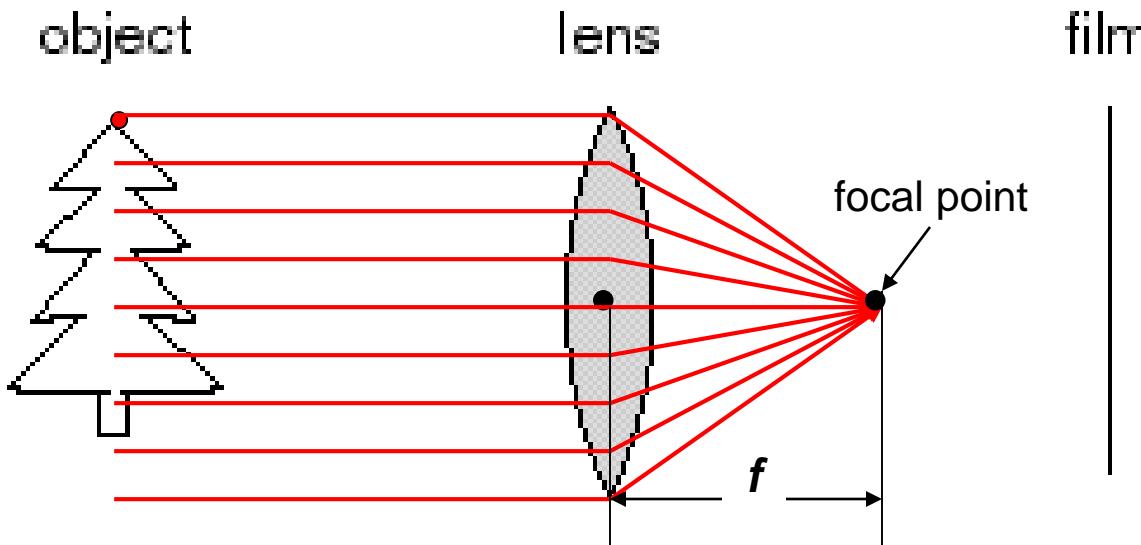


Let's add a lens...



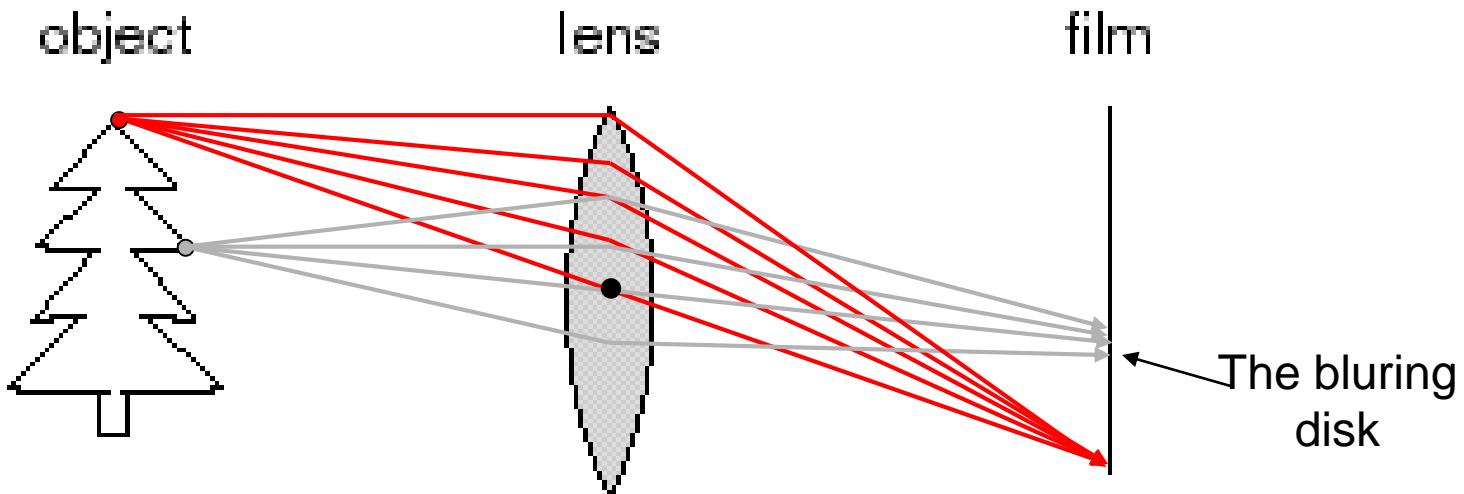
- The lens focuses light to film
 - The rays that travel through the **center** do not refract.

Let's add a lens...



- The lens focuses light to film
 - The rays that travel through the **center** do not **refract**.
 - **Parallel rays** intersect in a point that is at distance f from the lens (**focal point**).

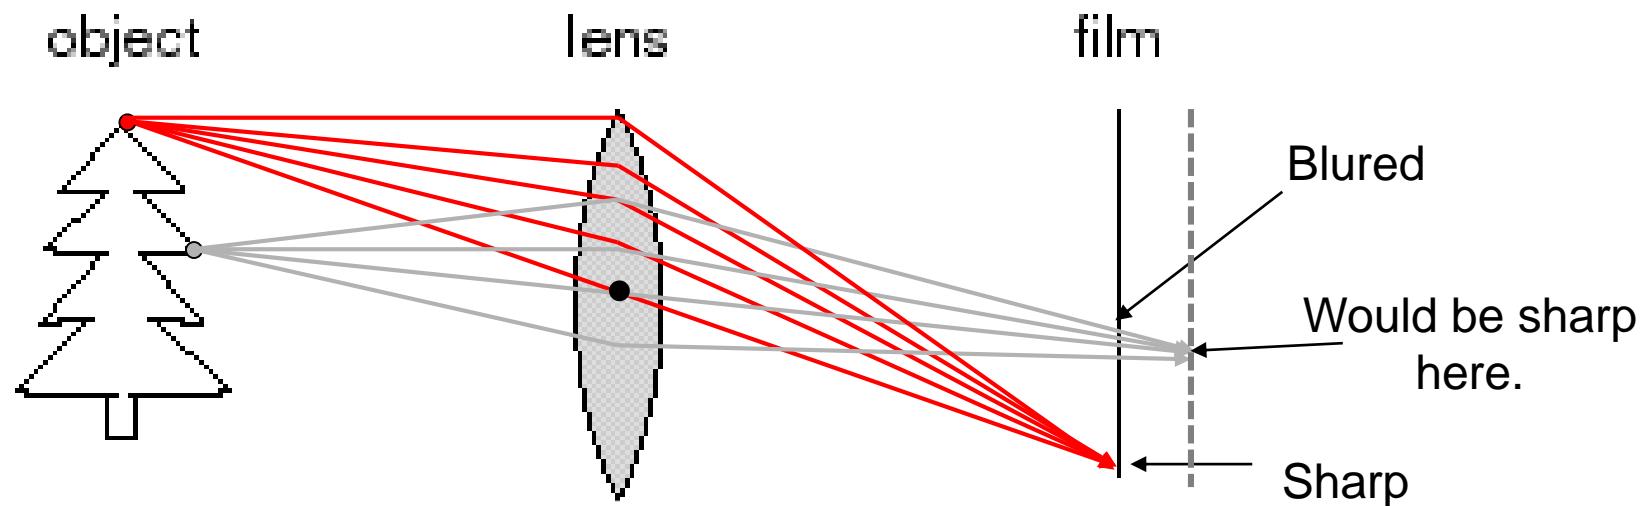
Let's add a lens...



- The lens focuses light to film
 - The rays that travel through the **center** do not refract.
 - Points at particular distance remain in-focus.
 - Points at other distances are blurred.

Focus and the depth-of-field

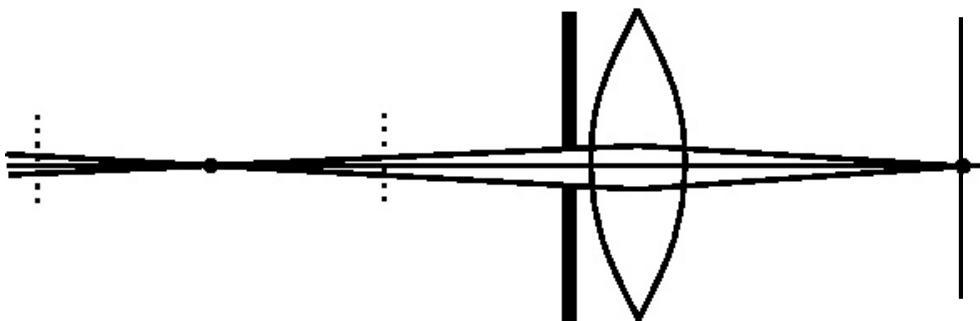
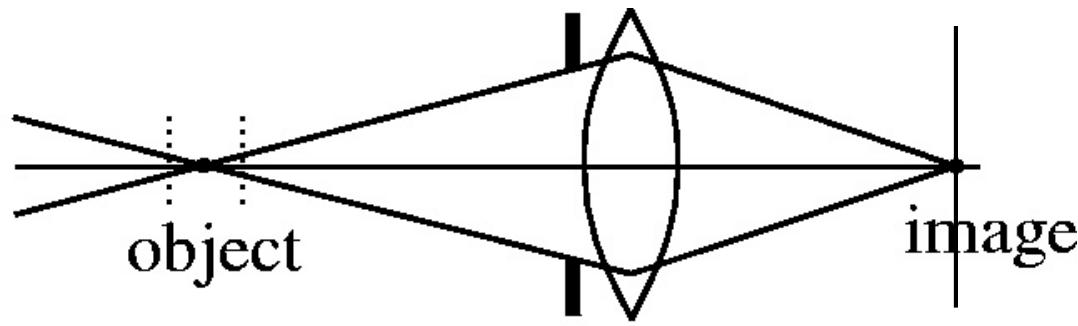
- **Thin lens:** Points at different depths get focused at different depths of image plane.
(Real-world lens have a greater depth of field)



- **Depth of field:** distance between image planes at which the blurring effect is sufficiently small..

Focus and the depth-of-field

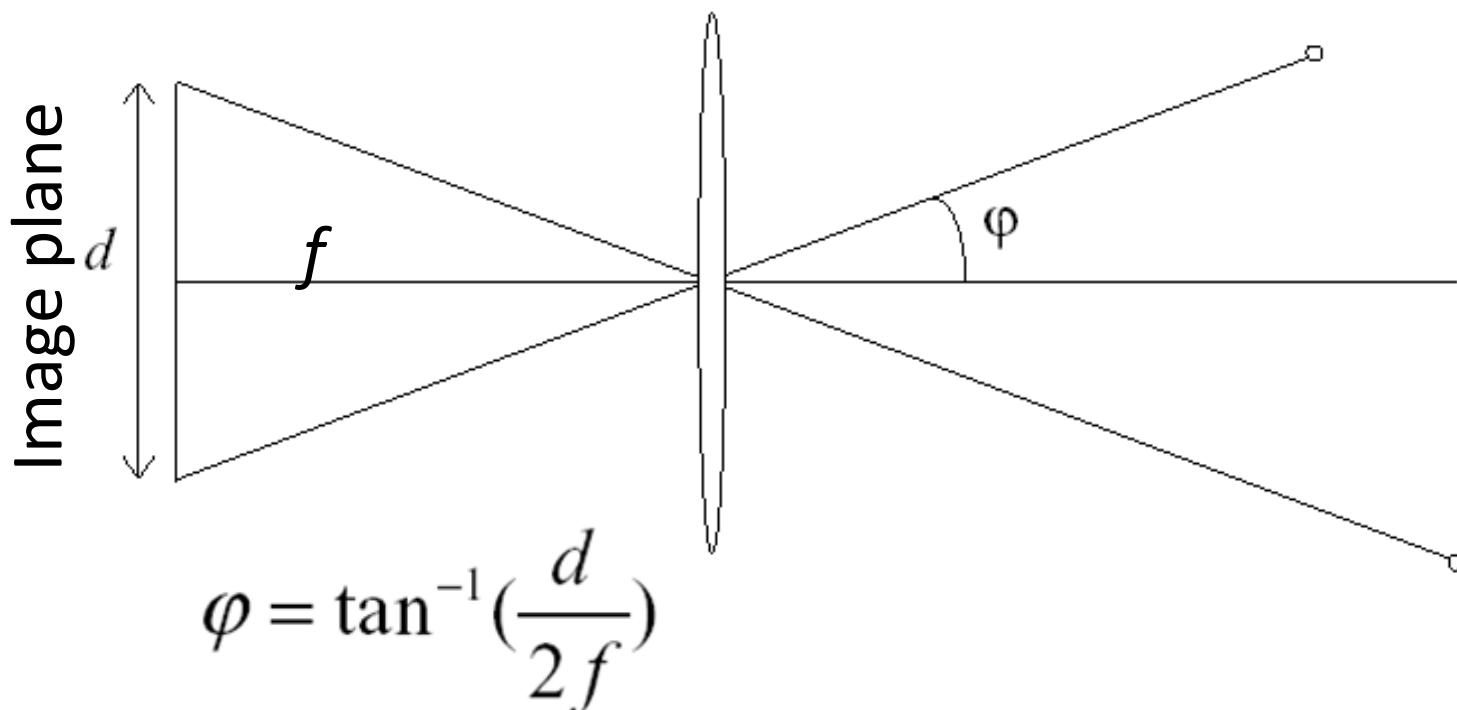
- Effects of aperture on the depth-of-field



- Small aperture increases the depth-of-field.
- But due to small illumination we have to increase the exposure time.

Field of view

- Field of view (FOV) ($2 \times \varphi$) is an angular measure of space perceived by the camera.



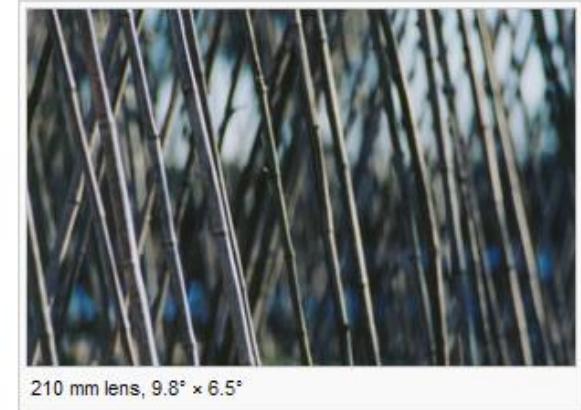
- Larger focal length → Smaller field of view

Field of view

- Small f results in wide-angle image
(Large field of view) →
 - More 3D points project to the sensor.

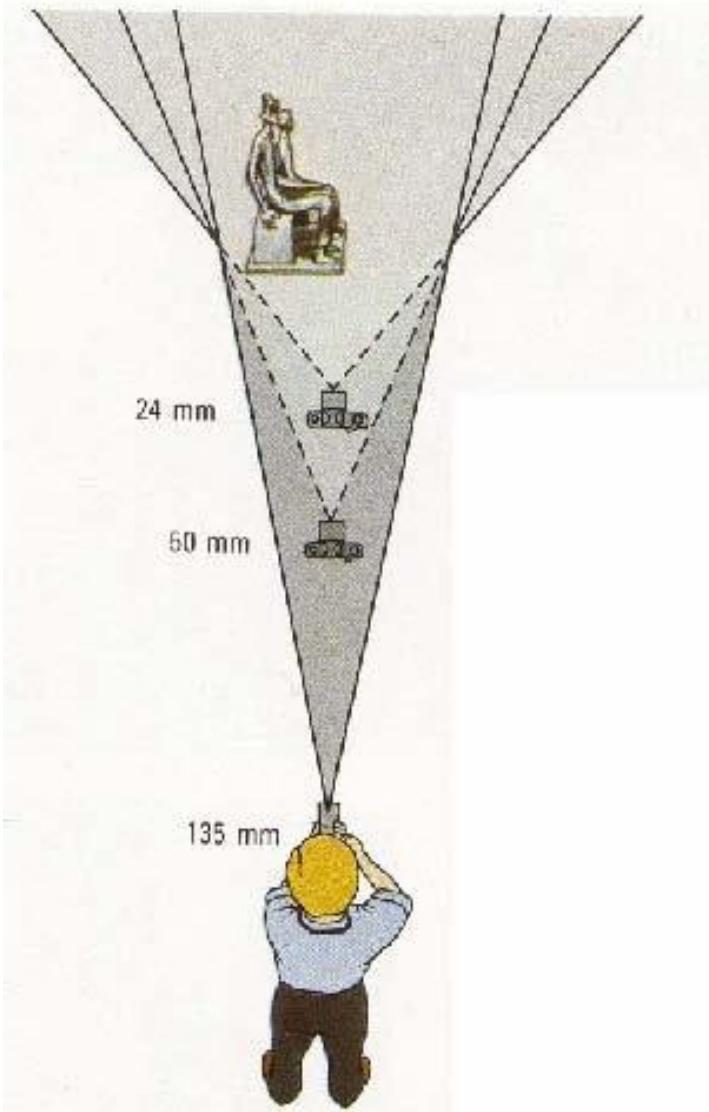


- Large f results in a telescopic image
(small FOV) →
 - Smaller portion of 3D scene is projected to the sensor.

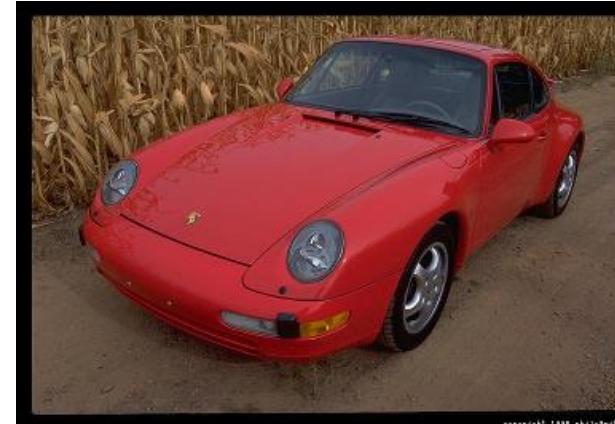


$$\varphi = \tan^{-1}\left(\frac{d}{2f}\right)$$

Field of view and focal length



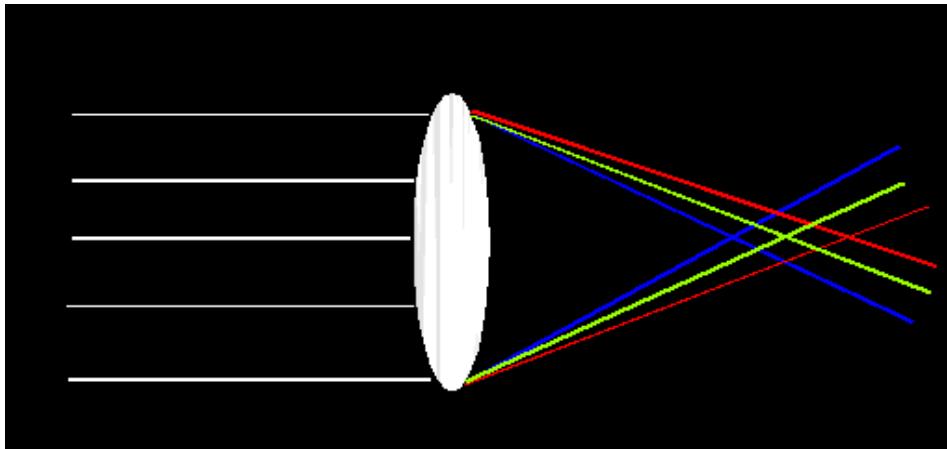
Large FOV, small f
Camera close to the car



Small FOV, large f
Camera far away from the car

Chromatic aberration

- Different wave-lengths refract at different angle and focus at slightly different distances:



Close to image center

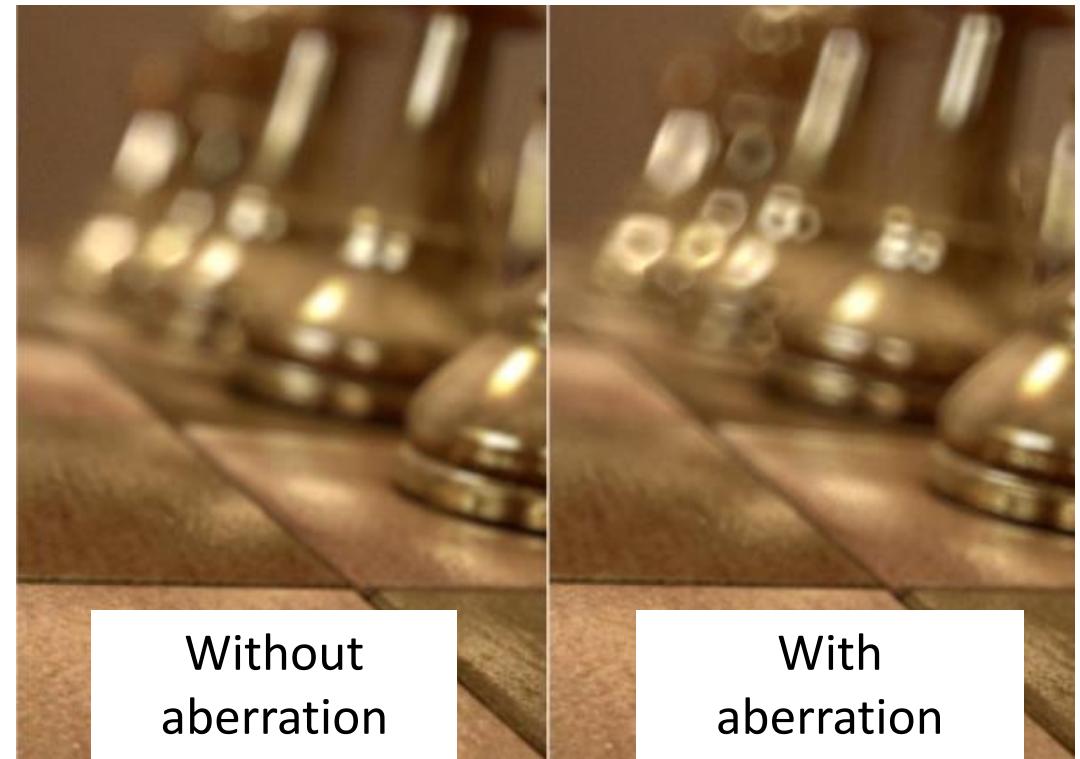
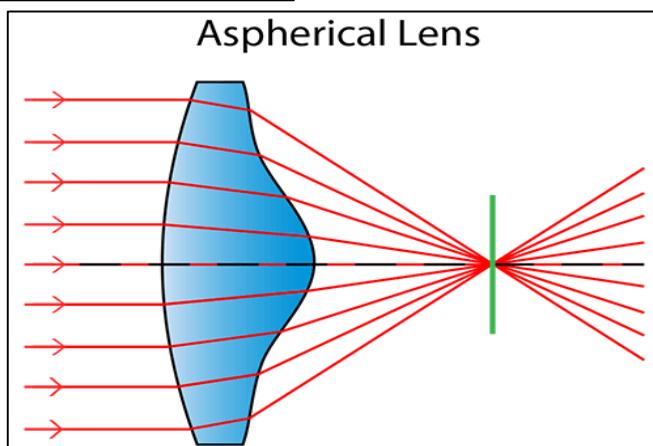
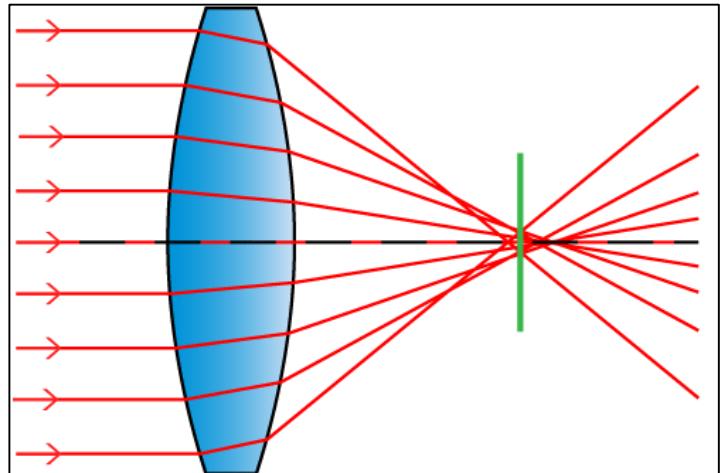


Close to image edge



Spherical aberration

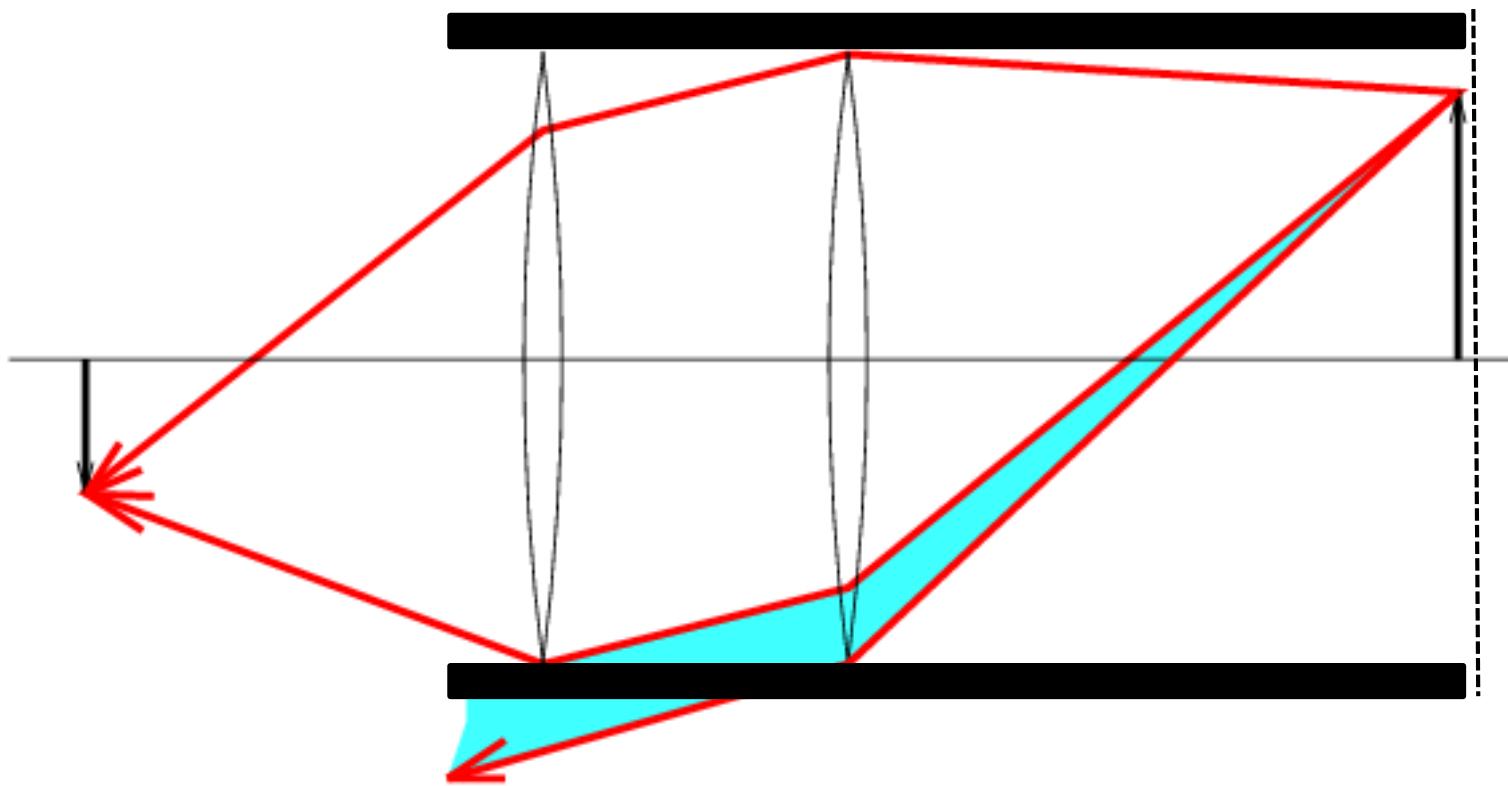
- Spherical lenses **do not focus** the light perfectly.
- Rays close to lens **edge** focus **closer** than those at the **center**.



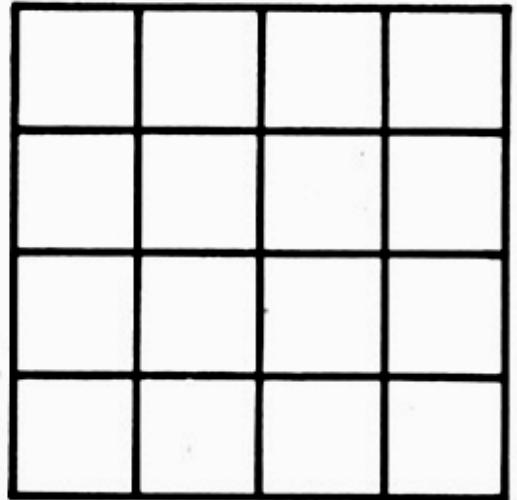
<http://photographylife.com/what-is-spherical-aberration>

<http://www.dofpro.com/sagallery.htm>

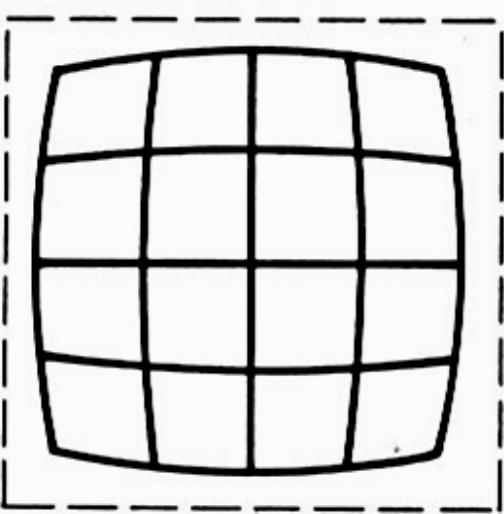
Vignetting



Radial distortion



Without distortion

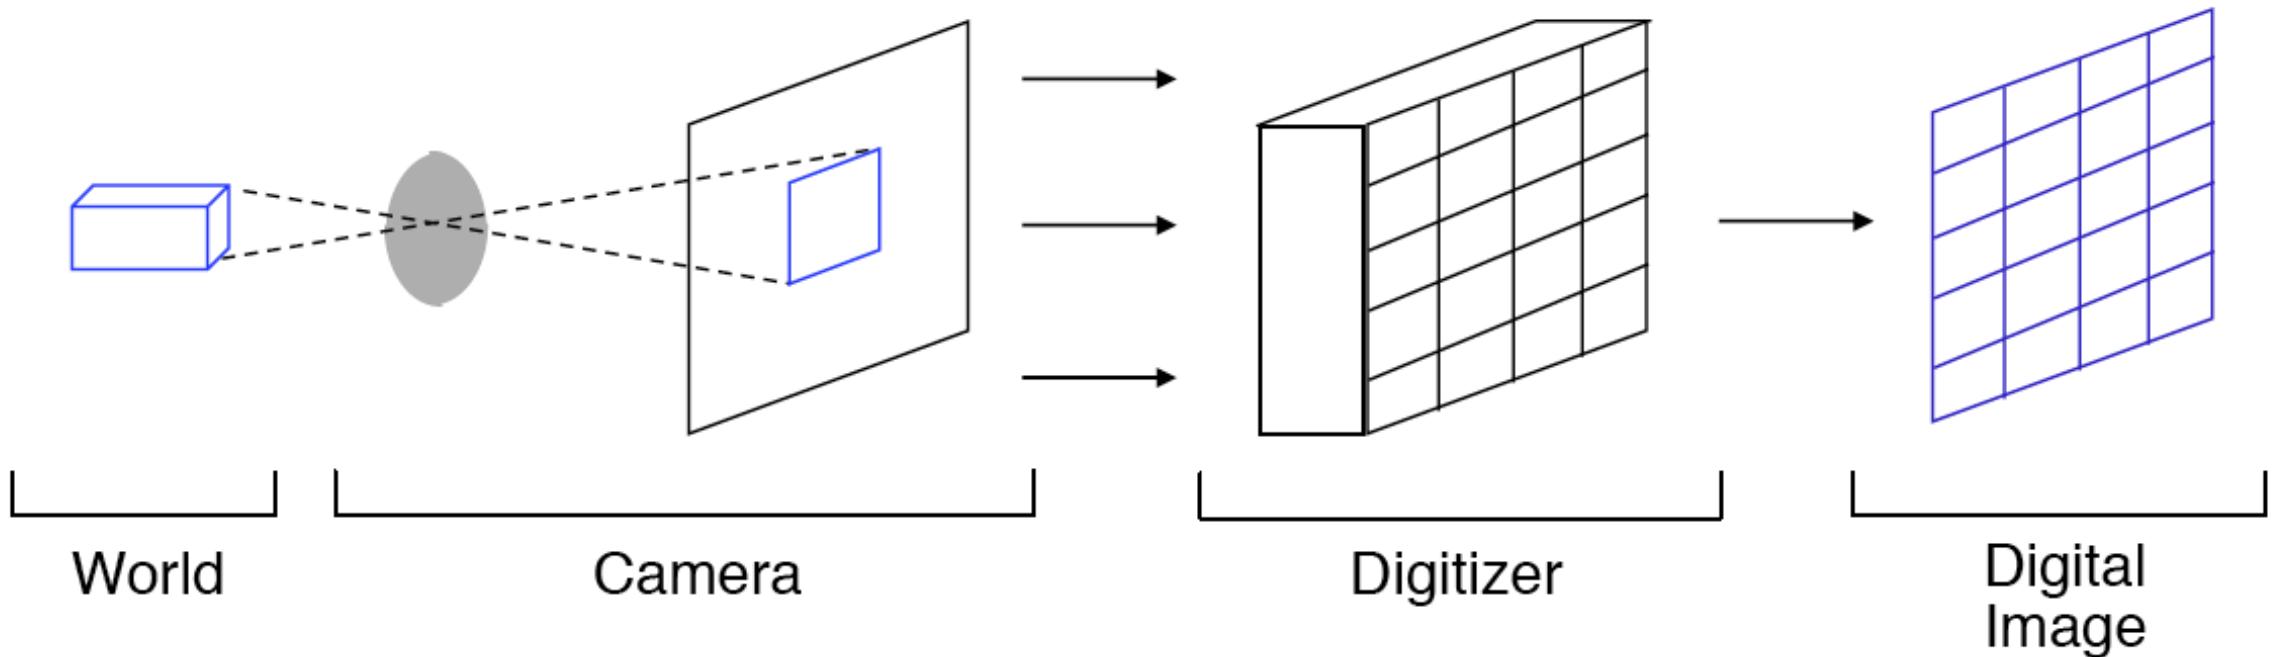


Barrel distortion



- Due to lens imperfections or fisheye.
- Most apparent at the edge of the image.

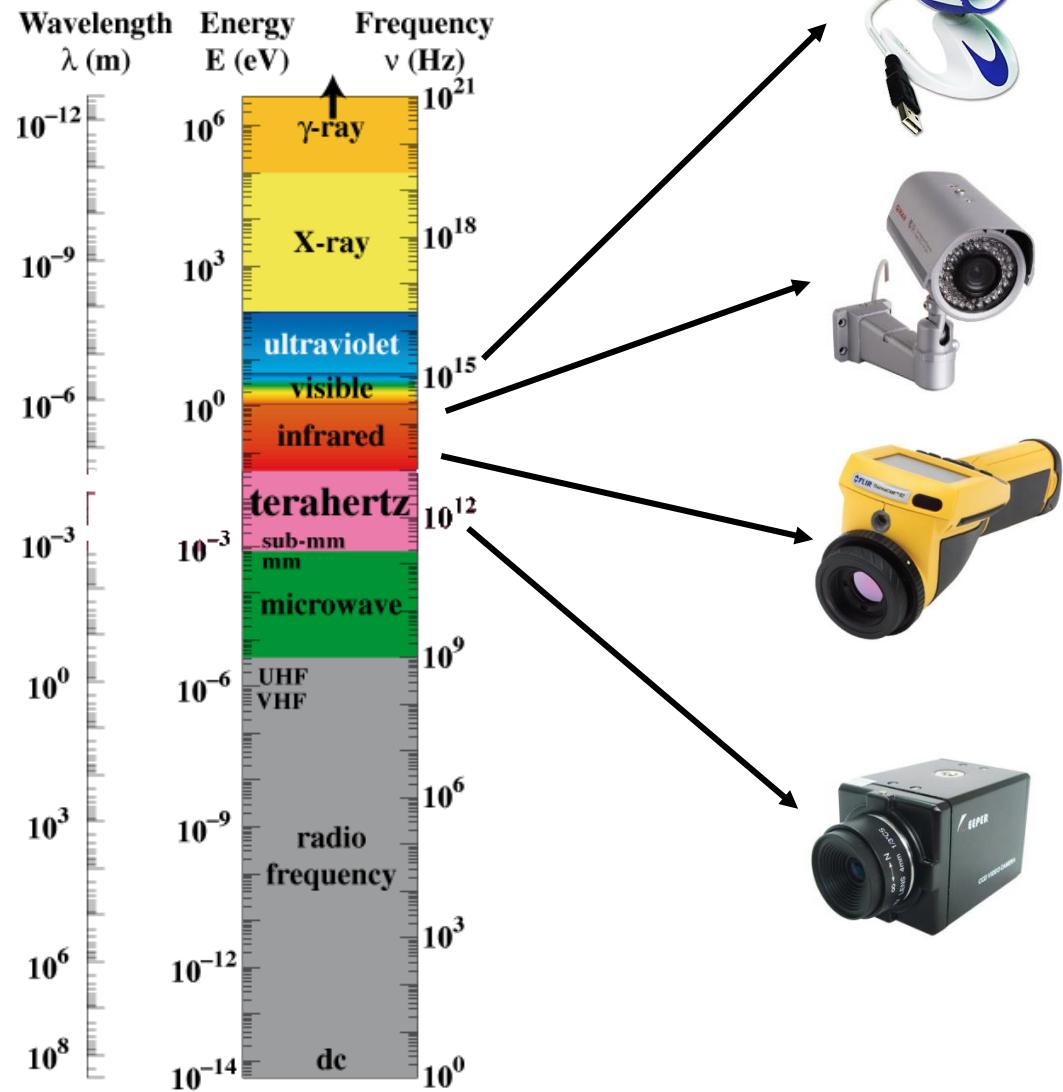
Digital image



- Instead of film, use matrix (array) of sensors.
- *Discretize* image into pixels.
- *Quantize* light into intensity levels.

Sensor: Camera

Electromagnetic spectrum



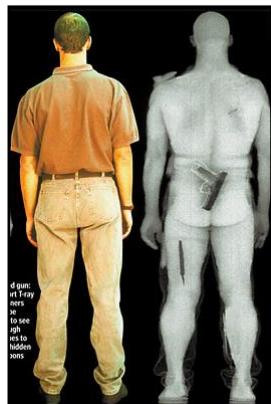
Visible light



Near-infrared light



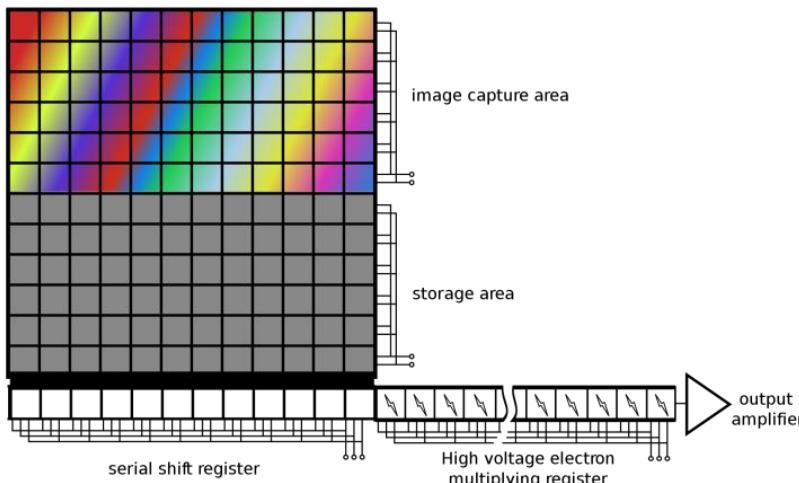
Far-infrared light



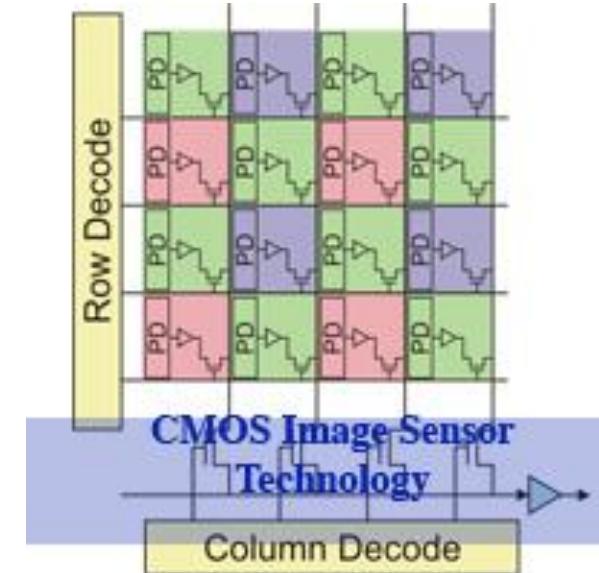
Terahertz light

Visible light cams: CCD vs CMOS

Charge coupled device (CCD)



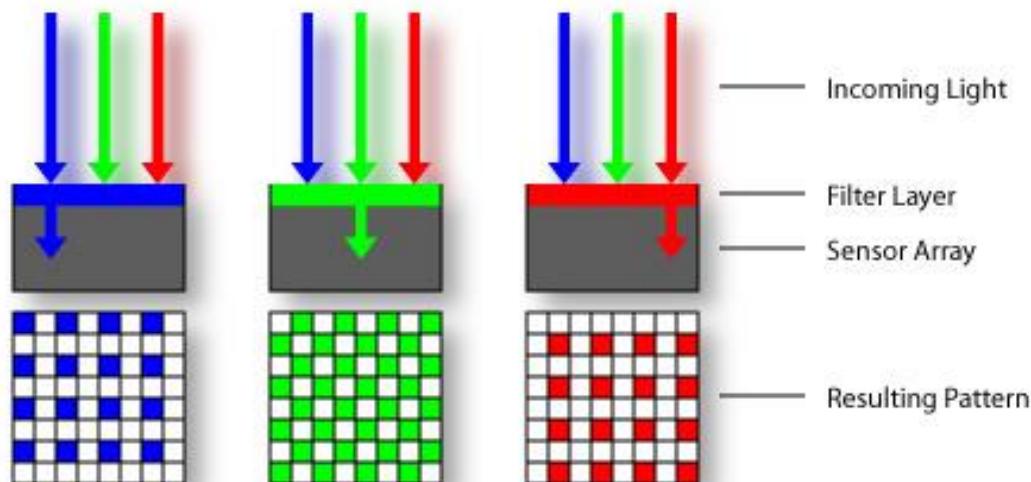
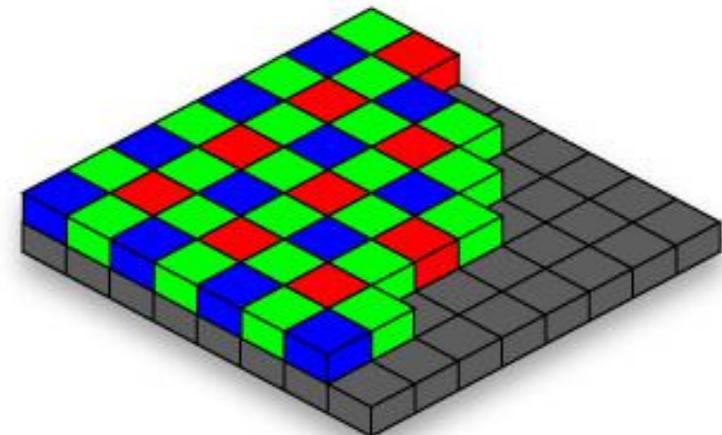
Complementary metal–oxide–semiconductor (CMOS)



- In both: Photons **cause charge** on each sensor „cell“.
- CCD **reads out the charge** (FIFO) **serially** and digitizes.
- CMOS performs **digitization on each cell** separately.
- CCD delivers better images.
- CMOS is cheaper to produce and is thus wide-spread.

Color perception in digital cameras

Bayer sensor



Why twice as many greens compared to blue and red?

Luminance is mostly determined by the green values.

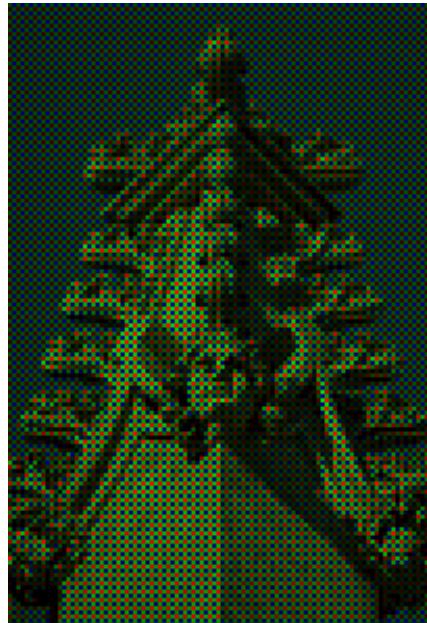
Human visual system much more sensitive to changes in intensity than in chroma (color).

Color perception in digital cameras

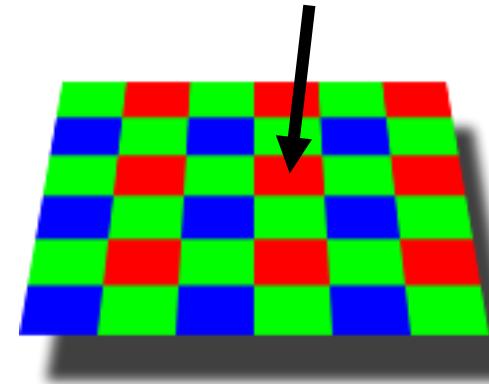
What you see



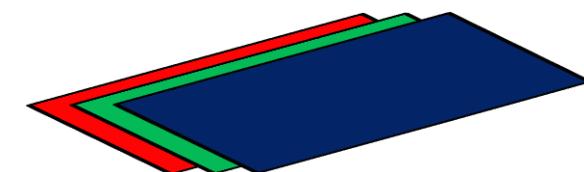
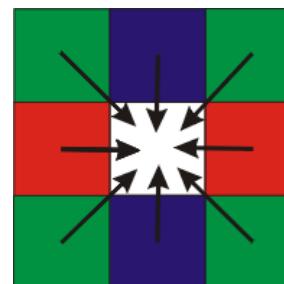
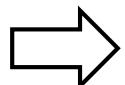
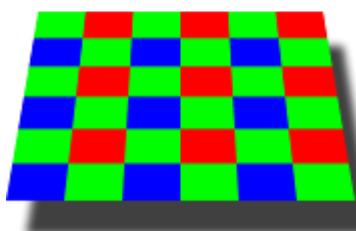
Your camera sees



Missing green!

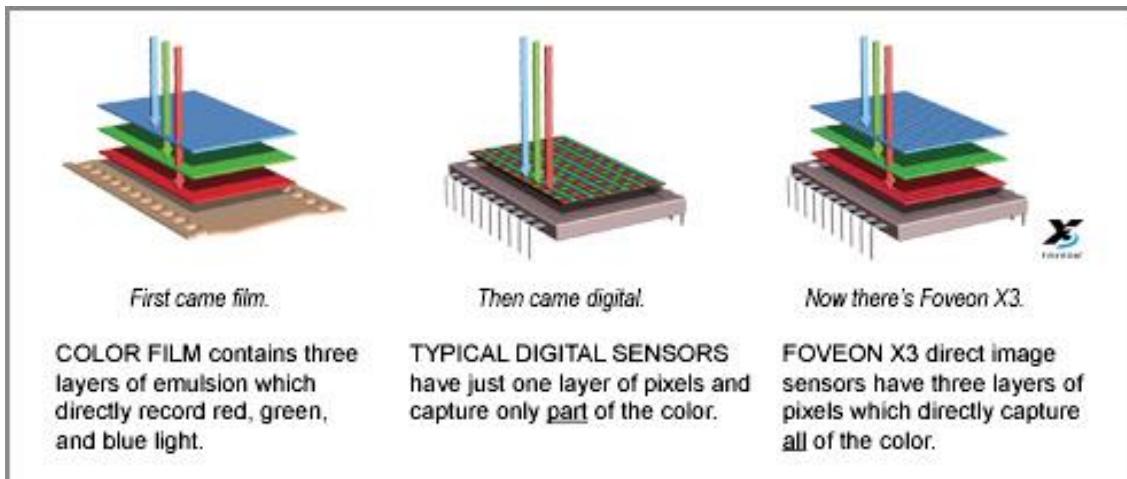


De-mosaicking: The missing color channels at a pixel need to be interpolated!

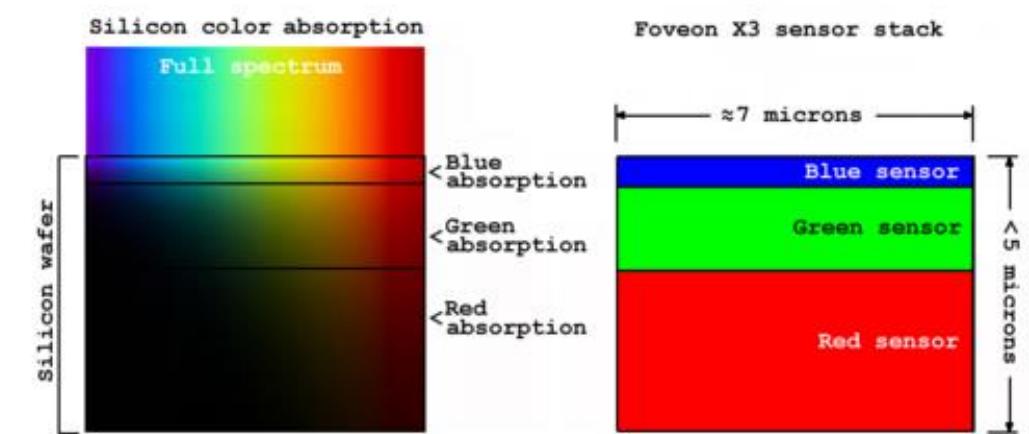


Color perception : Foveon X3

- CMOS-based sensor.
- Based on the fact, that red, green and blue color penetrate the silicon at different depths.



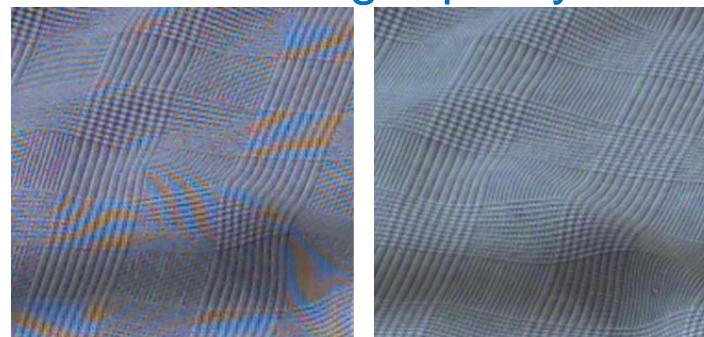
<http://www.foveon.com/article.php?a=67>



http://en.wikipedia.org/wiki/Foveon_X3_sensor

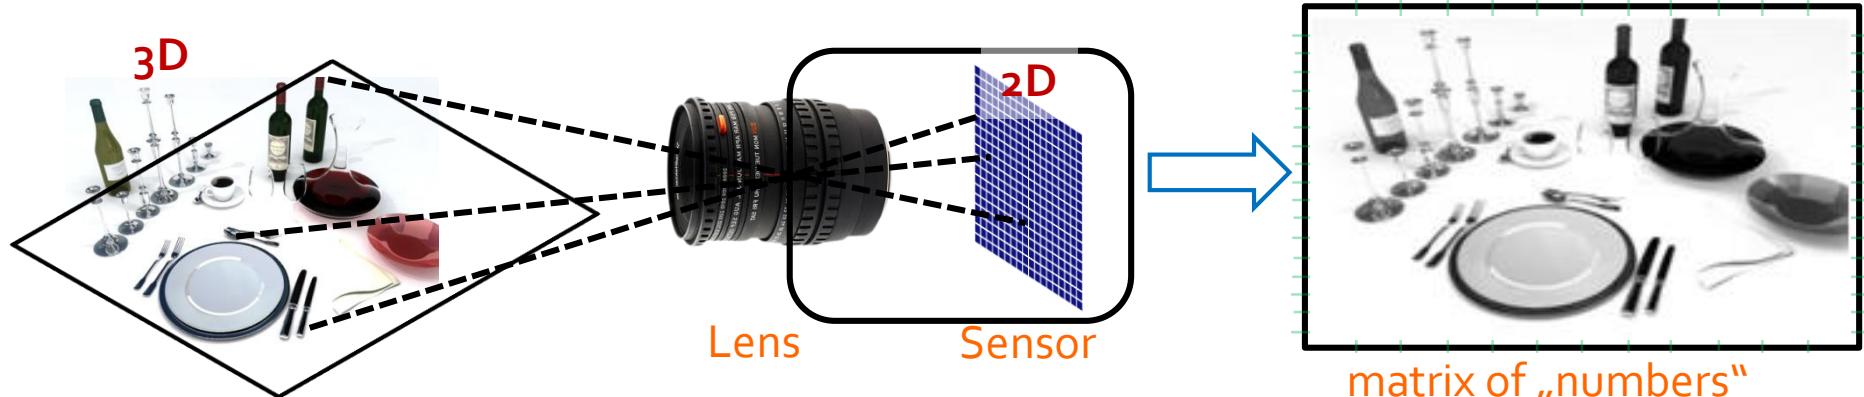
Better image quality

Bayer-like



Foveon X3

From camera to perception



- How does a human perceive the bottles, plates, forks,..., using only brightness?
- How do we perceive depth?
- Can a computer program do that?

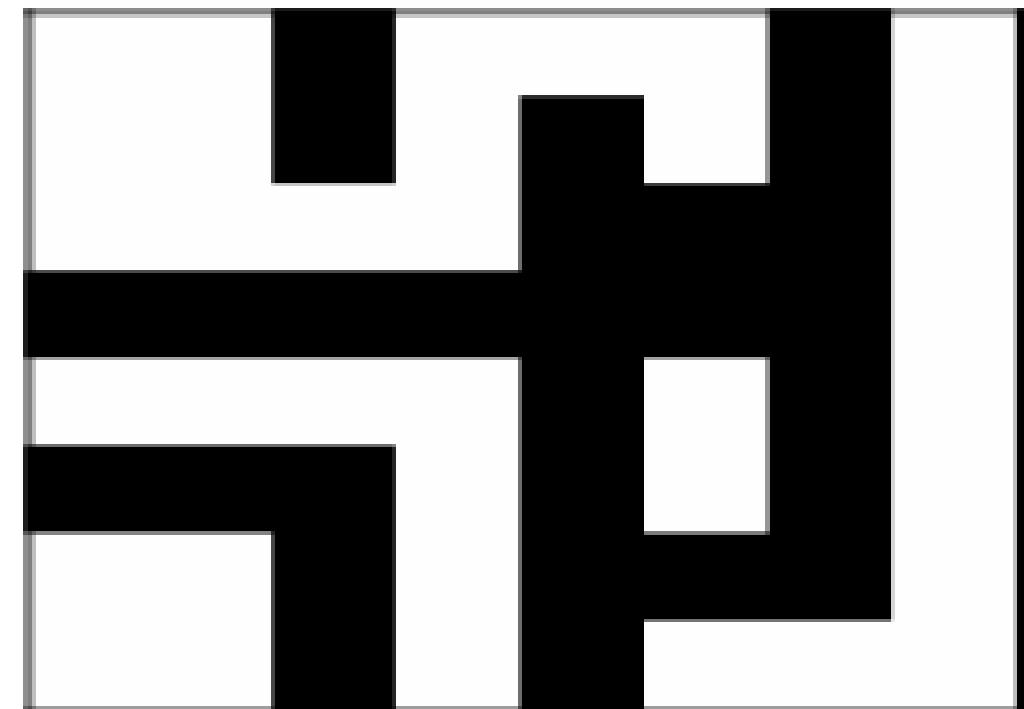
Machine perception

IMAGE PROCESSING 1

Binary images

- Only two possible gray levels
- Foreground vs. background

1	1	0	1	1	1	0	1
1	1	0	1	0	1	0	1
1	1	1	1	0	0	0	1
0	0	0	0	0	0	0	1
1	1	1	1	0	1	0	1
0	0	0	1	0	1	0	1
1	1	0	1	0	0	0	1
1	1	0	1	0	1	1	1



Usage: Machine vision

Fig. 3 Schematic diagram of marking inspection setup at Texas Instruments

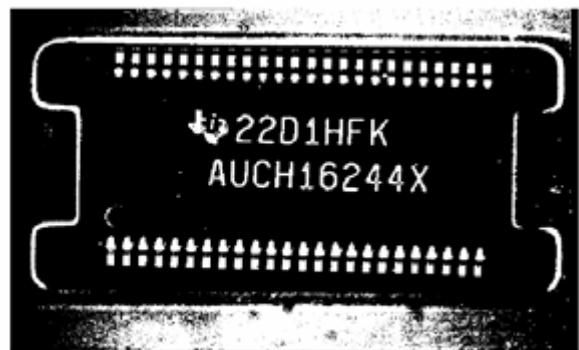
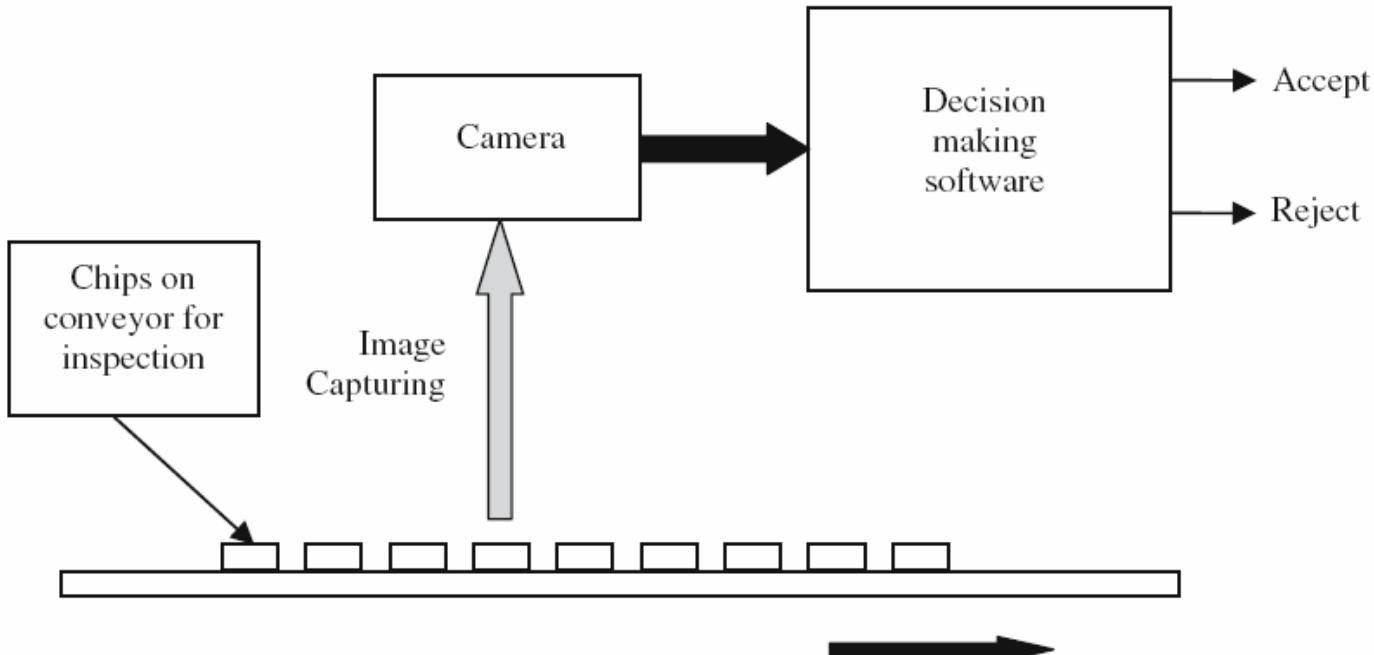
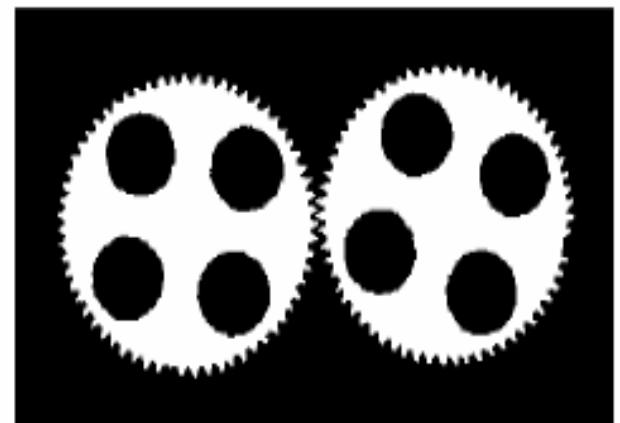


Fig. 7 Binarized image



Fig. 9 Row sum for separating a row

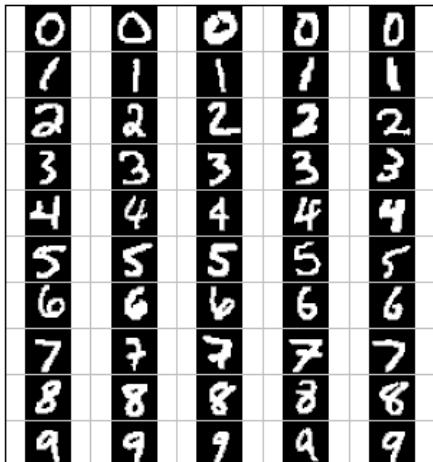


Source: Bastian Leibe

Usage : Text analysis



OCR on documents



Hand written numbers

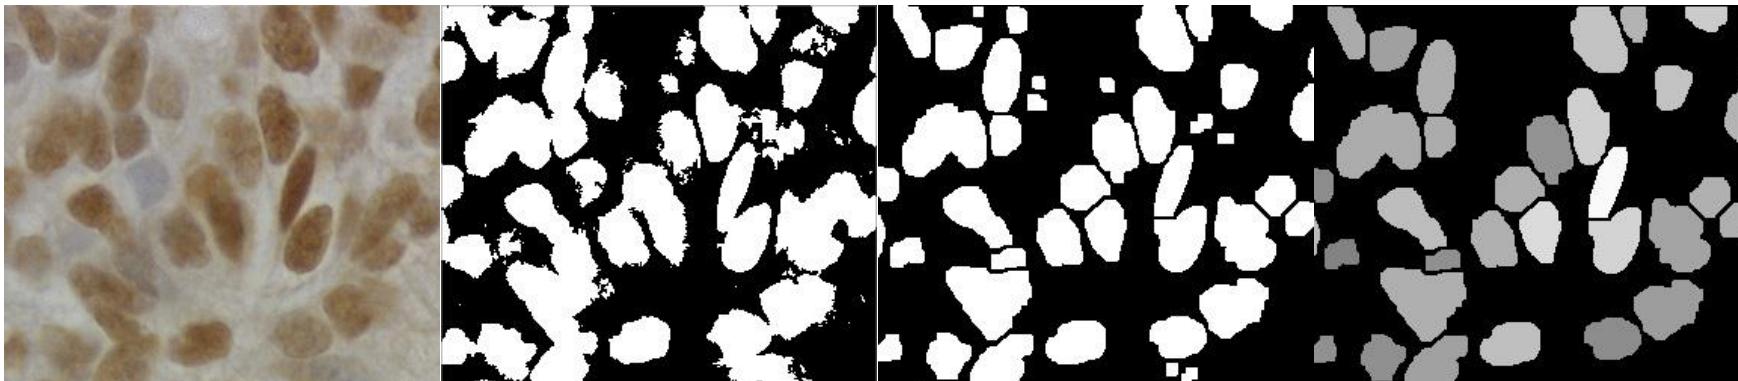
Text in the wild: after detection



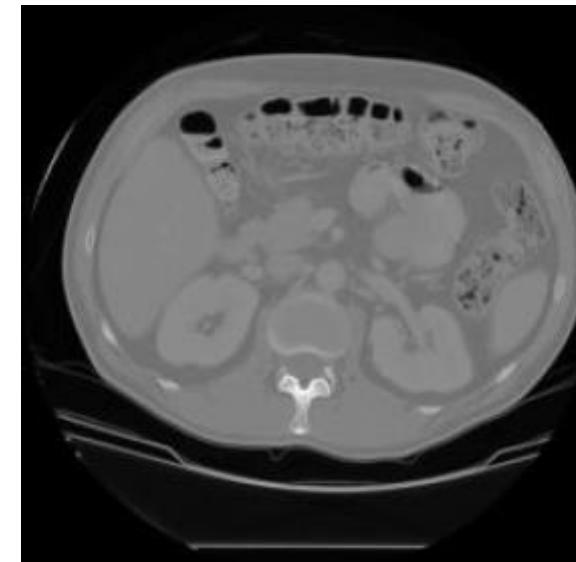
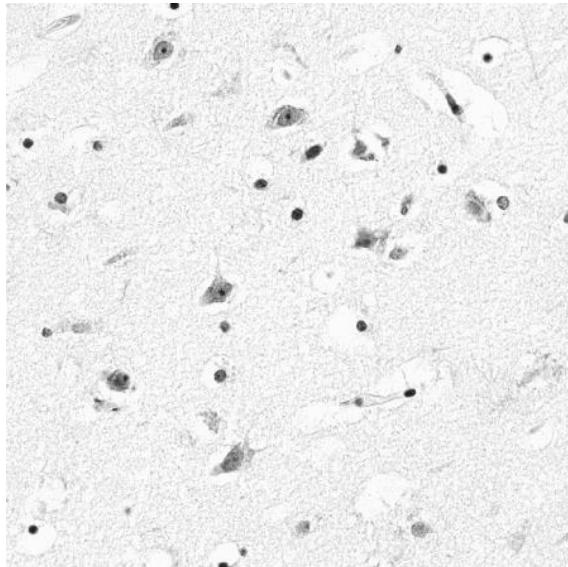
Source: Till Quack, Martin Renold

Source: Bastian Leibe

Usage: Medical imaging



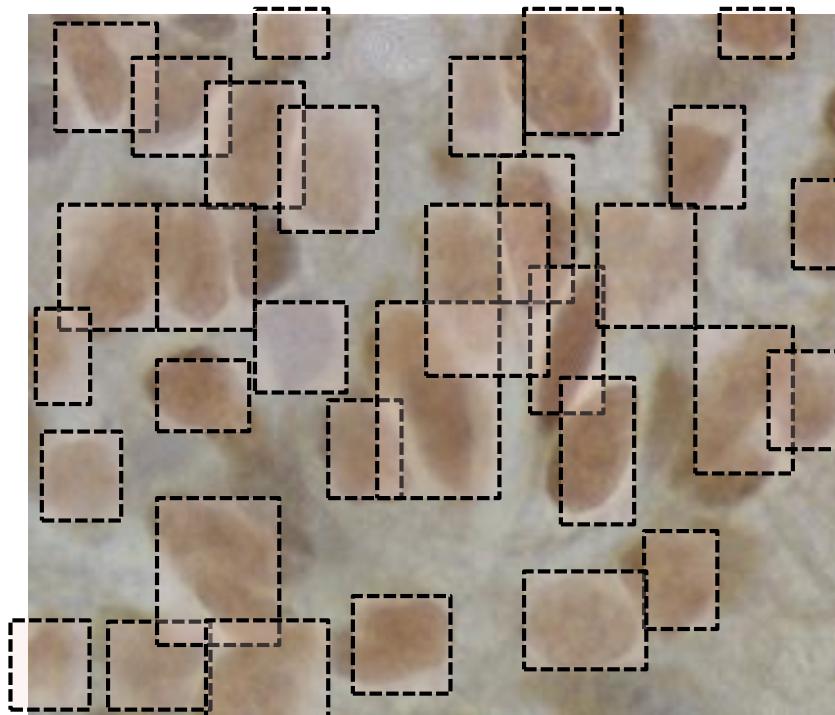
Source: D. Kim et al., Cytometry 35(1), 1999



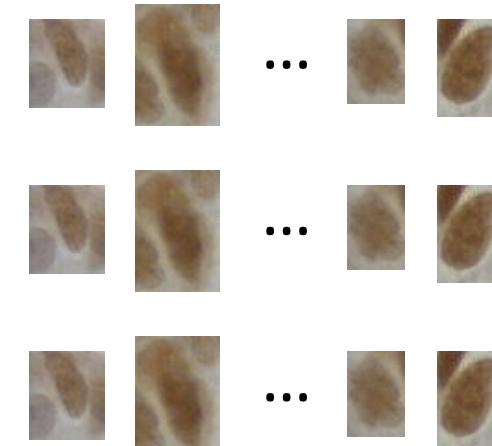
Source: Bastian Leibe

Task: Detect “round” cells

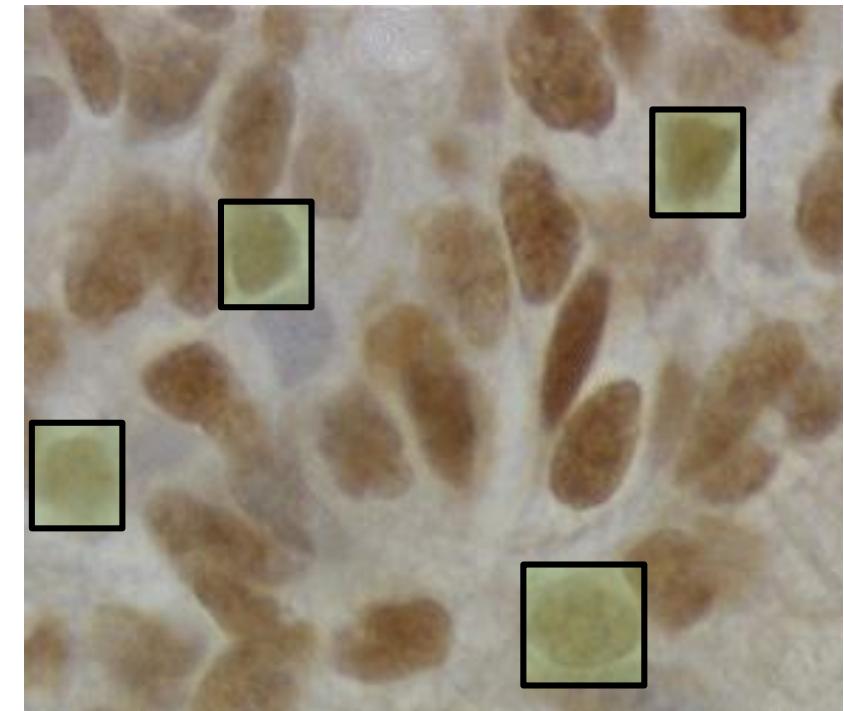
Generate hypotheses



Classify each region
into a “round” and
“not round”



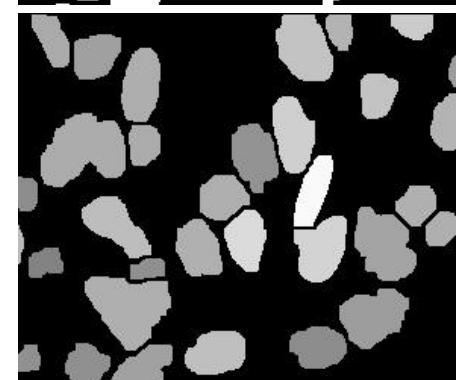
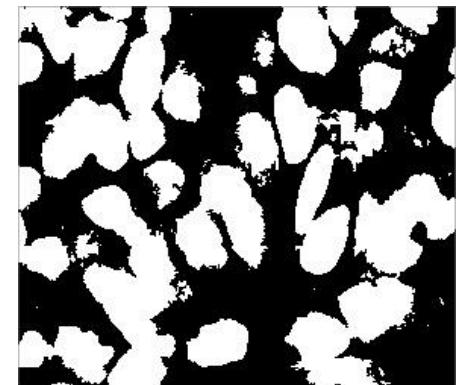
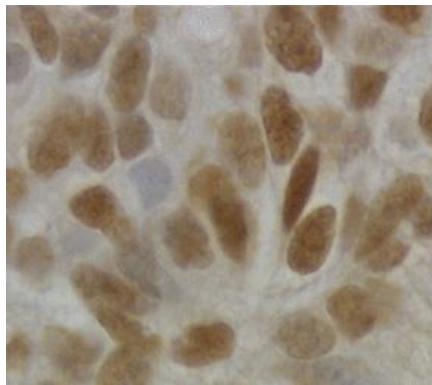
Keep “round” regions



Localize, Describe, Classify

Localize: Sequence of processing steps

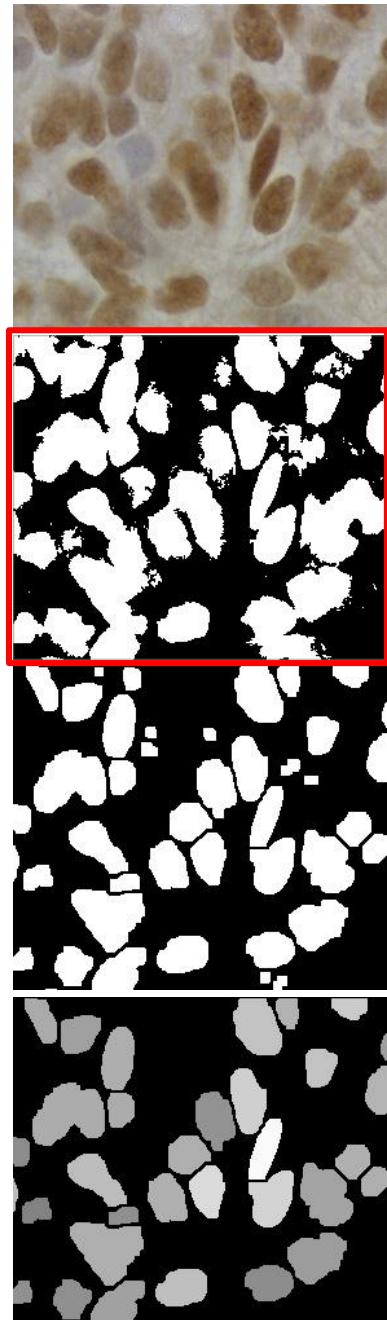
- Convert gray image to a binary image
 - Thresholding
- Clean binary image
 - Morphologic filtering
- Extract individual regions
 - Connected components



... then describe each localized region and classify

Machine perception

IMAGE THRESHOLDING



Thresholding

- Change image into a Binary Mask
- Various approaches
 - Apply a **single** threshold

$$F_T[i, j] = \begin{cases} 1, & \text{if } F[i, j] \geq T \\ 0, & \text{otherwise} \end{cases}$$

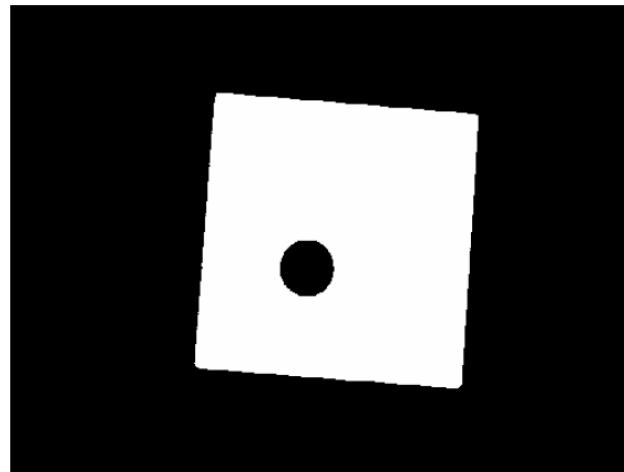
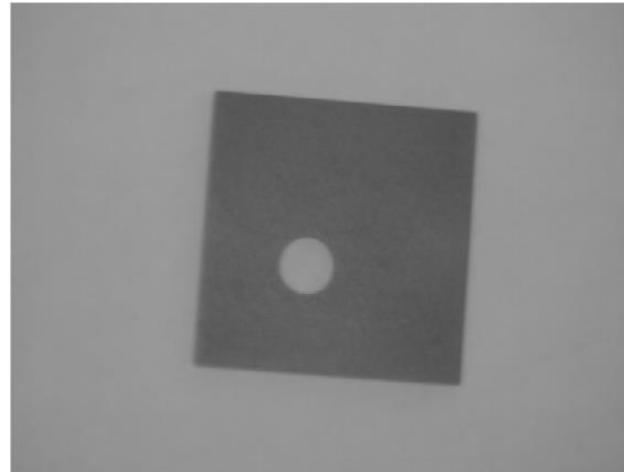
- Apply **two** thresholds

$$F_T[i, j] = \begin{cases} 1, & \text{if } T_1 \leq F[i, j] \leq T_2 \\ 0, & \text{otherwise} \end{cases}$$

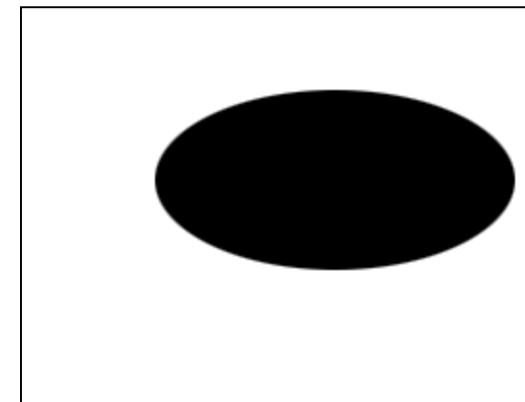
- Apply a **general** classifier

$$F_T[i, j] = \begin{cases} 1, & \text{if } F[i, j] \in Z \\ 0, & \text{otherwise} \end{cases}$$

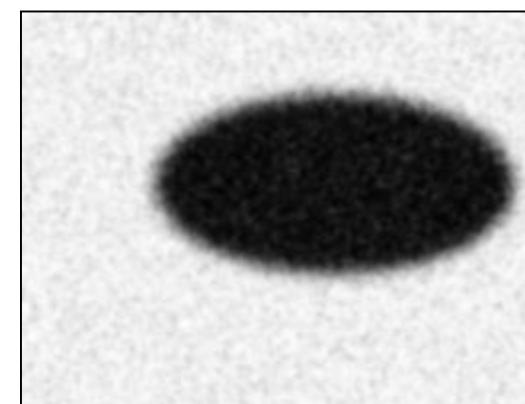
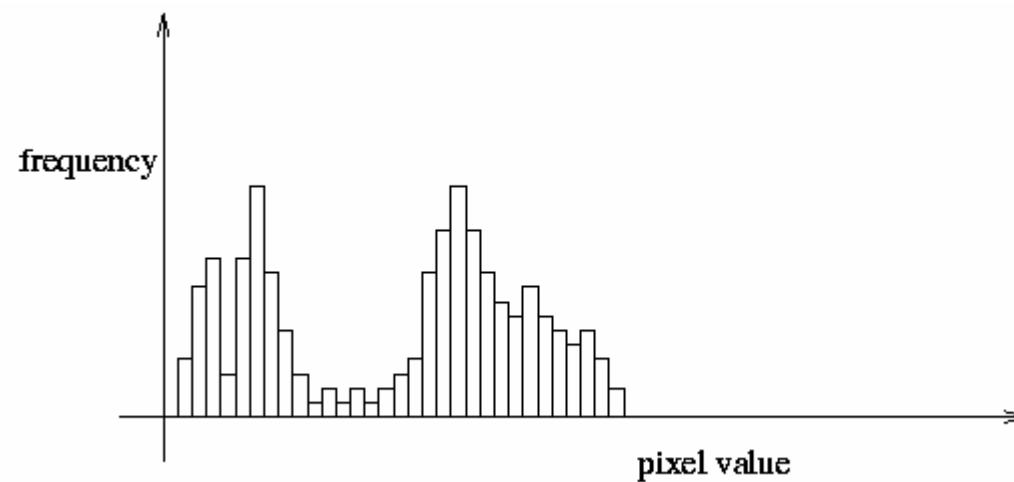
Object/background separation



A simple example: Bimodal histogram



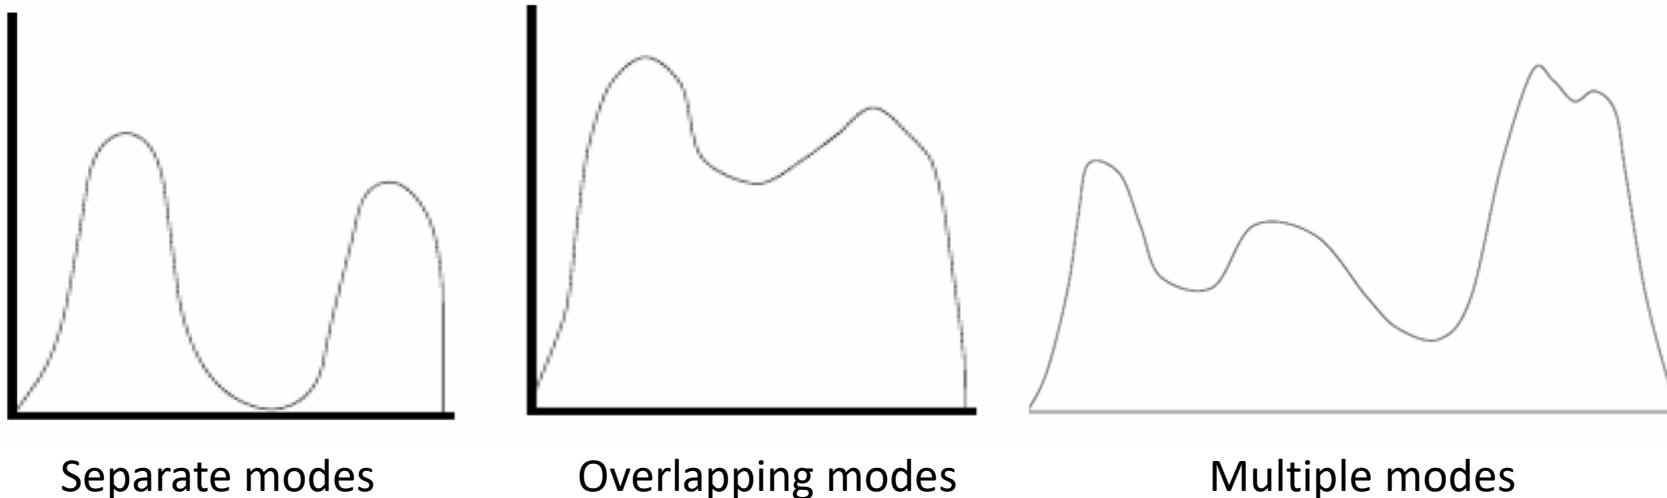
Ideal case:
bright object on
dark background.



A more realistic noisy
image.

A not so simple example...

- What to do here?



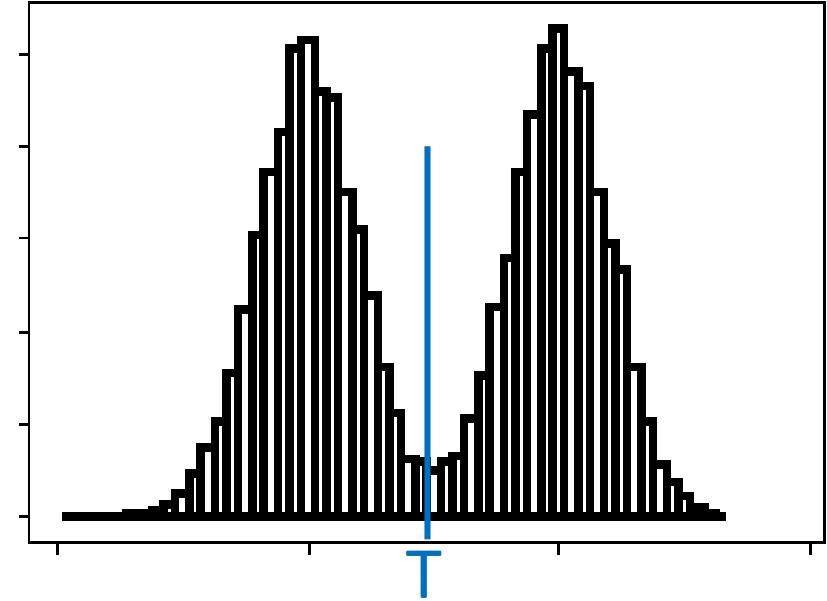
- Generally thresholding is a **difficult** problem
 - Domain knowledge helps a great deal.
 - E.g., the portion on letters on a page.
 - E.g., size of the structure we want to detect...

Global binarization [Otsu '79]

- Find a threshold T , that minimizes intensity variances within classes separated by T :

$$\sigma_{within}^2(T) = n_1(T)\sigma_1^2(T) + n_2(T)\sigma_2^2(T)$$

$$n_1(T) = |\{I_{(x,y)} < T\}|, n_2(T) = |\{I_{(x,y)} \geq T\}|$$



- This equals to maximization of between class variance $\sigma_{between}$:

$$\begin{aligned}\sigma_{between}^2(T) &= \sigma^2 - \sigma_{within}^2(T) \\ &= n_1(T)n_2(T)[\mu_1(T) - \mu_2(T)]^2\end{aligned}$$

Otsu, N (1979), "[A threshold selection method from gray-level histograms](#)", IEEE SMC

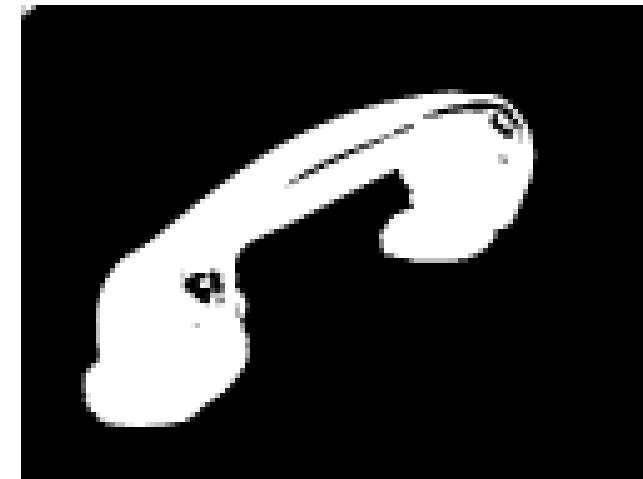
Otsu's Algorithm

For threshold value T

1. Separate the pixels into two groups by intensity threshold T
2. For each group get an average intensity and calculate $\sigma_{between}^2$.

Select the T^* , that maximizes the variance:

$$T^* = \arg \max_T [\sigma_{between}^2(T)]$$



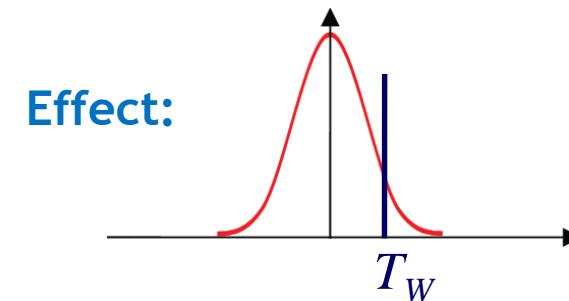
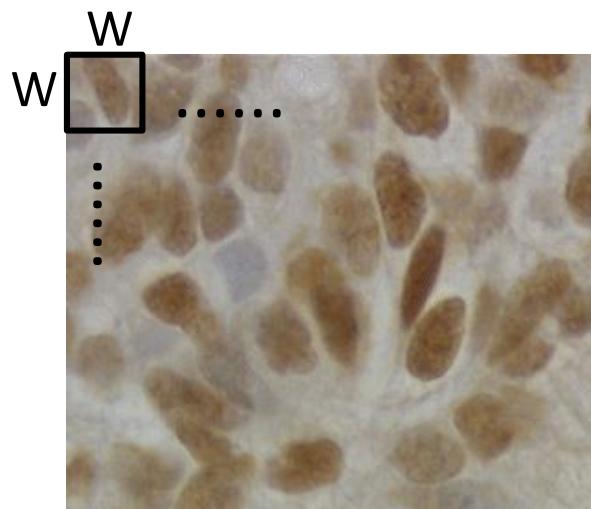
Local binarization [Niblack'86]

- Estimate a local threshold in neighborhood W:

$$T_W = \mu_W + k \cdot \sigma_W$$

with $k \in [-1,1]$ set by user.

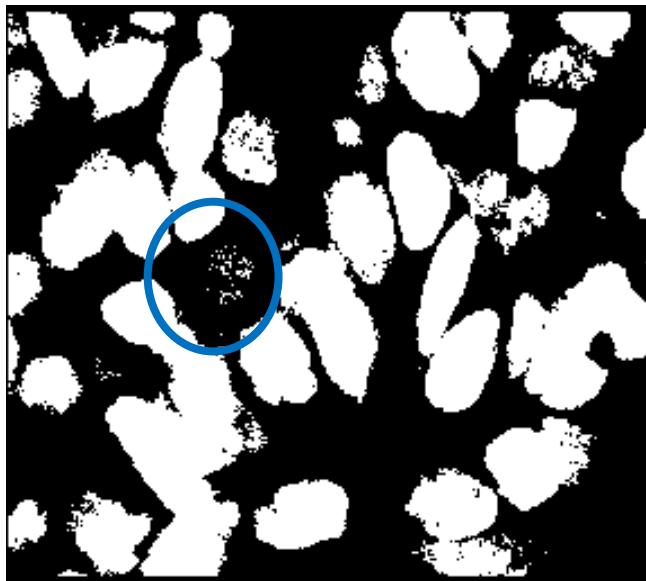
- Calculate the threshold separately for each pixel.



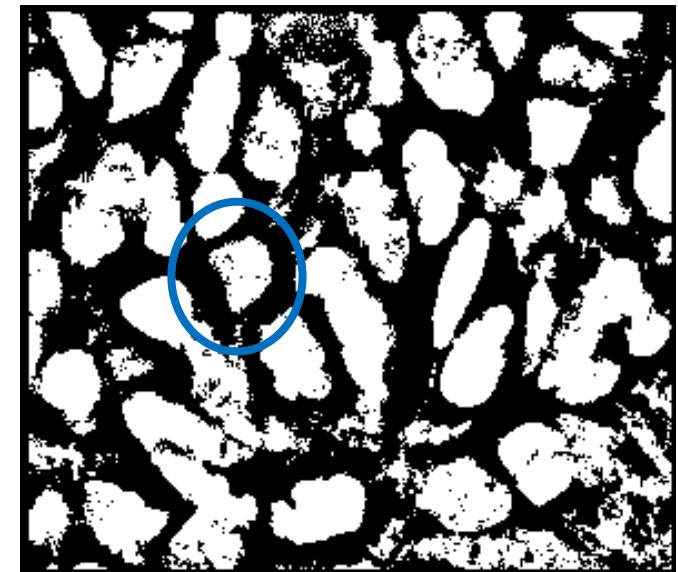
Examples of thresholding



Original



Global (Otsu)

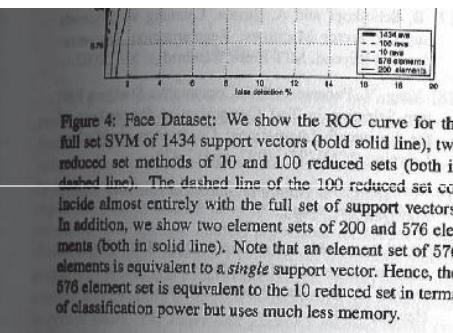


Local (Niblack)

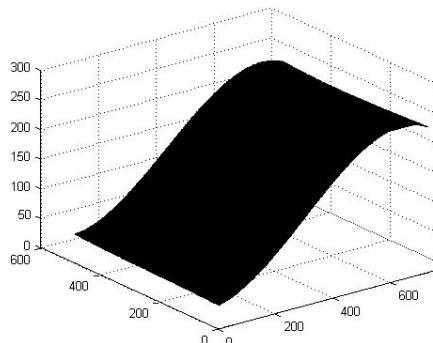
Additional improvements

- The shade in documents is often smooth...

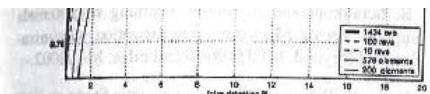
⇒ Try to model it by a polynomial!



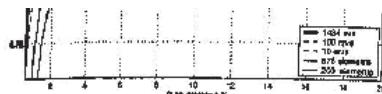
Original



Fitted surface



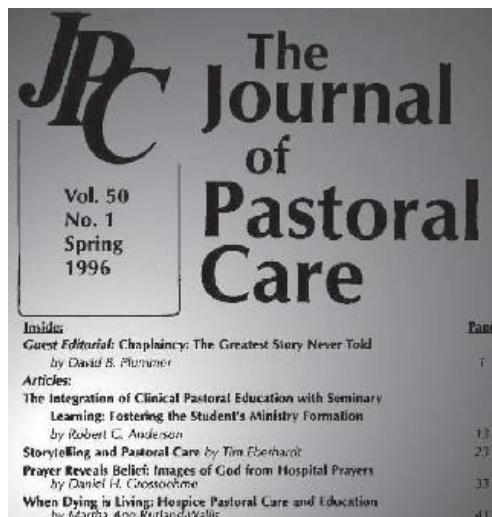
Shadow compensation



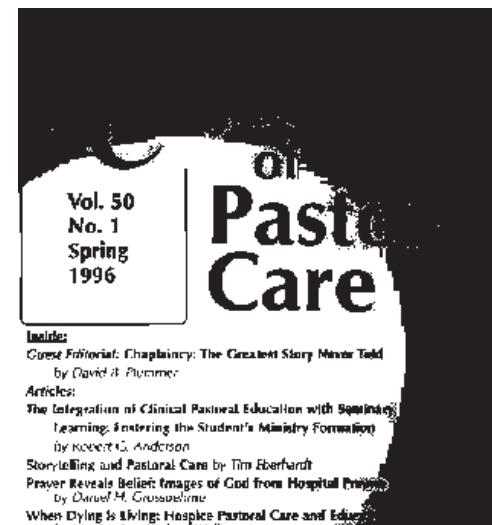
Binarized result

Comparison of results

Original image

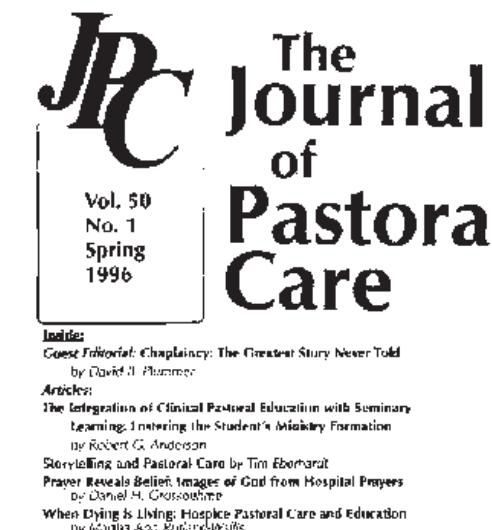


Local (Sauvola)



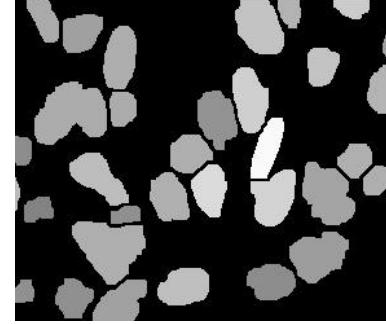
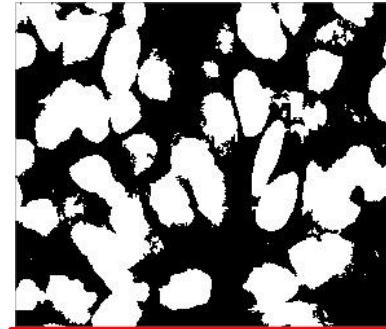
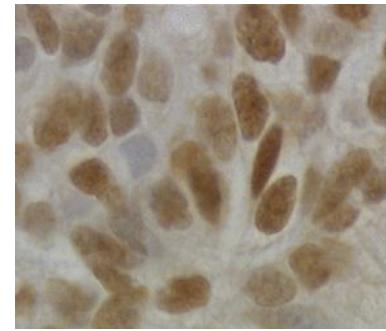
Global (Otsu)

Polynomial
+ global



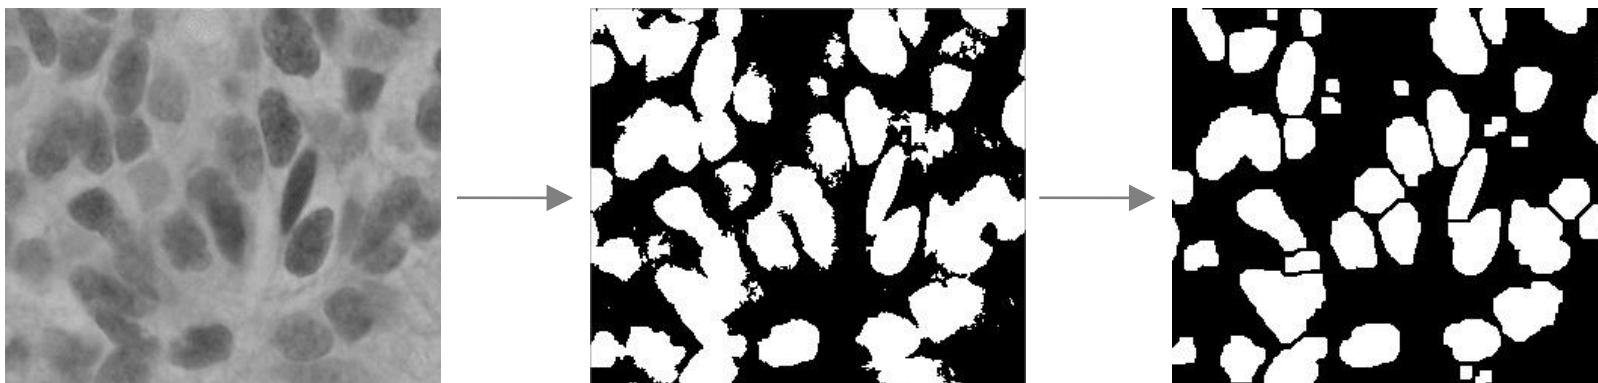
Machine perception

CLEANING THE IMAGE



Cleaning the binary image

- Thresholded image still includes noise



- Require post-processing to remove artefacts
- Morphological operators
 - Remove isolated points and small structures
 - Fill holes

Dilation: A sneak peak preview

- Dilate the regions of „white“ pixels
- Increases the size of the structures
- Fills holes in regions



Before dilation



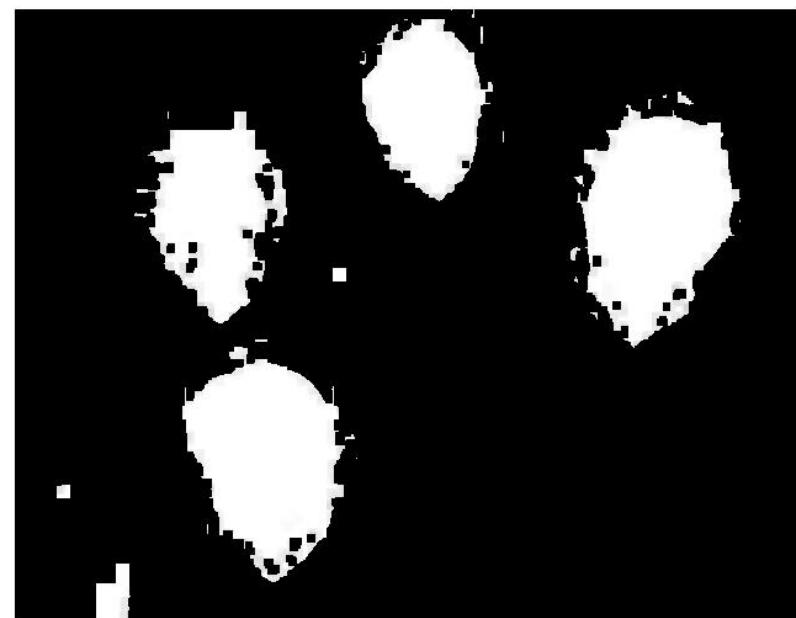
After dilation

Erosion: A sneak peak preview

- Erode the regions of „white“ pixels
- Reduce the size of structures
- Remove bridges, branches, noise



Before erosion



After erosion

Central to morphology: Structuring element (SE)

- Can be any shape and content:

1	1	1
1	1	1
1	1	1

0	1	0
1	1	1
0	1	0

0	0	1	0	0
0	1	1	1	0
1	1	1	1	1
0	1	1	1	0
0	0	1	0	0

Origin of the SE

- Fit: All “1” pixels in SE cover “1” pixels in the image.
- Hit: Any “1” pixels in SE cover “1” pixels in the image.

	0	1	2	3	4	5	6	...
0	0	0	0	0	0	0	0	0
1	0	0	0	1	1	0	0	0
2	0	0	1	1	1	1	1	0
3	0	1	1	1	1	1	1	1
4	0	1	1	1	1	1	1	1
5	0	0	1	1	1	1	1	1
:

Image courtesy: Brian Mac Namee

Fitting & Hitting

0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	0	0	0	0	0	0	0	0	0
0	0	1	B	1	1	1	0	C	1	0	0	0	0
0	1	1	1	1	1	1	1	0	0	0	0	0	0
0	1	1	1	1	1	1	1	0	0	0	0	0	0
0	0	1	1	1	1	1	1	0	0	0	0	0	0
0	0	1	1	1	1	1	1	1	0	0	0	0	0
0	0	1	1	1	1	1	A	1	1	1	0	0	0
0	0	0	0	0	1	1	1	1	1	1	1	1	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0

0	1	0
1	1	1
0	1	0

Structuring
Element

Fit / Hit?

A:

B:

C:

Fit : All “1” elements in SE are 1

Hit: Any “1” element in SE is 1

Erosion

- Erosion of image f by structuring element s is given by $g = f \ominus s$.
- The structuring element s is positioned with its origin at (x, y) and the new pixel value is determined using the rule:

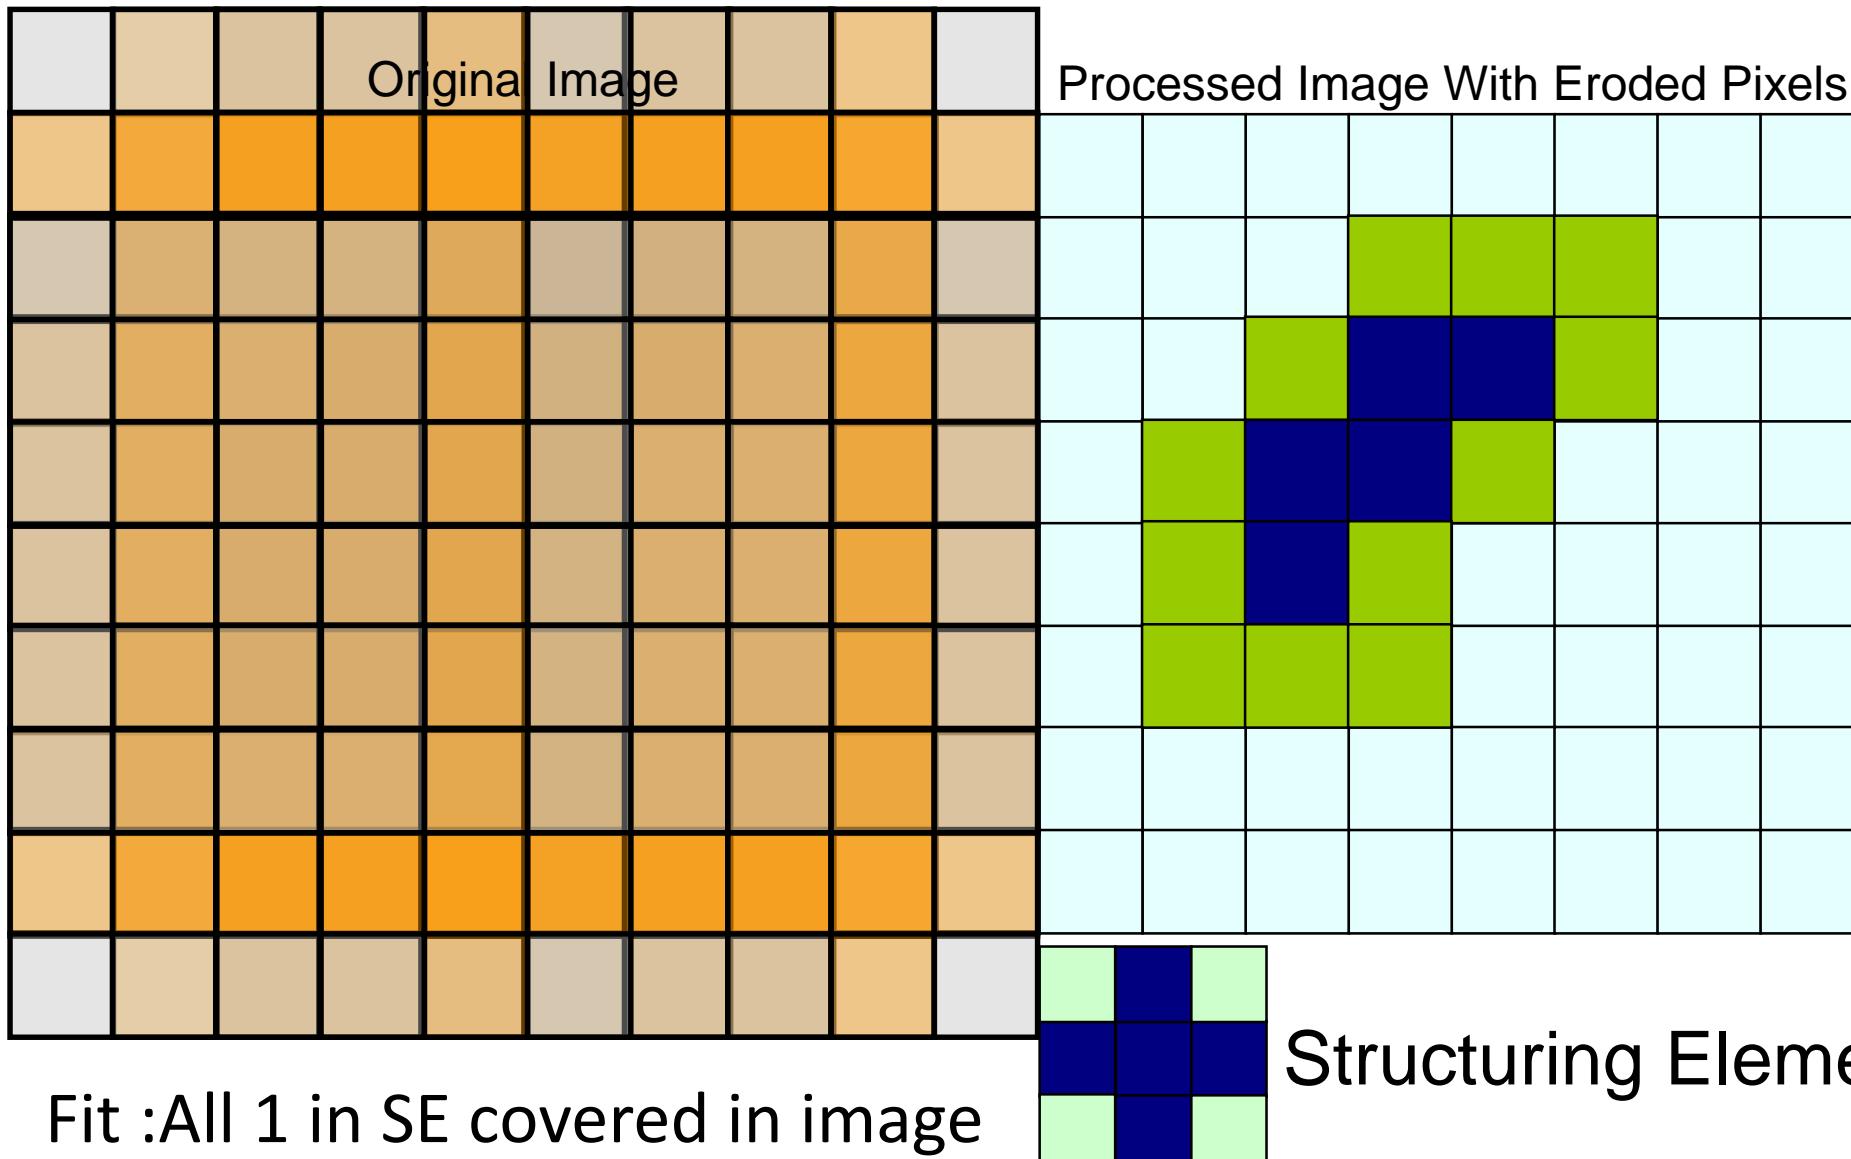
$$g(x, y) = \begin{cases} 1 & \text{if } s \text{ fits } f \\ 0 & \text{otherwise} \end{cases}$$

S	0	1	0
0	0	1	1
1	1	0	0
0	0	1	0

Fit: All “1” pixels in SE cover “1” pixels in the image.

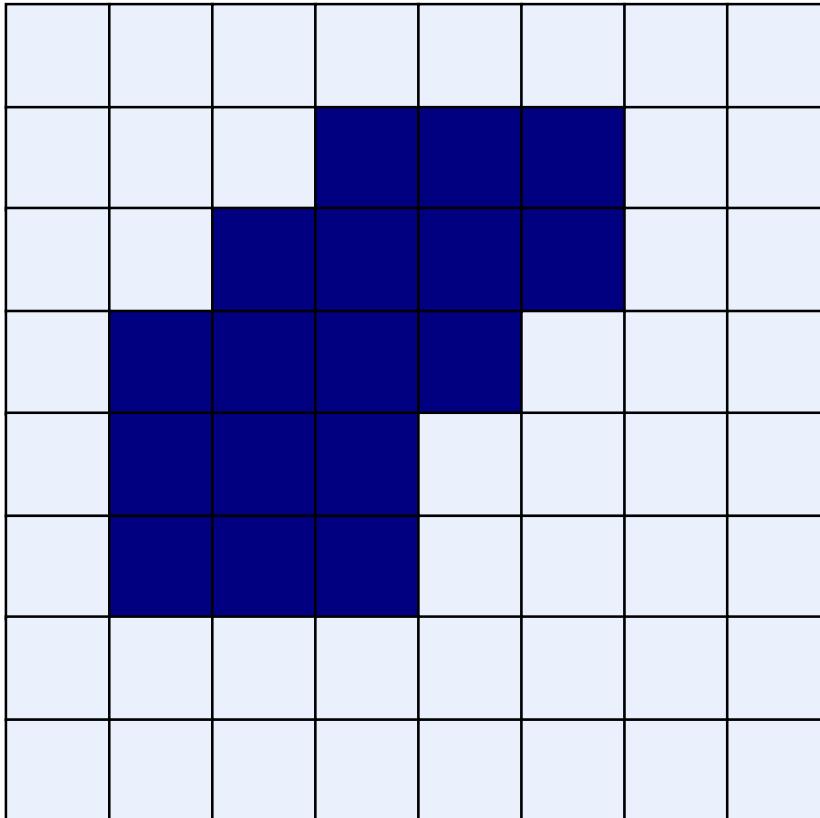
SE placed on image at (2,2)							
0	1	2	3	4	5	6	...
0	0	0	0	0	0	0	0
1	0	0	0	1	0	0	0
2	0	0	1	1	1	0	
3	0	1	1	1	1	1	1
4	0	1	1	1	1	1	1
5	0	0	1	1	1	1	1
:	⋮	⋮	⋮	⋮	⋮	⋮	⋮

Erosion Example



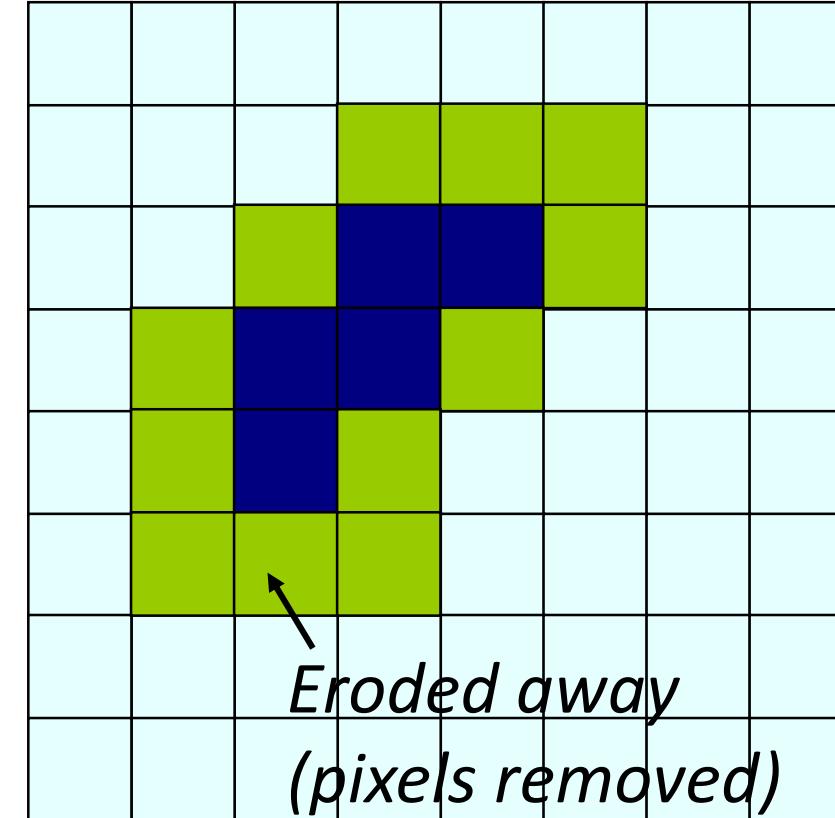
Erosion Example

Original Image

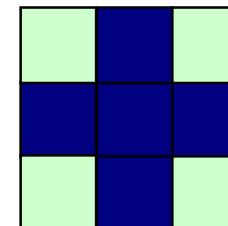


Fit :All 1 in SE covered in image

Processed Image With Eroded Pixels

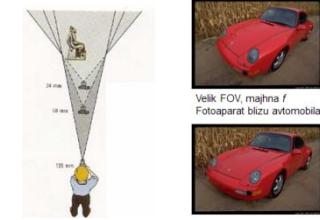
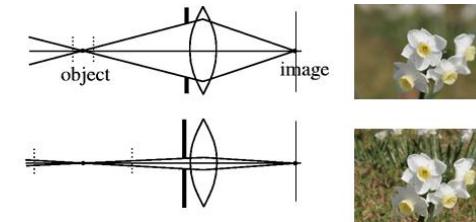
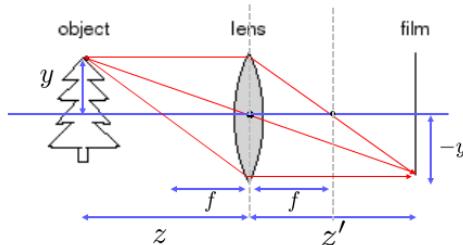
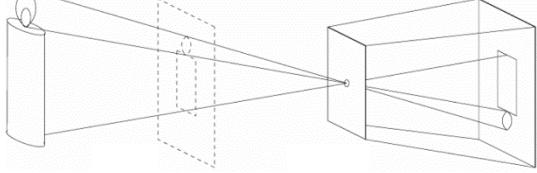


Structuring Element

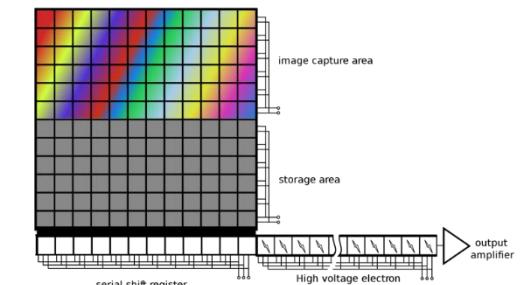
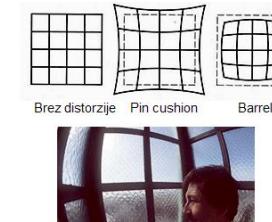
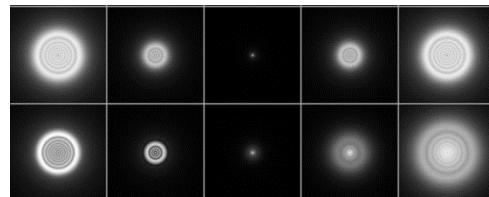
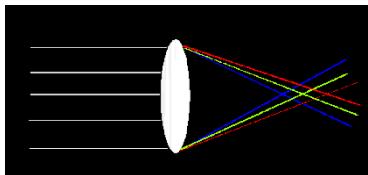


Previously at MP...

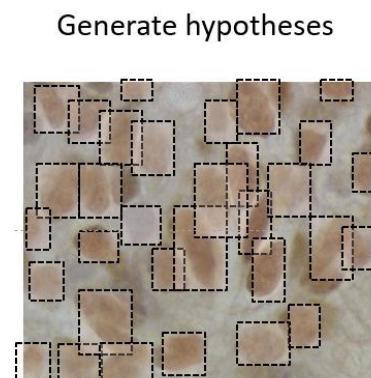
- Pinhole camera and image formation



- Image aberrations + digitization

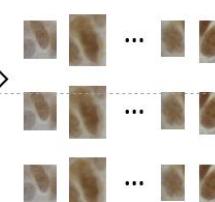


- Image processing 1

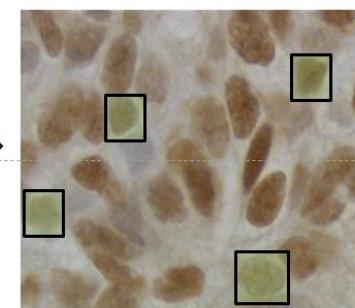


Generate hypotheses

Classify each region into a “round” and “not round”

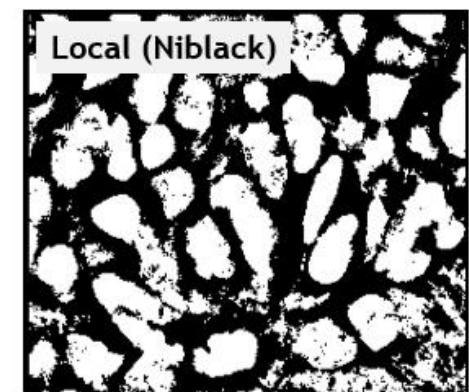
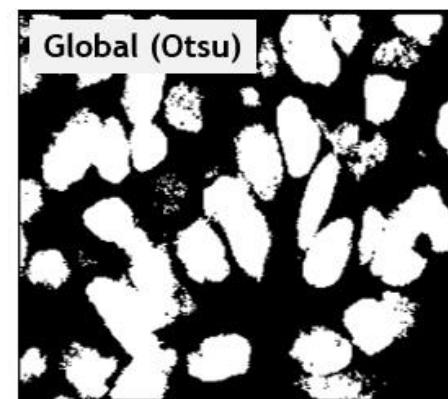
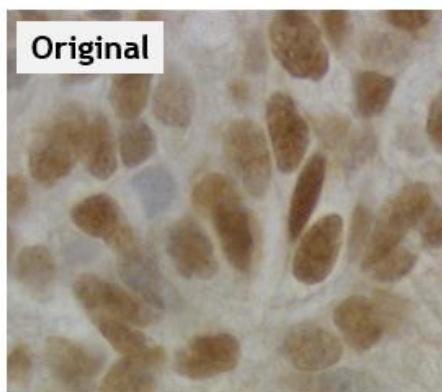
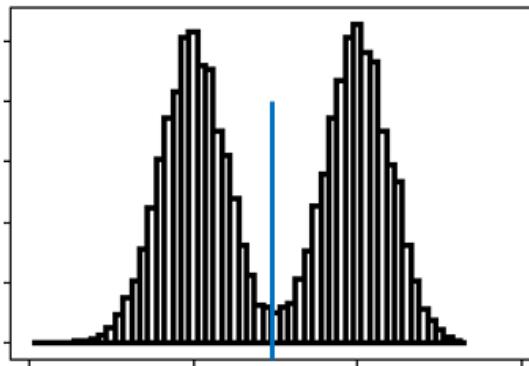


Keep “round” regions



Previously at MP...

- Separate foreground from background by thresholding

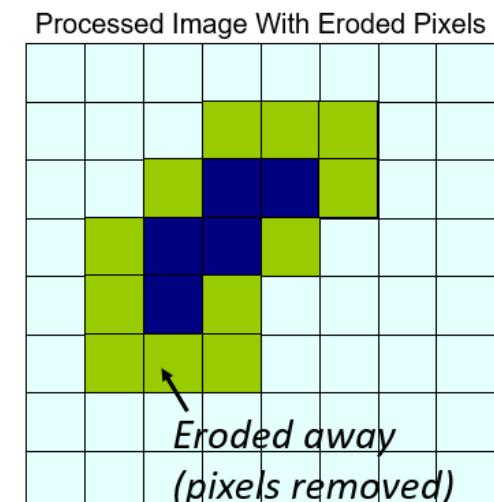
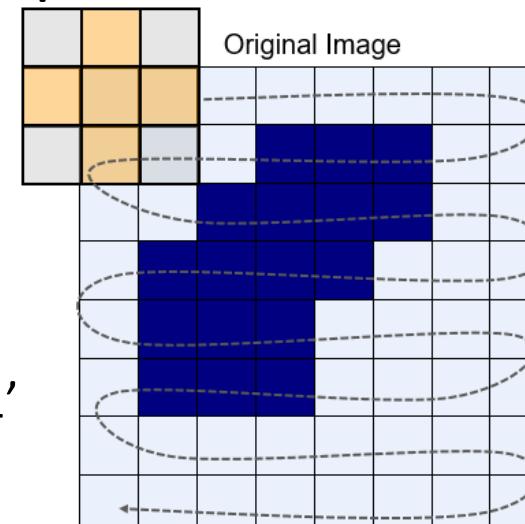


- Clean binary images by morphological operations



0	1	0
1	1	1
0	1	0

Fit: All “1” pixels in SE cover “1” pixels in the image.



Dilation

- Dilation of image f by structuring element s is given by $g = f \oplus s$.
- The structuring element s is positioned with its origin at (x, y) and the new pixel value is determined using the rule:

$$g(x, y) = \begin{cases} 1 & \text{if } s \text{ hits } f \\ 0 & \text{otherwise} \end{cases}$$

S

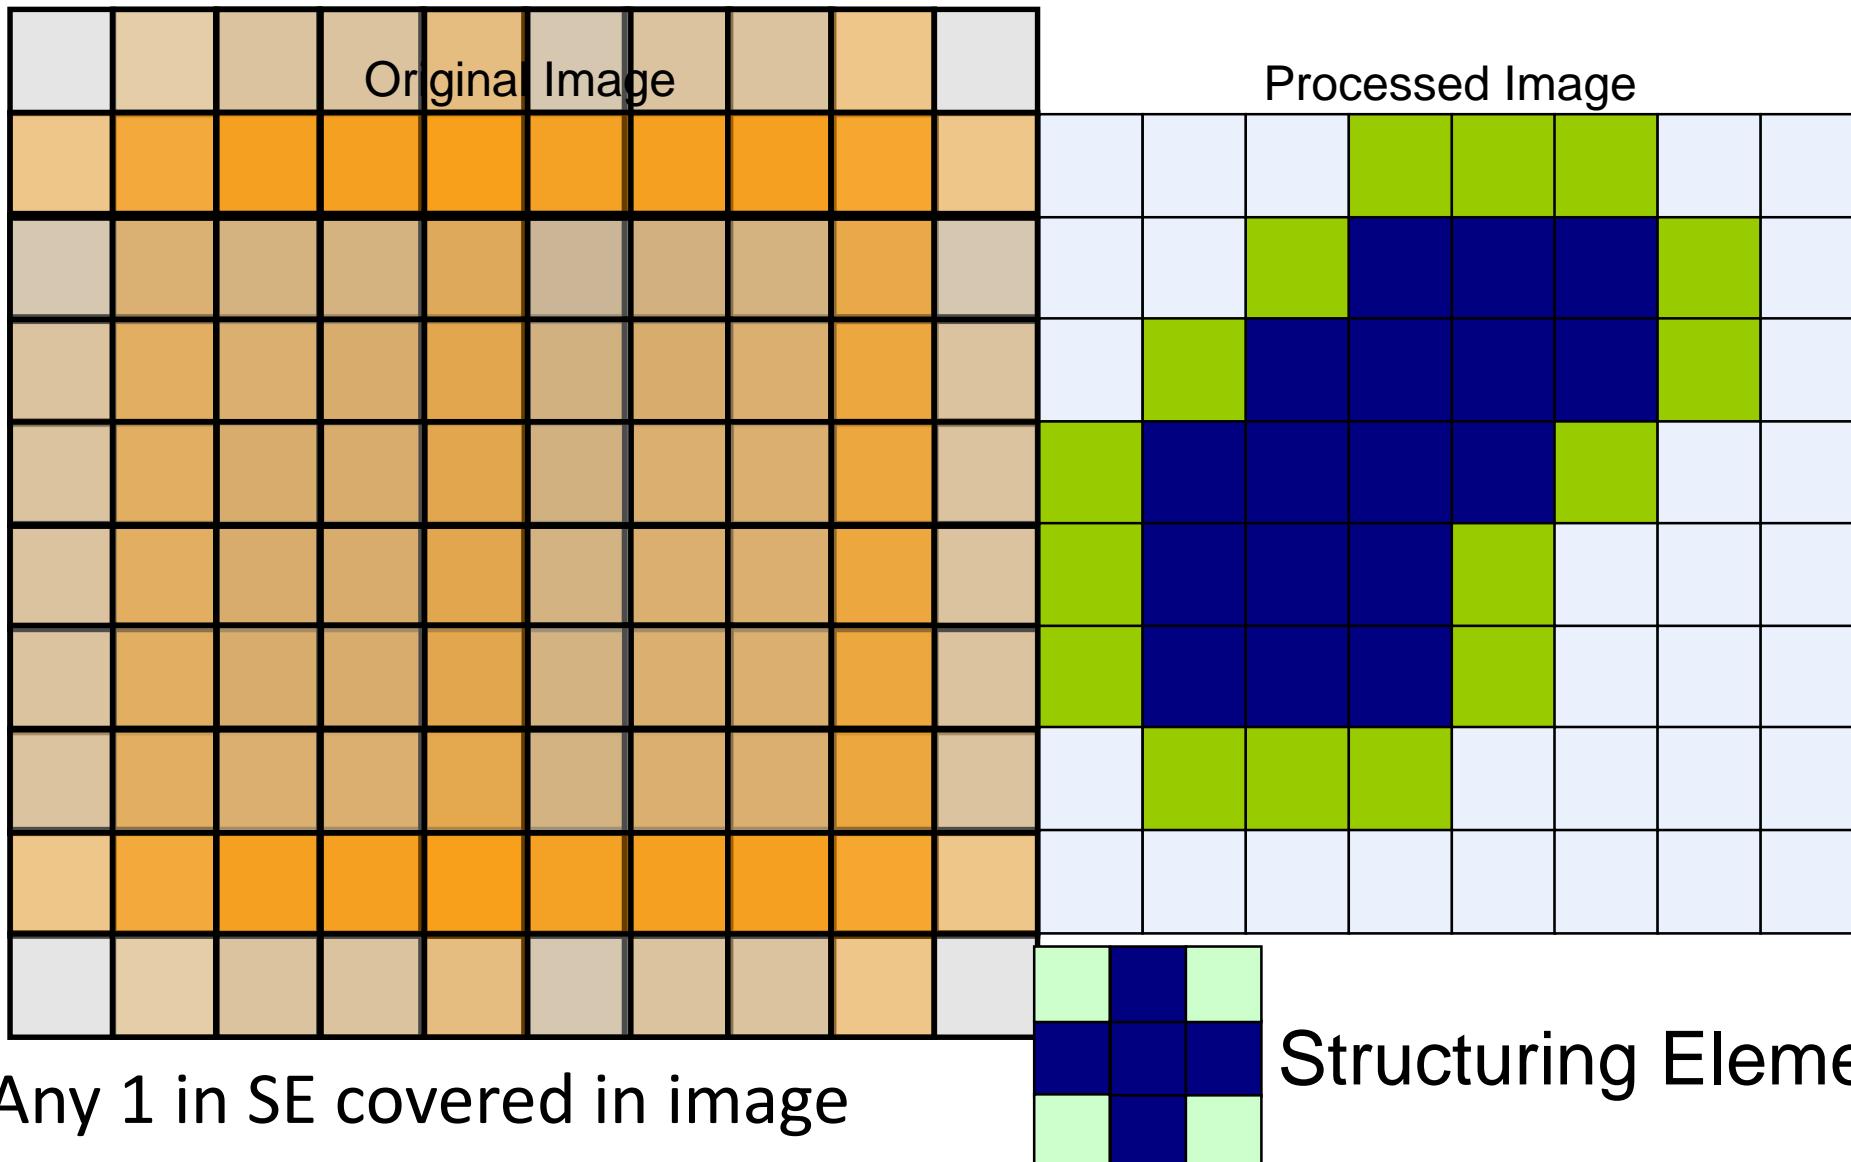
0	1	0
1	1	1
0	1	0

Hit: Any “1” pixels in SE cover “1” pixels in the image.

SE placed on image at (2,2)

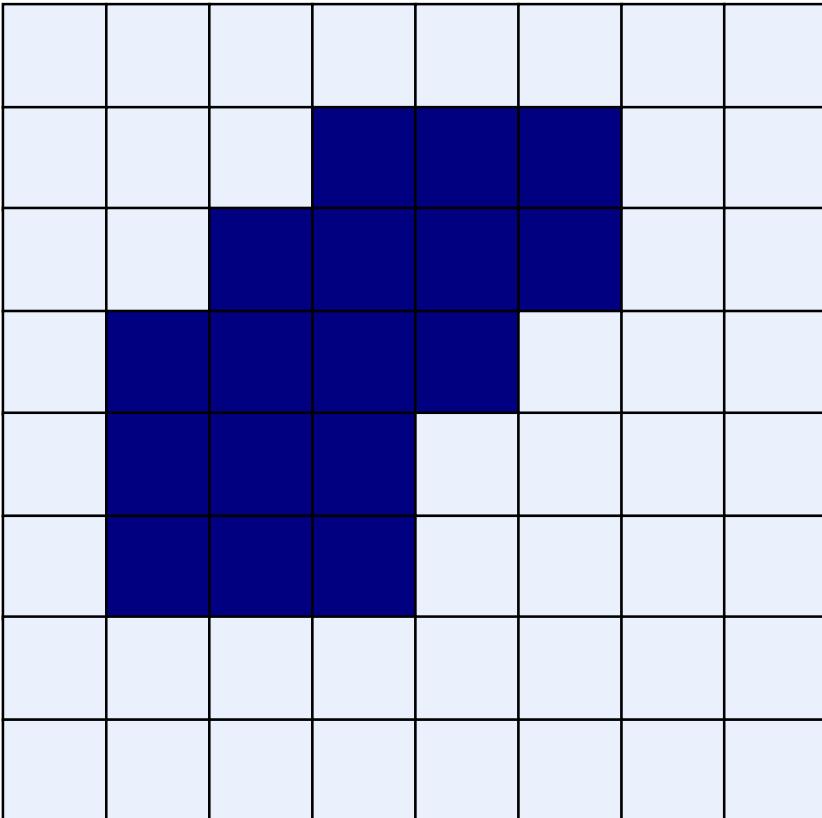
	0	1	2	3	4	5	6 ...
0	0	0	0	0	0	0	0
1	0	0	0	1	0	0	0
2	0	0	1	1	1	0	0
3	0	1	1	1	1	1	1
4	0	1	1	1	1	1	1
5	0	0	1	1	1	1	1
:

Dilation Example

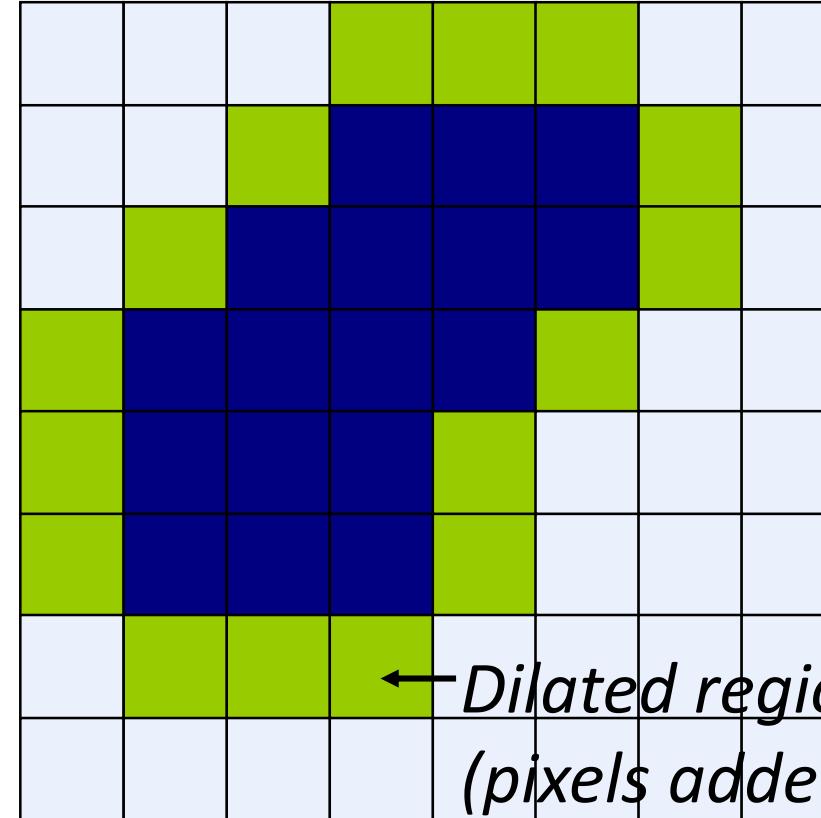


Dilation Example

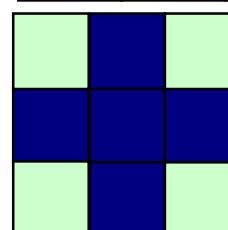
Original Image



Processed Image



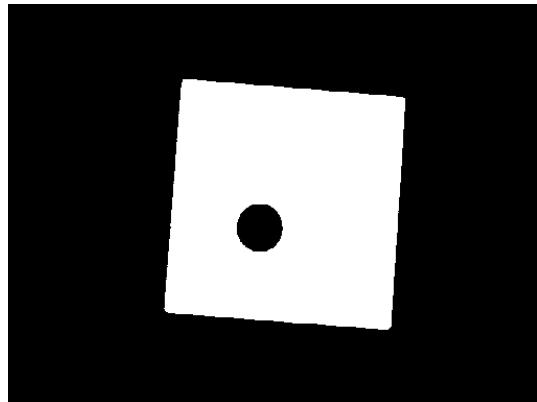
← *Dilated region
(pixels added)*



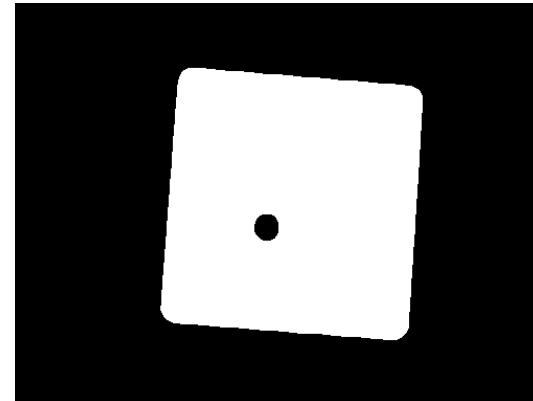
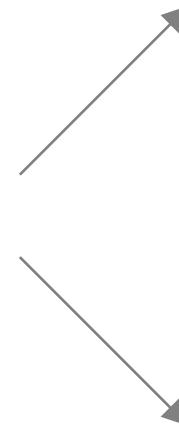
Hit: Any 1 in SE covered in image

Structuring Element

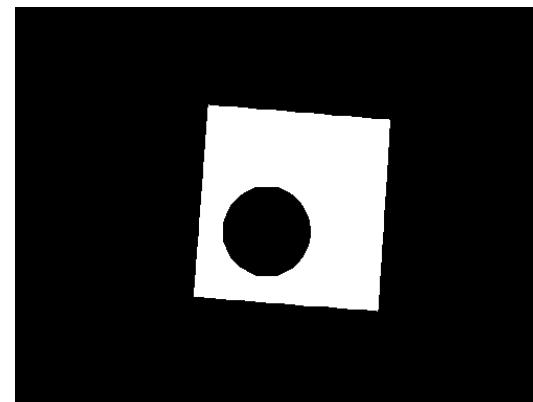
Effects of erosion and dilation



Original



Dilation by a round structuring element.

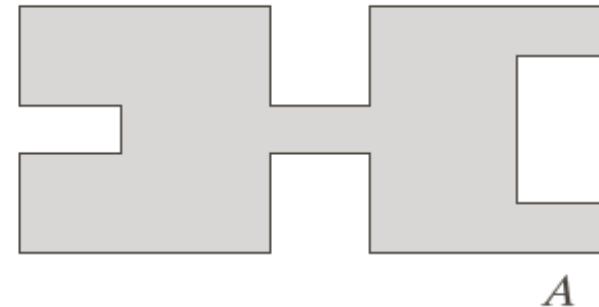


Erosion by a round structuring element.

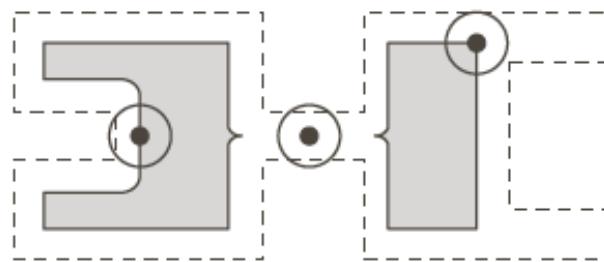
Combined operations: Opening

- Definition
 - Apply erosion then dilation

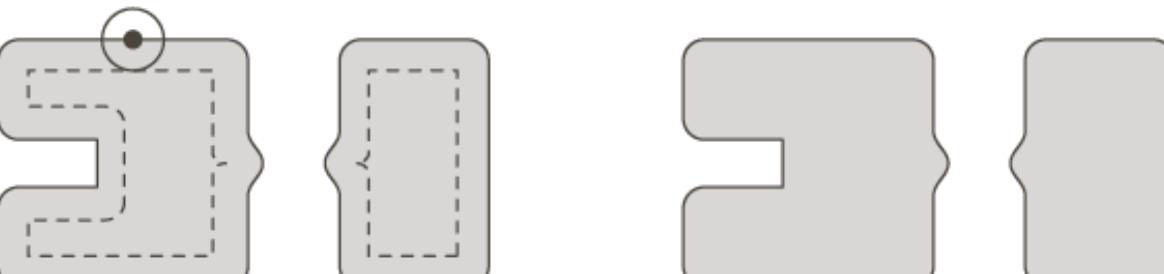
$$A \circ B = (A \ominus B) \oplus B$$



A



$A \ominus B$

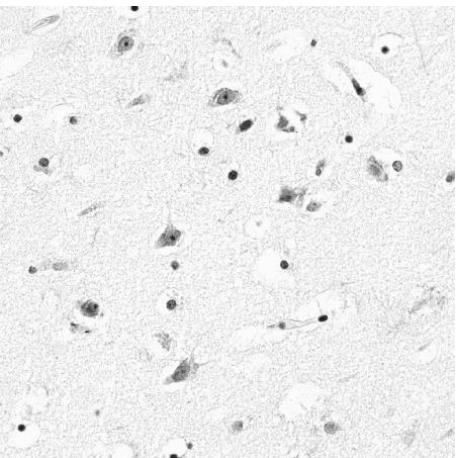


$A \circ B = (A \ominus B) \oplus B$

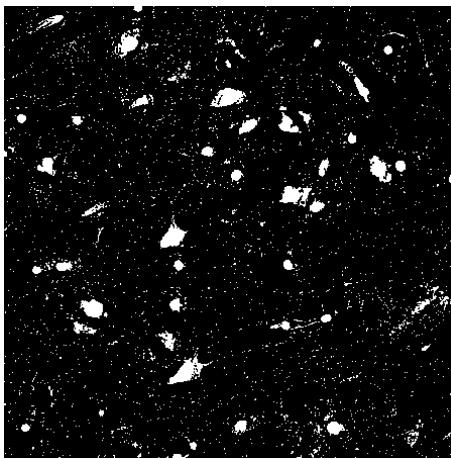
- Effect:
 - ⇒ Removes small objects,
preserves rough shape.

Effects of opening

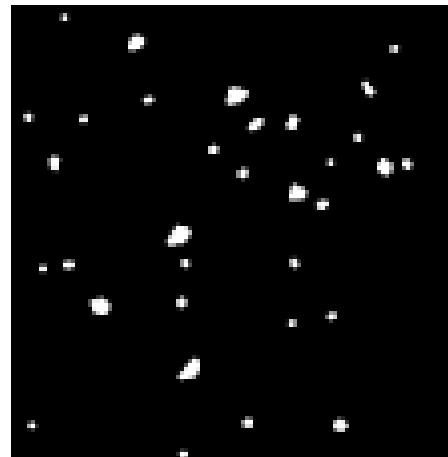
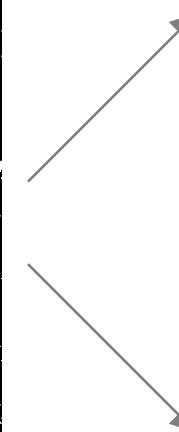
- Can filter out structures by selecting the size of structuring element.



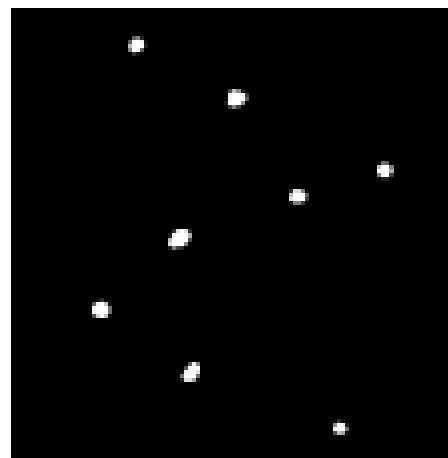
Original



Thresholded



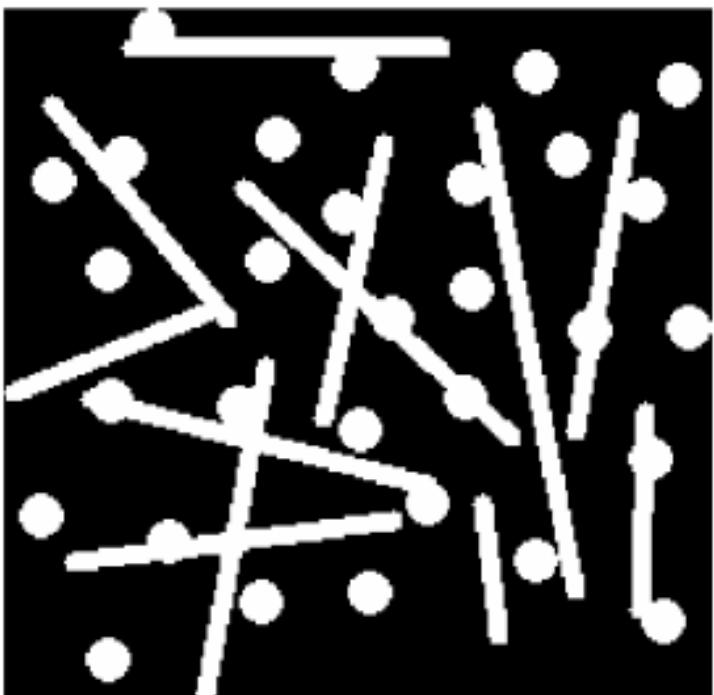
Opening by a small structuring element



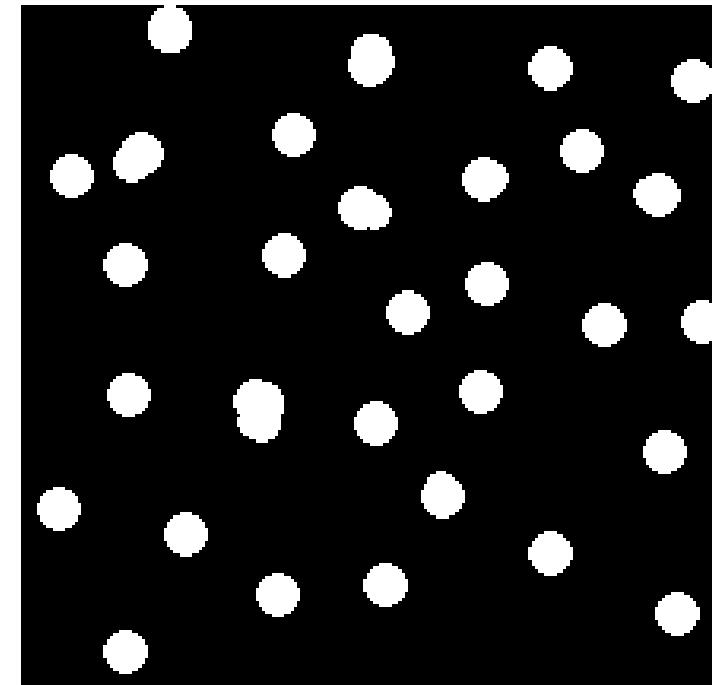
Opening by a large structuring element

Effects of opening

- Choose the structure in image by choosing the shape of the structuring element.



Original image

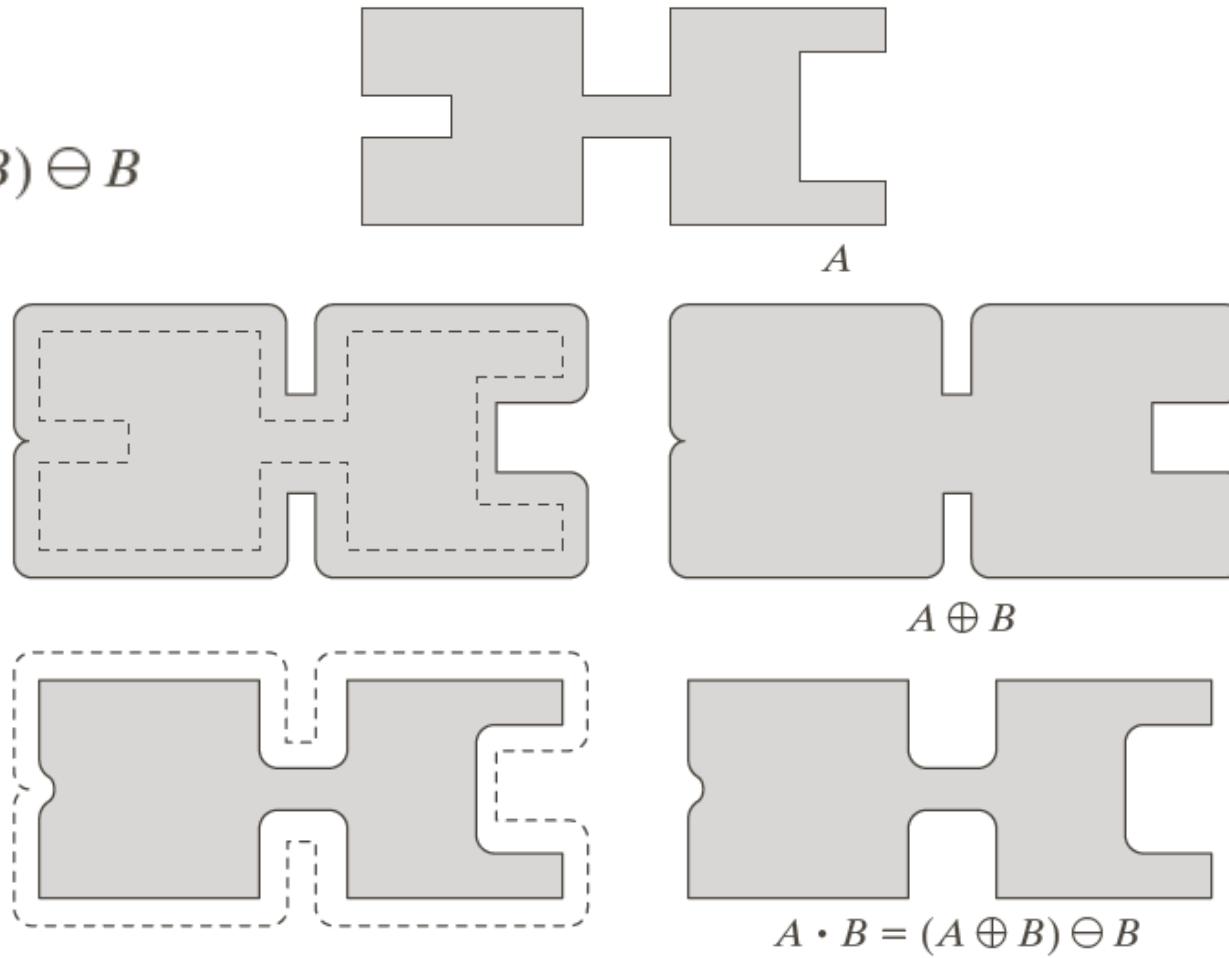


Opening by a round structuring element

Combined operations: Closing

- Definition
 - Apply dilation then erosion

$$A \cdot B = (A \oplus B) \ominus B$$



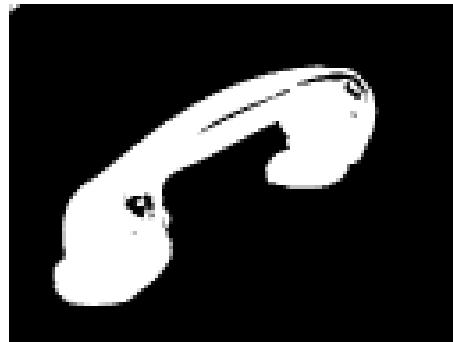
- Effect
 - ⇒ Fill holes, preserves the original shape.

Effects of closing

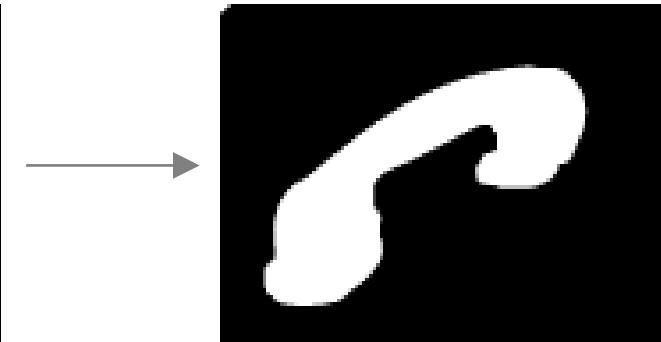
- Fill holes in thresholded image
(e.g., reflections)



Original

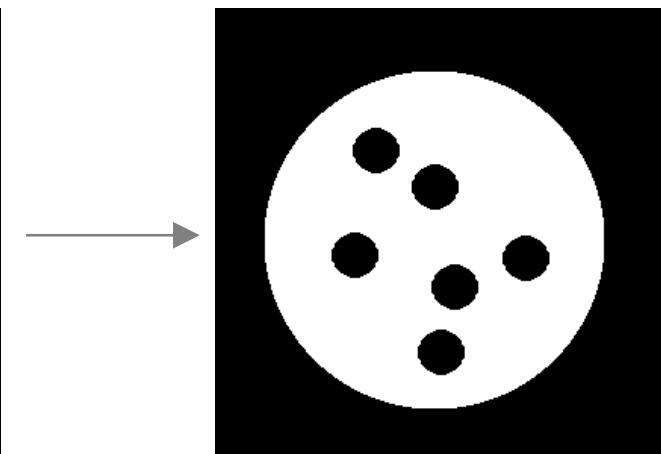
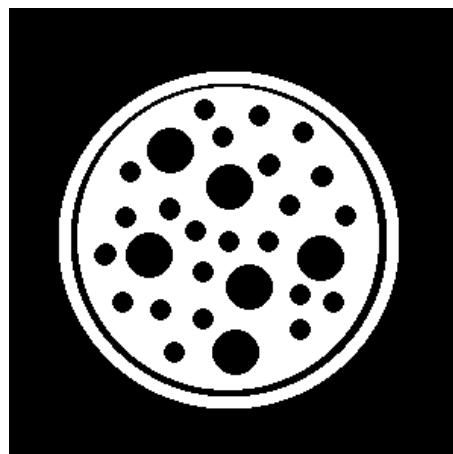


Thresholded



Closing by a round structuring element

The size of structuring element determines the maximal size of holes that will be filled.



Example: opening + closing

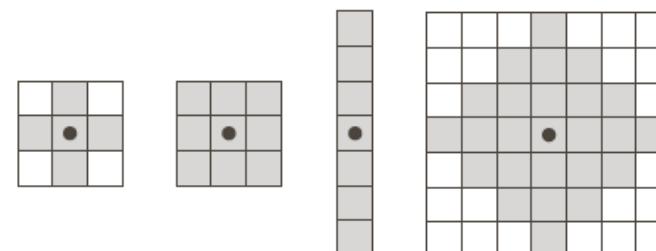


Morphological operators in Matlab

- Main operations
 - Dilation (Matlab: `imdilate`)
 - Erosion (Matlab: `imerode`)
- Several important combinations
 - Opening (Matlab: `imopen`)
 - Closing (Matlab: `imclose`)
 - Boundary extraction
- Much more available
(see help)



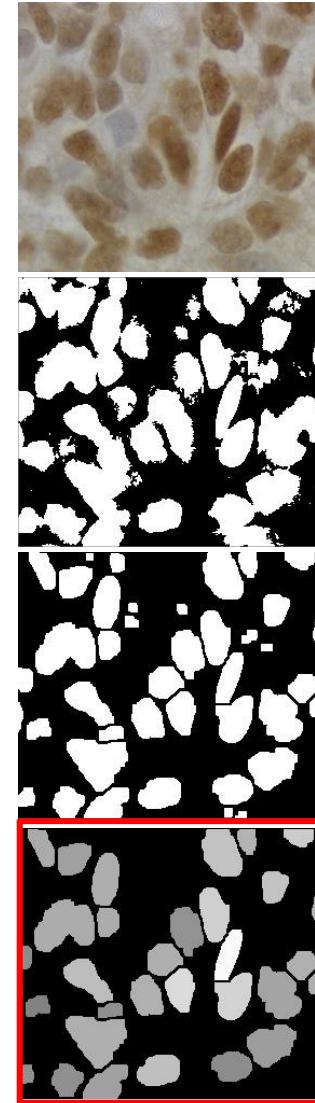
Examples of structuring elements:



Matlab:
`>> help strel`

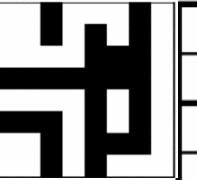
Machine perception

LABELLING REGIONS



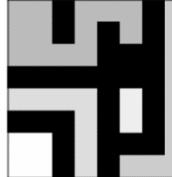
Connected components for labeling

- Goal: find separate connected regions



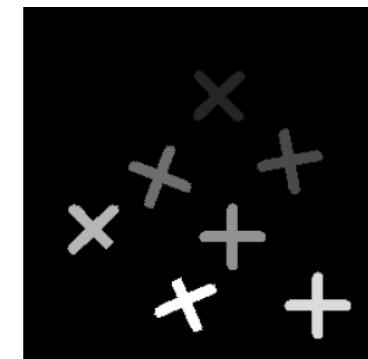
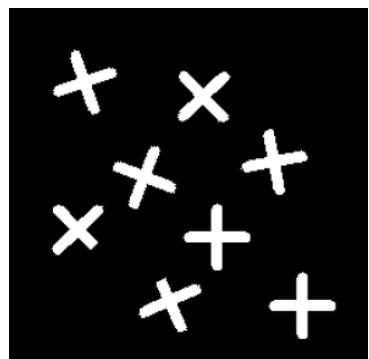
1	1	0	1	1	1	0	1
1	1	0	1	0	1	0	1
1	1	1	1	0	0	0	1
0	0	0	0	0	0	0	1
1	1	1	1	0	1	0	1
0	0	0	1	0	1	0	1
1	1	0	1	0	0	0	1
1	1	0	1	0	1	1	1

Binary image

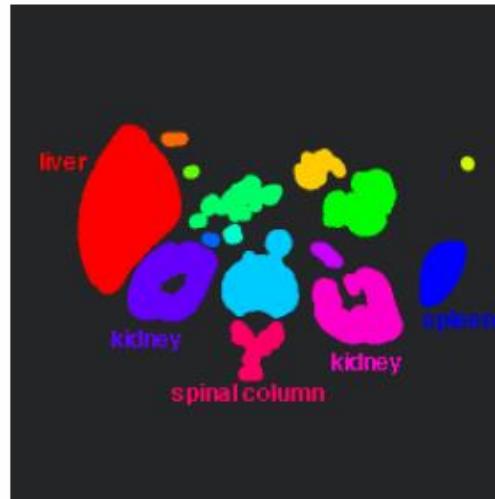


1	1	0	1	1	1	0	2
1	1	0	1	0	1	0	2
1	1	1	1	0	0	0	2
0	0	0	0	0	0	0	2
3	3	3	3	0	4	0	2
0	0	0	3	0	4	0	2
5	5	0	3	0	0	0	2
5	5	0	3	0	2	2	2

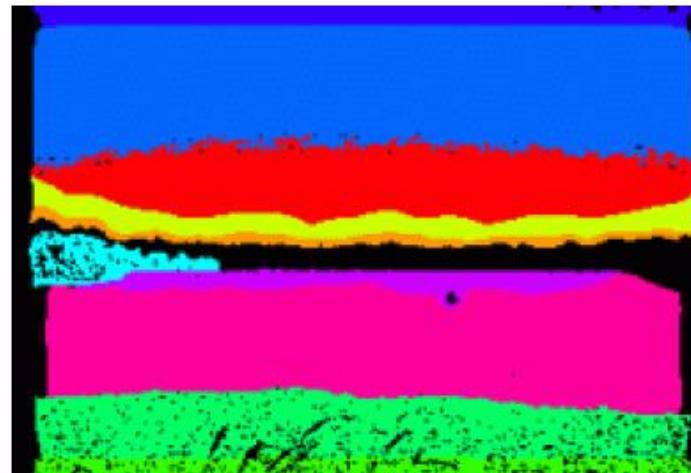
connected components



Examples of connected components



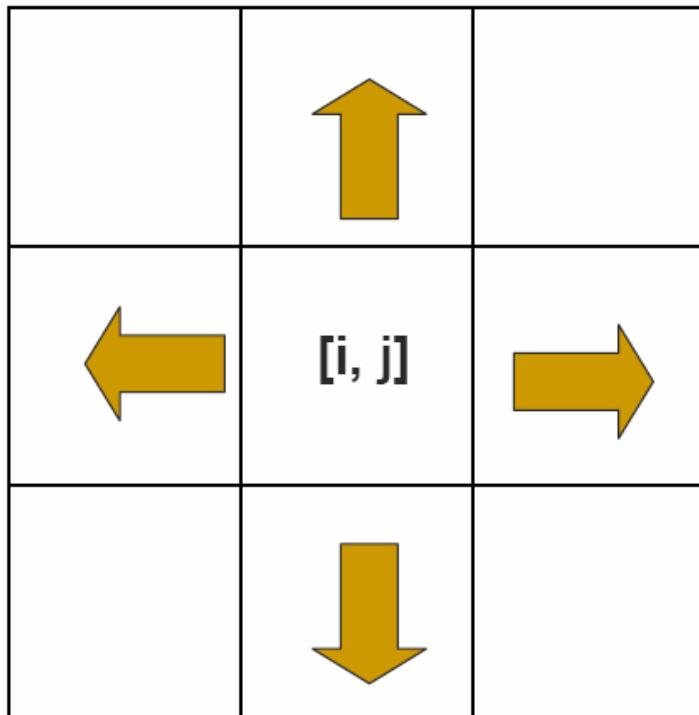
connected
components
of 1's from
thresholded
image



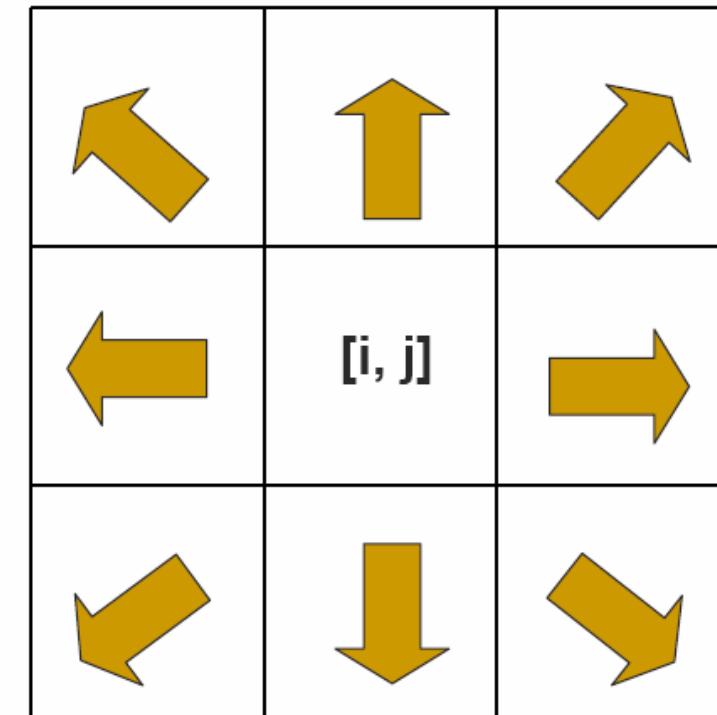
connected
components
of cluster
labels

Connectivity

- Determines which pixels are considered neighbors.



4-neighborhood



8-neighborhood

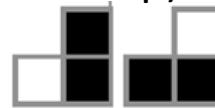
Sequential connected components

- Process image from left to right, from top to bottom:

1.) If the current pixel value is 1

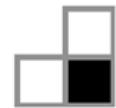
i.) If only one neighbor (left or top) is 1,

copy its label.



ii.) If both neighbors are 1 and have same label,

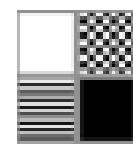
copy that label.



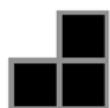
iii.) If they have different labels

– Copy label from the left.

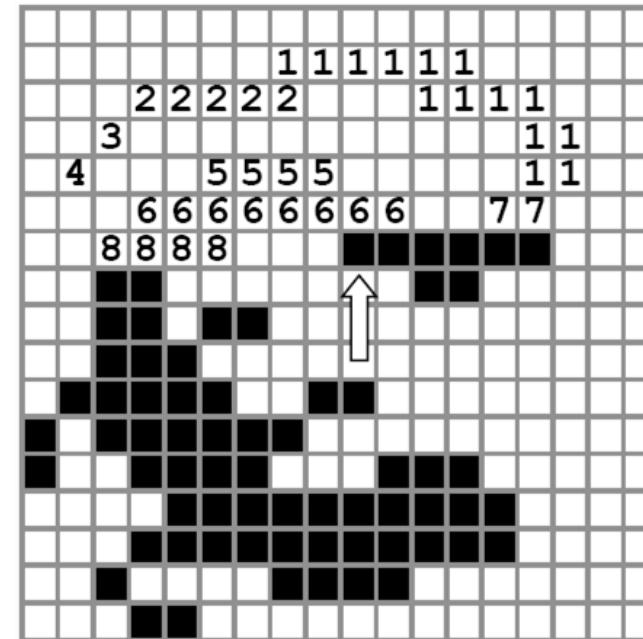
– Update the table of equivalent labels.



iv.) Otherwise form a new label.



- Relabel with the smallest equivalent labels.



{1, 2, 7}
{3, }
{4, }
{5, 6, 8}
{ }{ }

Example SCC: 8-connectivity

(a)

0 Background

1 Foreground

(b) only background neighbors

new label (2)

Example SCC: 8-connectivity

(c) exactly one neighbor label

neighbor label is propagated

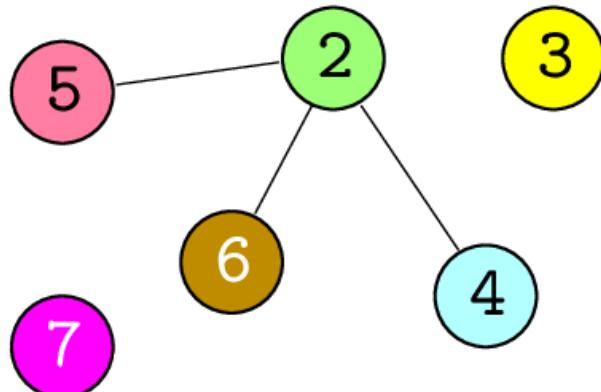
(d) two different neighbor labels

one of the labels (2) is propagated
(Update equivalency table {2,5})

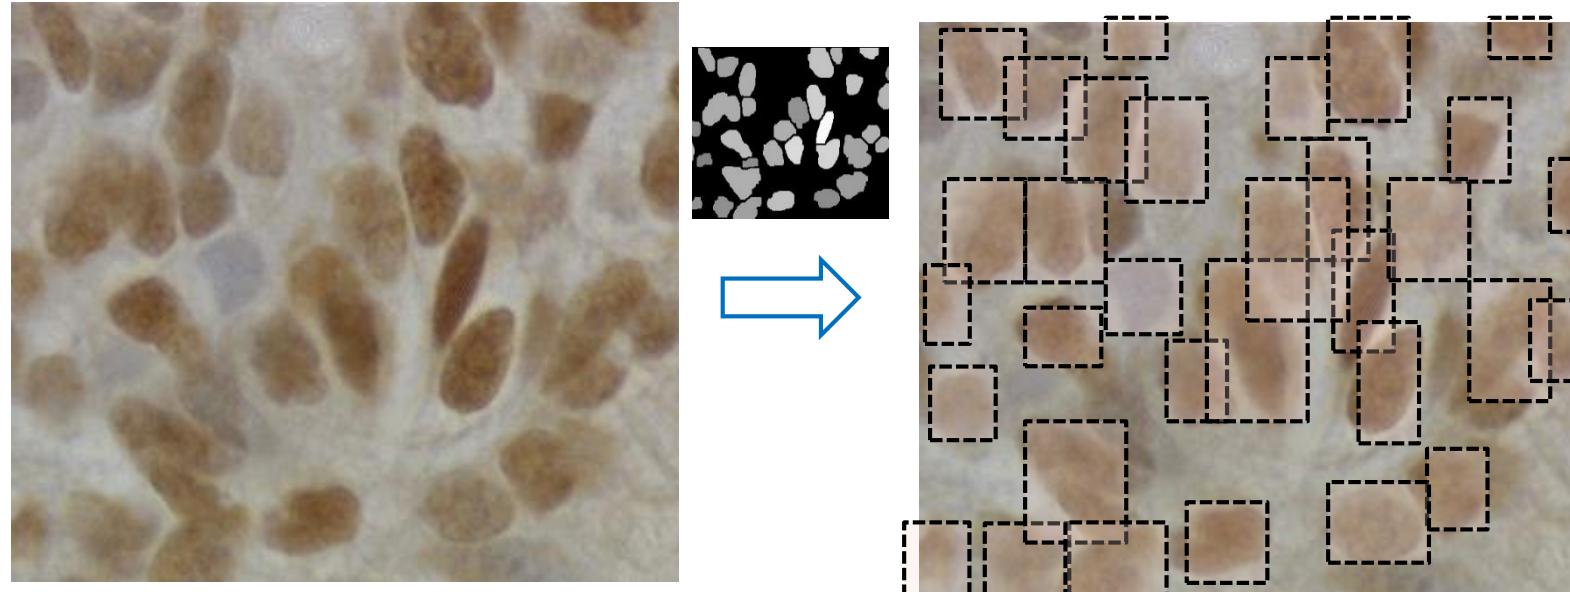
Example SCC: 8-connectivity

First pass: label

Equivalency table



Second pass: apply equivalences



Machine perception

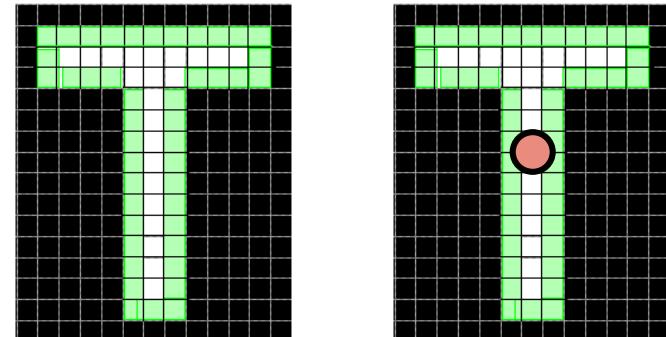
REGION DESCRIPTORS

Simple region descriptors

- A region can be detected using the connected components.
- How to describe it?
- Some examples:

- Area A
- Perimeter l
- Compactness $c=l^2/(4\pi A)$
- Circularity, roundness l/c
- Centroid (center of mass)
- Major and minor axes λ_1, λ_2
- Eccentricity $\|\lambda_1\|/\|\lambda_2\|$
- Minimal bounding box area $A_m = h b$
- Rectangularity A/A_m

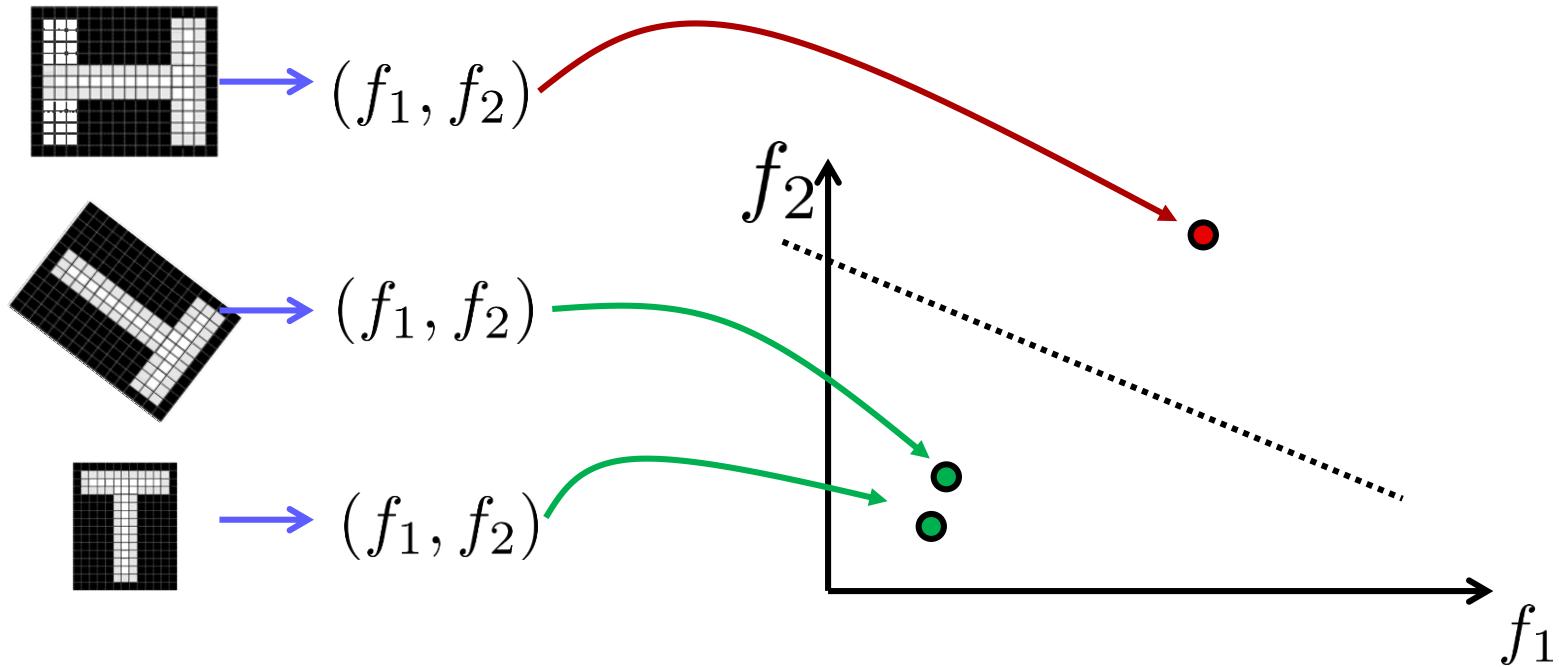
Matlab: *regionprops*



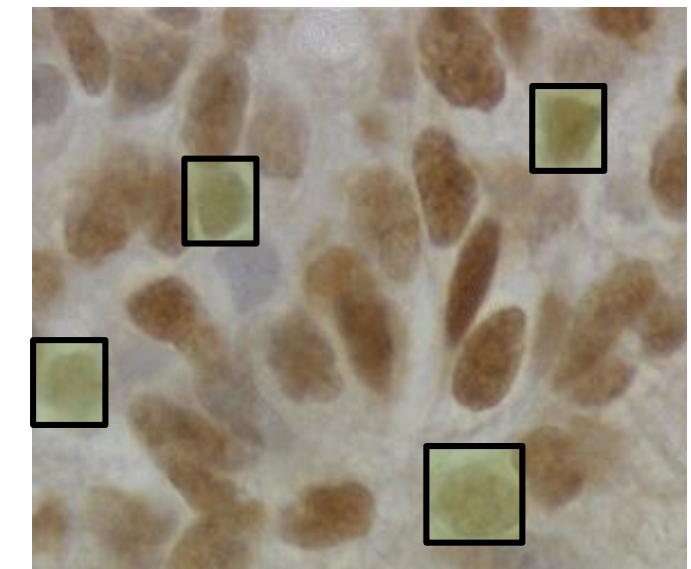
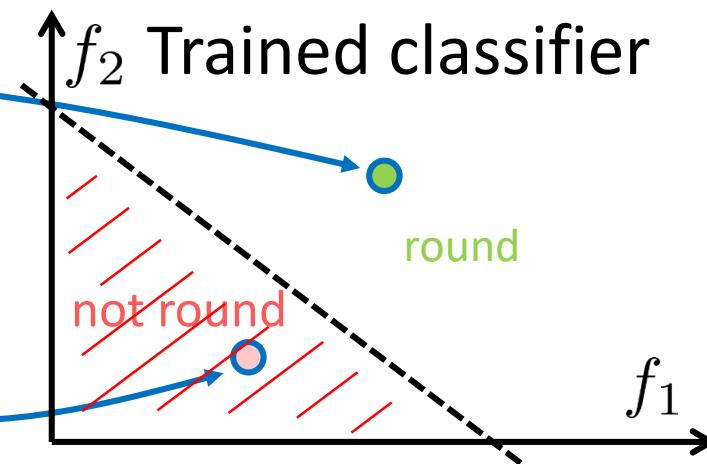
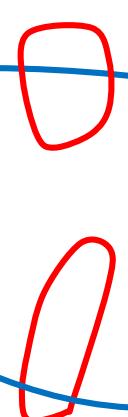
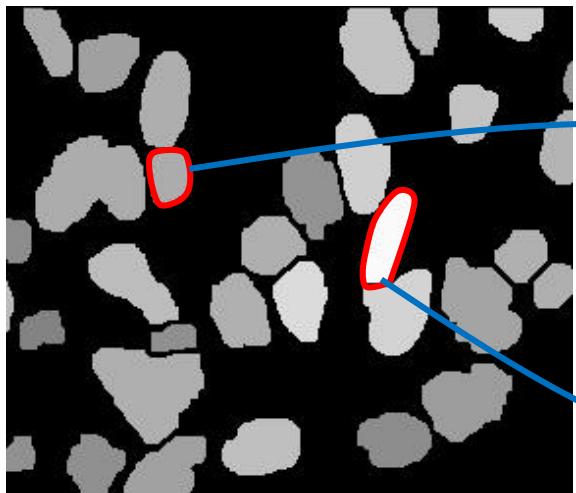
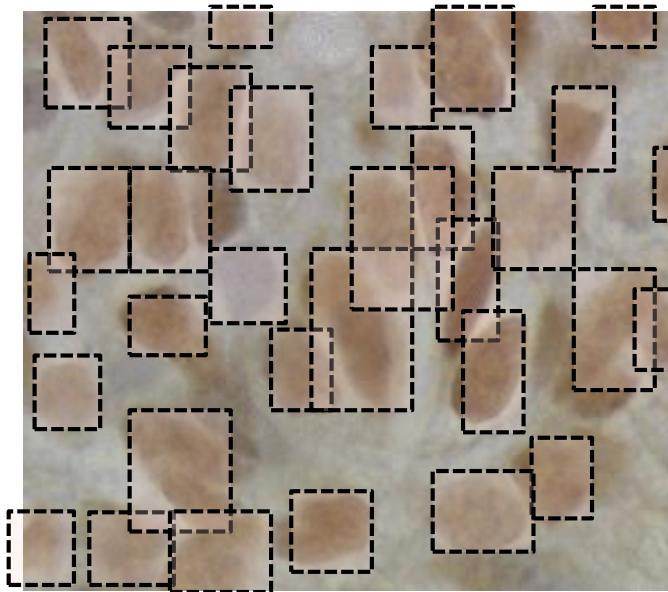
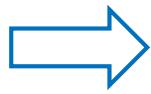
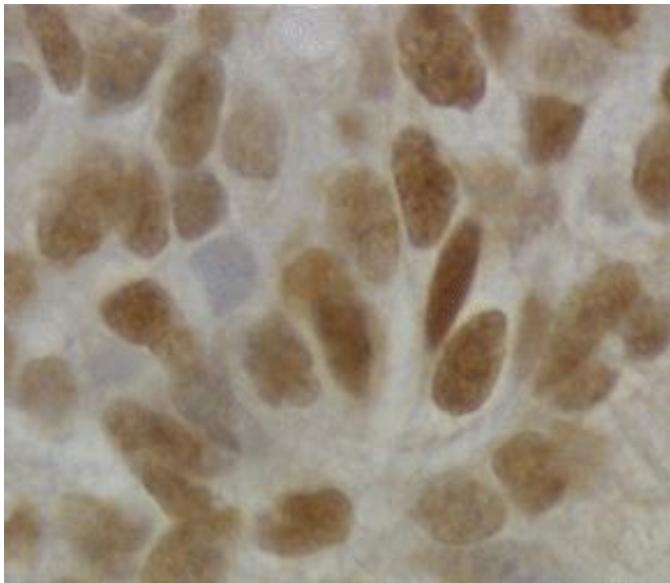
(Easy to come up with your own)

Require a level of invariance (App dependent)

- We seek a descriptor that maps:
 - Two images of the same object close-by in feature space.
 - Two images of different objects to points far between each other.



Task: Detect round cells

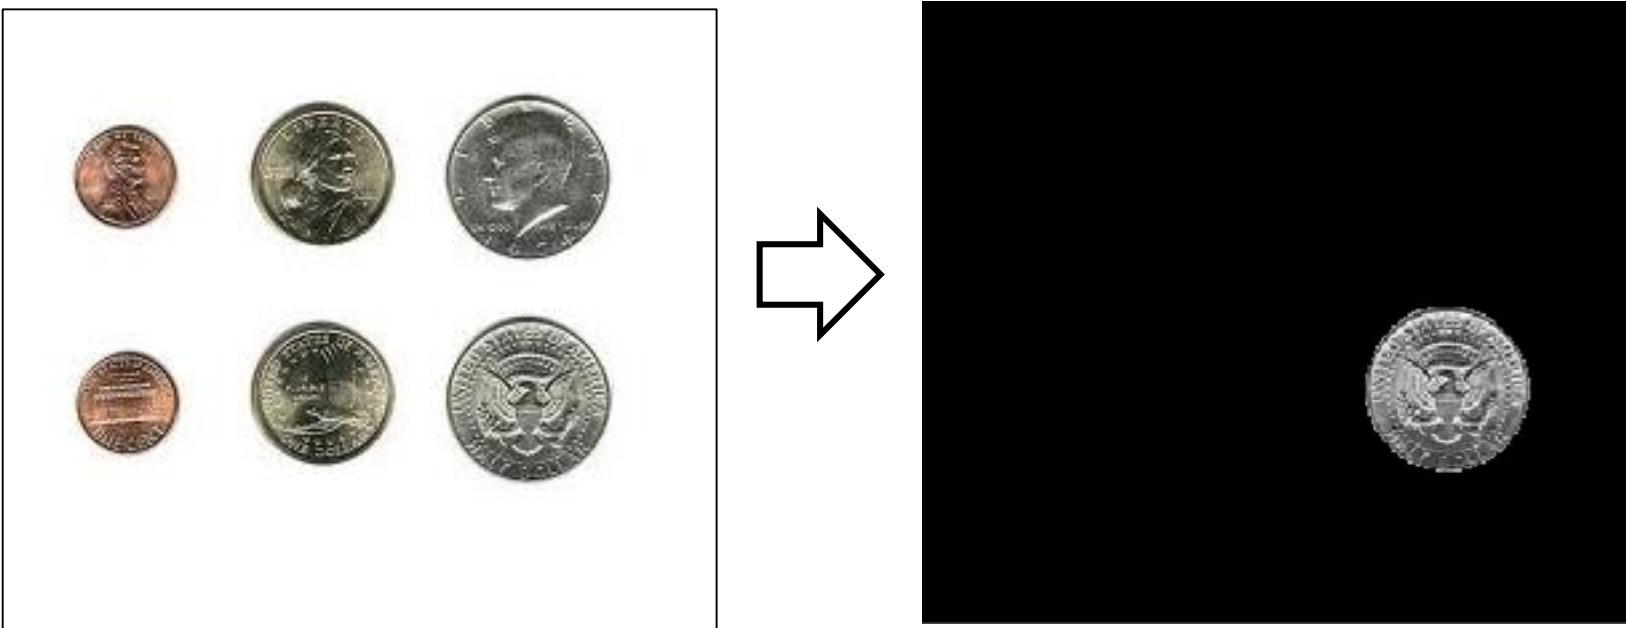


Summary: Binarization

- Pros
 - Fast, simple to store
 - Simple techniques
 - Works in constrained setups
- Cons
 - Difficult to get „clean“ shapes
 - Many real-world scenarios contain noise
 - Often too coarse representation
 - Not robust in changes of 3D view changes

Try and play with this at home:

- Task: “Automatically select only one of the coins”
 - Threshold, clean by morphology, find connected components, select a single connected component, overlay the input image with the selection.



Matlab code...

- `a = imread('coins.jpg') ;`
- `figure(1) ; clf ; imagesc(a);`
- `a = rgb2gray(a) ;`
- `figure(1) ; clf ; imagesc(a); colormap gray ;`
- `I = a < 200 ; figure(1) ; clf ; imagesc(I); colormap gray ;`
- `se = strel('disk',4) ;`
- `I2 = imclose(I,se) ; figure(1) ; clf ; imagesc(I2); colormap gray ;`
- `[I3,num] = bwlabel(I2);`
- `figure(2); clf ; imagesc(I3)`
- `figure(3); clf ; imagesc(I3==1)`
- `a2= double(a).*double(I3==1) ;`
- `figure(4); clf ; imagesc(a2)`

References

- R.C. Gonzales, R.E. Woods, *Digital Image Processing*. Prentice Hall, 2001
- Morfology – Spletni tutorial: <http://homepages.inf.ed.ac.uk/rbf/HIPR2/>
- David A. Forsyth, Jean Ponce, Computer Vision: A Modern Approach (2nd Edition), (prva izdaja dostopna na spletu)
- R. Szeliski, Computer Vision: Algorithms and Applications, 2010
- R. Hartley, A. Zisserman, Multiple View Geometry in Computer Vision, 2nd Edition, Cambridge University Press, 2004
- Rob Fergus, „Computer Vision“, lectures
- Bastian Leibe „Computer Vision“, lectures
- Kristen Grauman „Computer Vision“, Lectures
- Otsu, N., "A Threshold Selection Method from Gray-Level Histograms," *IEEE Transactions on Systems, Man, and Cybernetics*, Vol. 9, No. 1, 1979, pp. 62-66.