



# Machine perception

# Image processing 2

Matej Kristan



Laboratorij za Umetne Vizualne Spoznavne Sisteme,  
Fakulteta za računalništvo in informatiko,  
Univerza v Ljubljani

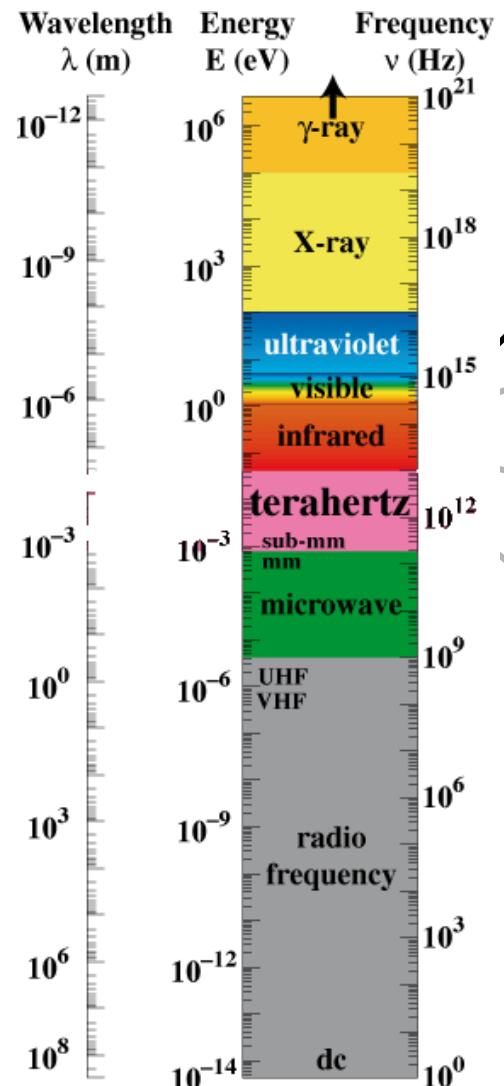


Machine perception

**COLOR**

# Sensor: Camera

Electromagnetic spectrum



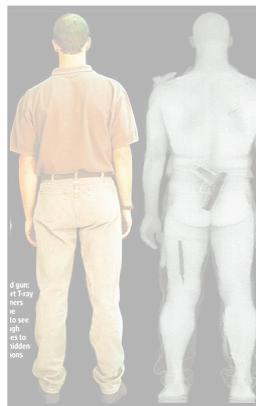
Visible light



Near-infrared light



Far-infrared light

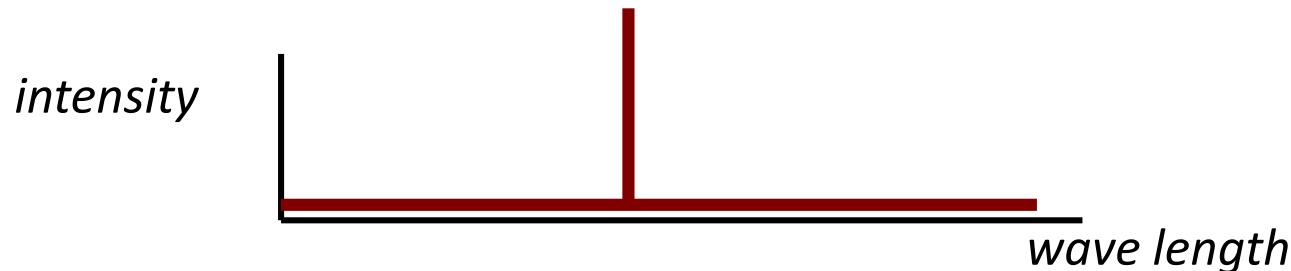


Terahertz light

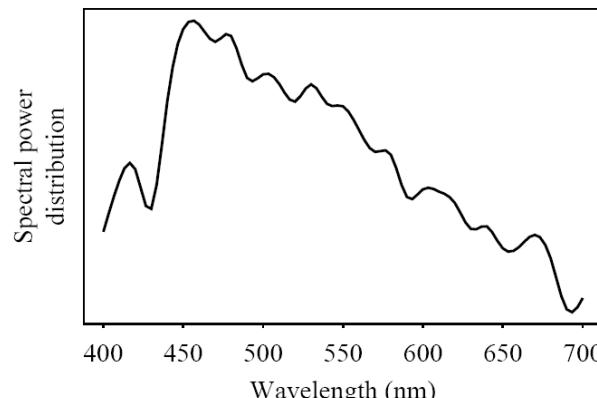
# Light

---

- *Light is an electromagnetic radiation composed of several frequencies*
- Properties are described by its spectrum (i.e., how much of each frequency is present)
- E.g., laser light contains only a narrow band of wavelengths (frequencies)



- Visible light contains radiation with wavelength of 400-700nm

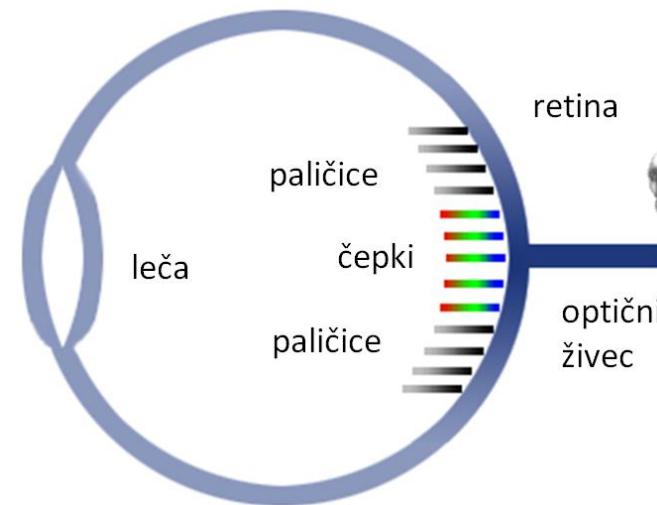
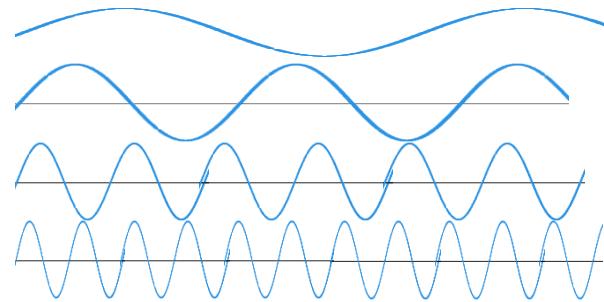
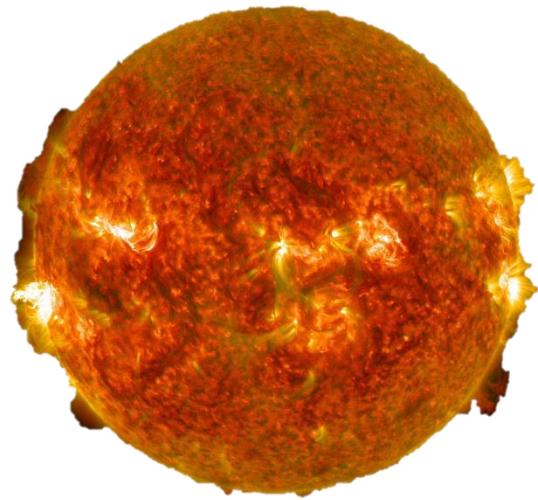


**Spectrum:** Energy radiated in a unit of time w.r.t. wavelength.

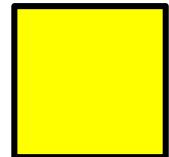
# Human color perception

---

- Human eye (retina) contains specialized cells that react to different wavelengths differently.
- Three types of cells called “cones”: R, G, B
- A type of cells called “rods”: intensity only

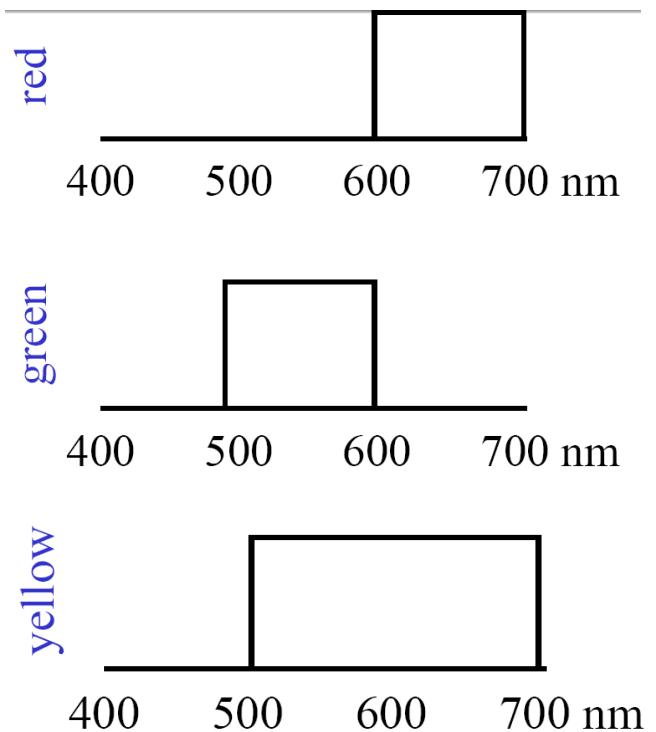
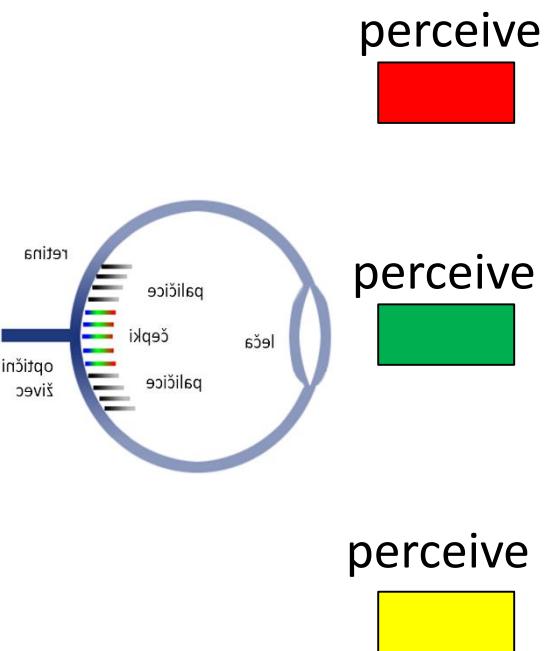


Yellow

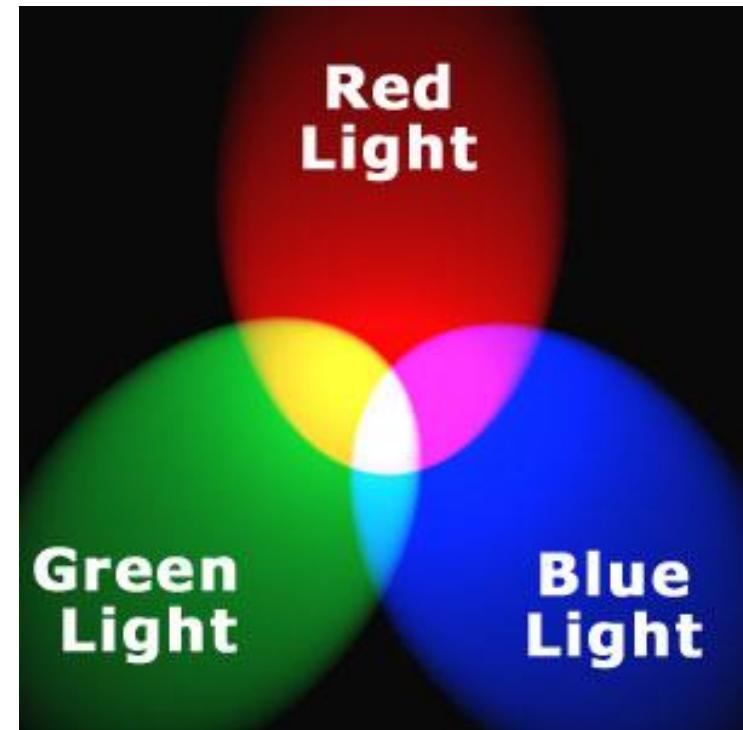


# Additive mixture model

- What color do we get if we shine a **red** and **green** light?



colors mix by summation of their spectra.



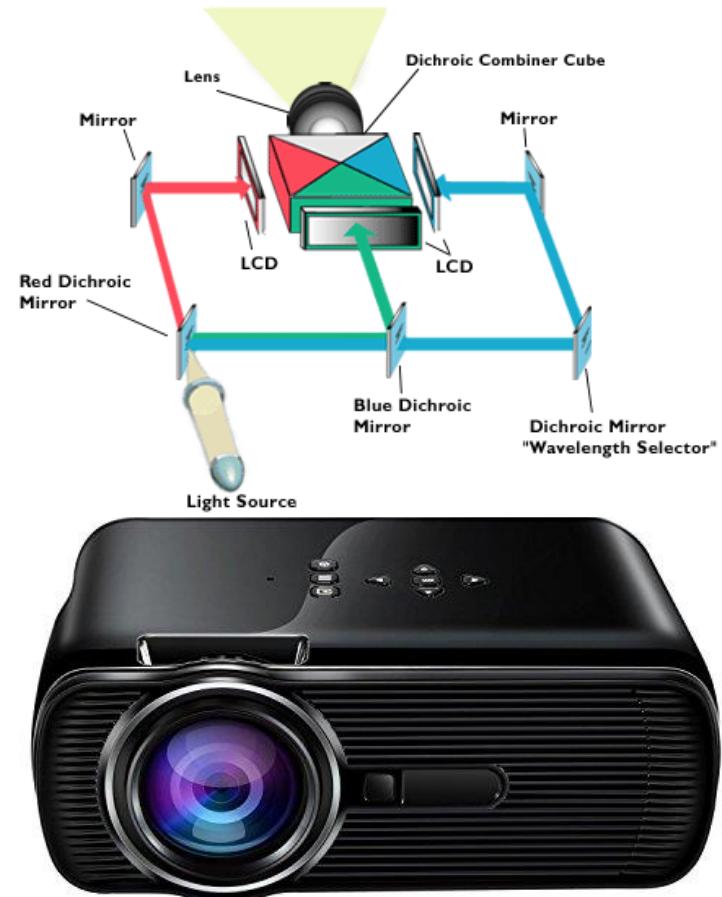
colors *added to black*.

# Systems using the additive model

---



Monitors

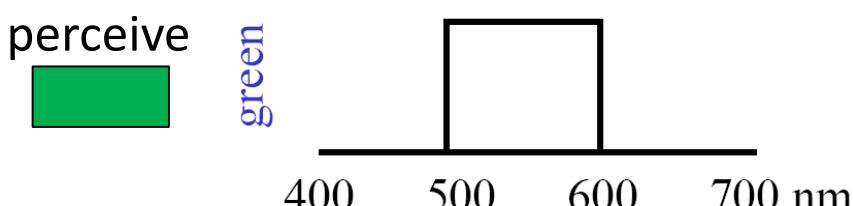
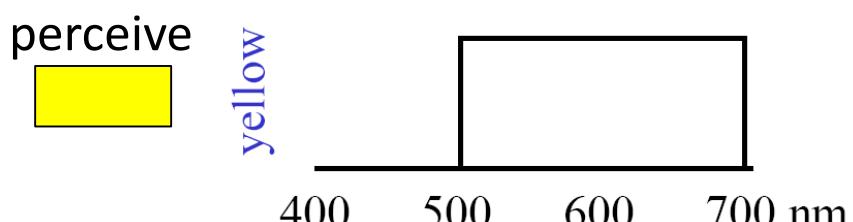
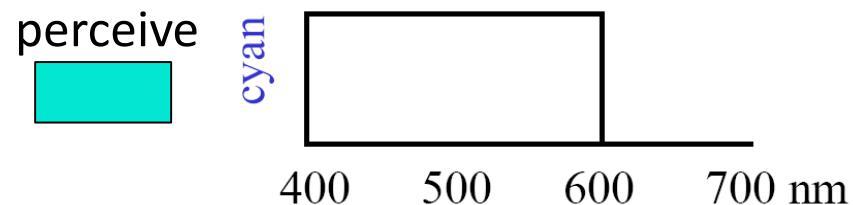
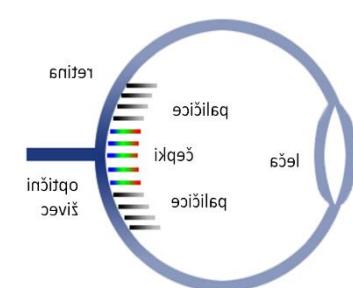
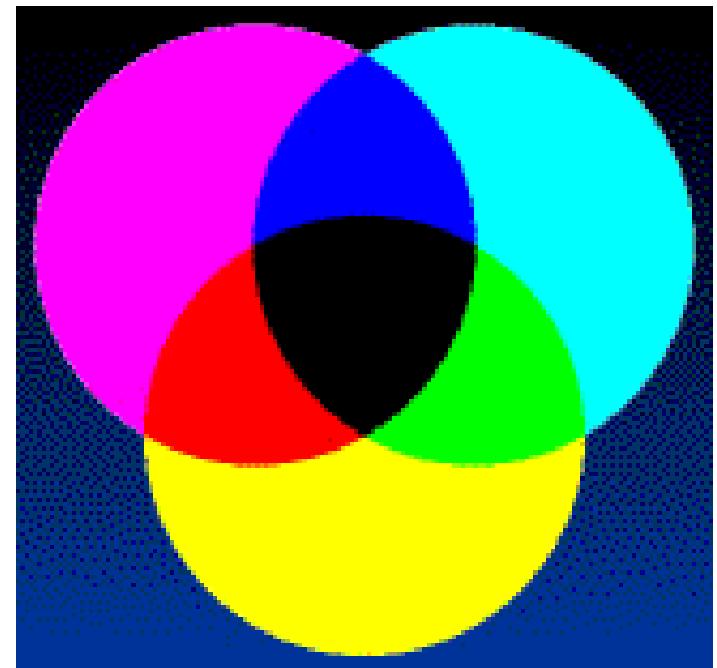


LCD projector

# Subtractive models

- What color do we get if applying cyan and yellow pigment to white paper?

colors mix by spectra  
*multiplication.*



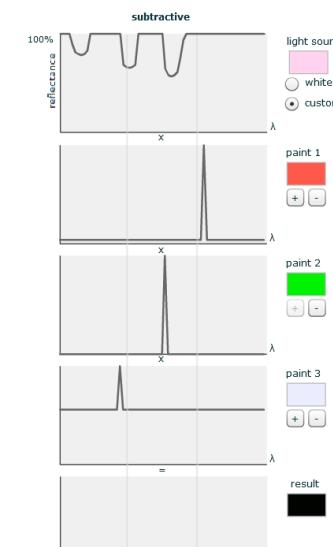
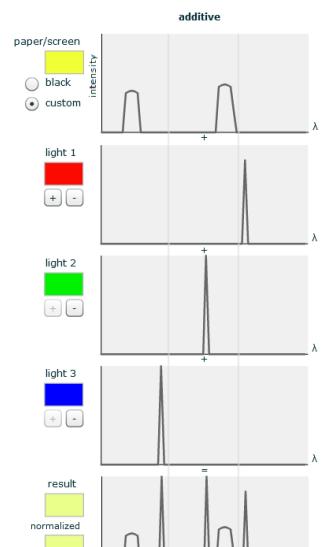
Pigments *remove* the color from the incident white light.

# Systems using a subtractive model

- Printing on paper
- Crayons
- Photographic film



- See this nice app and play with setups:  
<https://graphics.stanford.edu/courses/cs178/applets/colormixing.html>

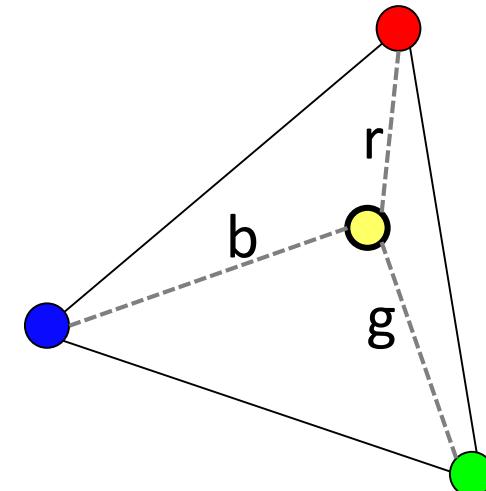


# Color spaces

---

- Role of colorspace: Unique color specification (e.g., for reproduction)
  - *Specifying a color in a color space allows accurate color reproduction on various media like photo, print and monitor.*
- Defined by the choice of **primary colors (primaries)**  
Recall: The human eye is equipped with sensory cells for the perception of the three primary colors (RGB)
- New color is a **weighted sum** of primaries

*By mixing the colours, we get any colour that lies within the triangle of primaries.*
- Mixing weights  $r,g,b$  to get any color were estimated on human subjects

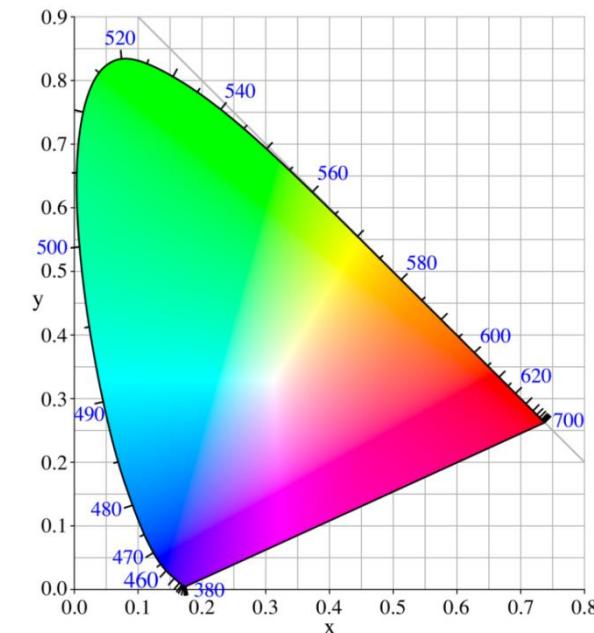
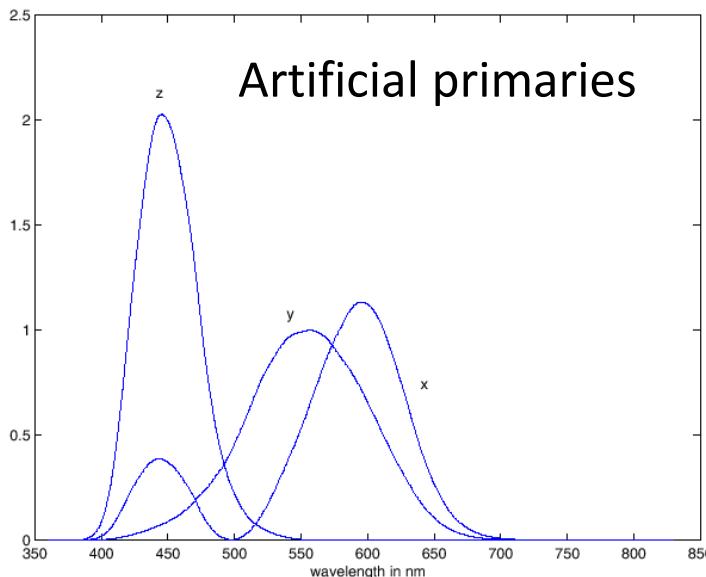


# Linear color space example: CIE XYZ

---

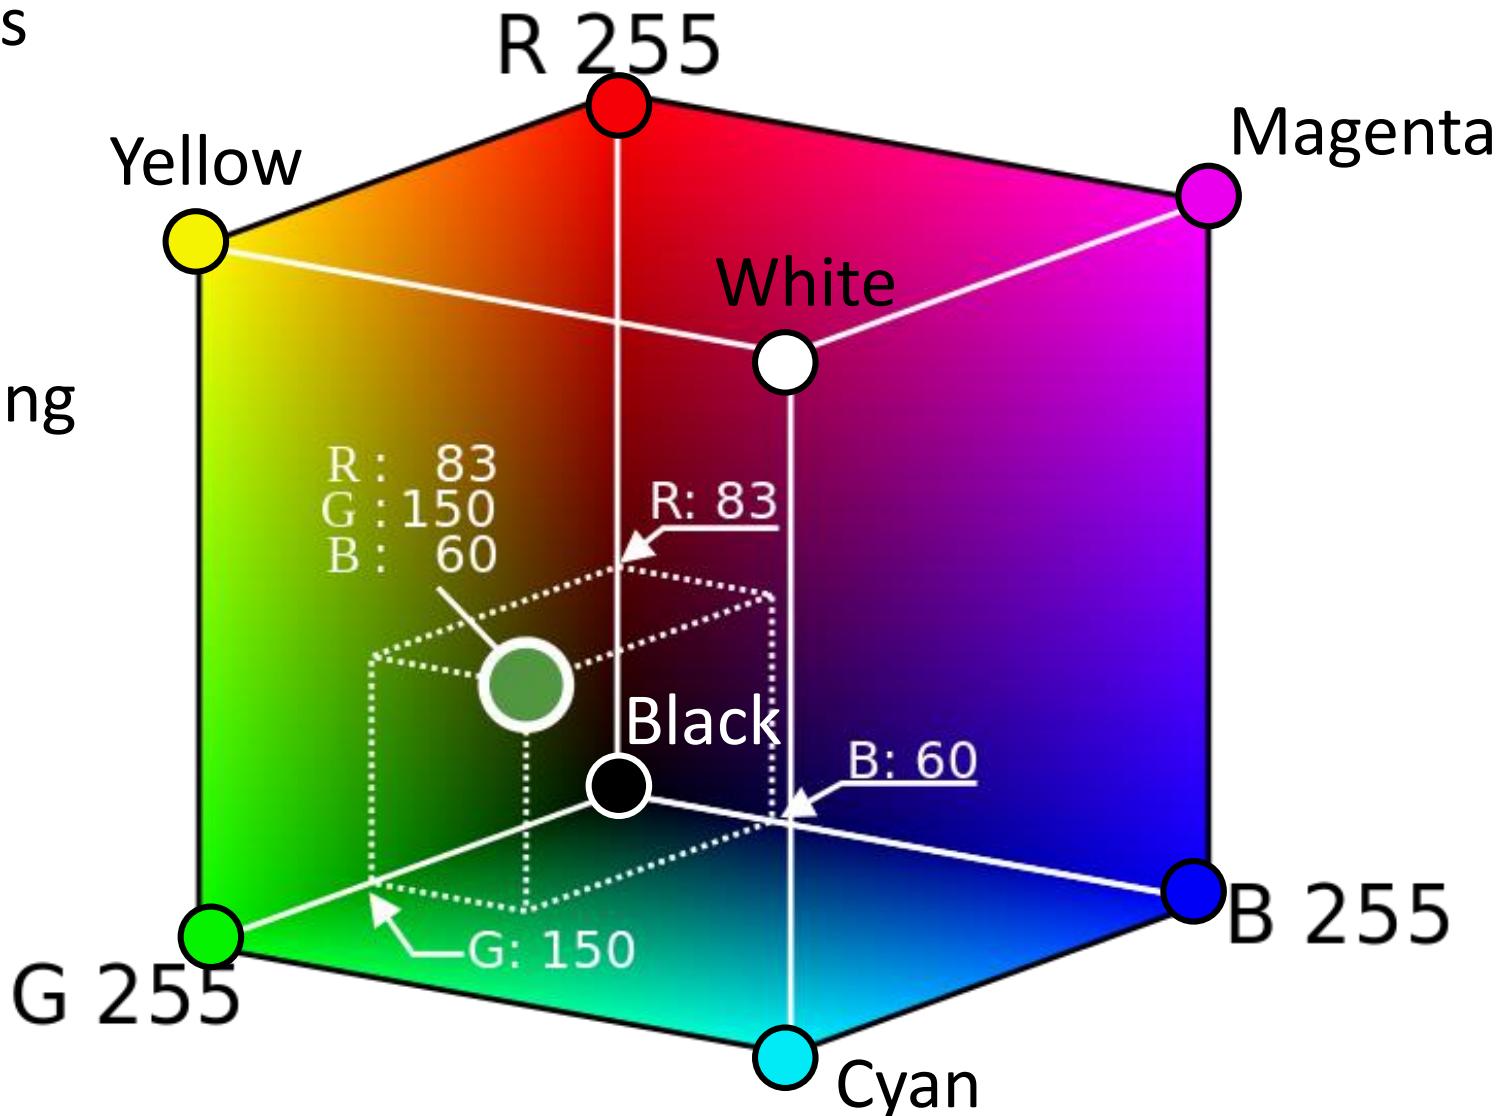
- Commission international d'éclairage (CIE), 1931
- Y is intensity
- Representation by xy colorspace:

$$x = \frac{X}{X+Y+Z}, \quad y = \frac{Y}{X+Y+Z}, \quad z = \frac{Z}{X+Y+Z}$$
$$x + y + z = 1$$



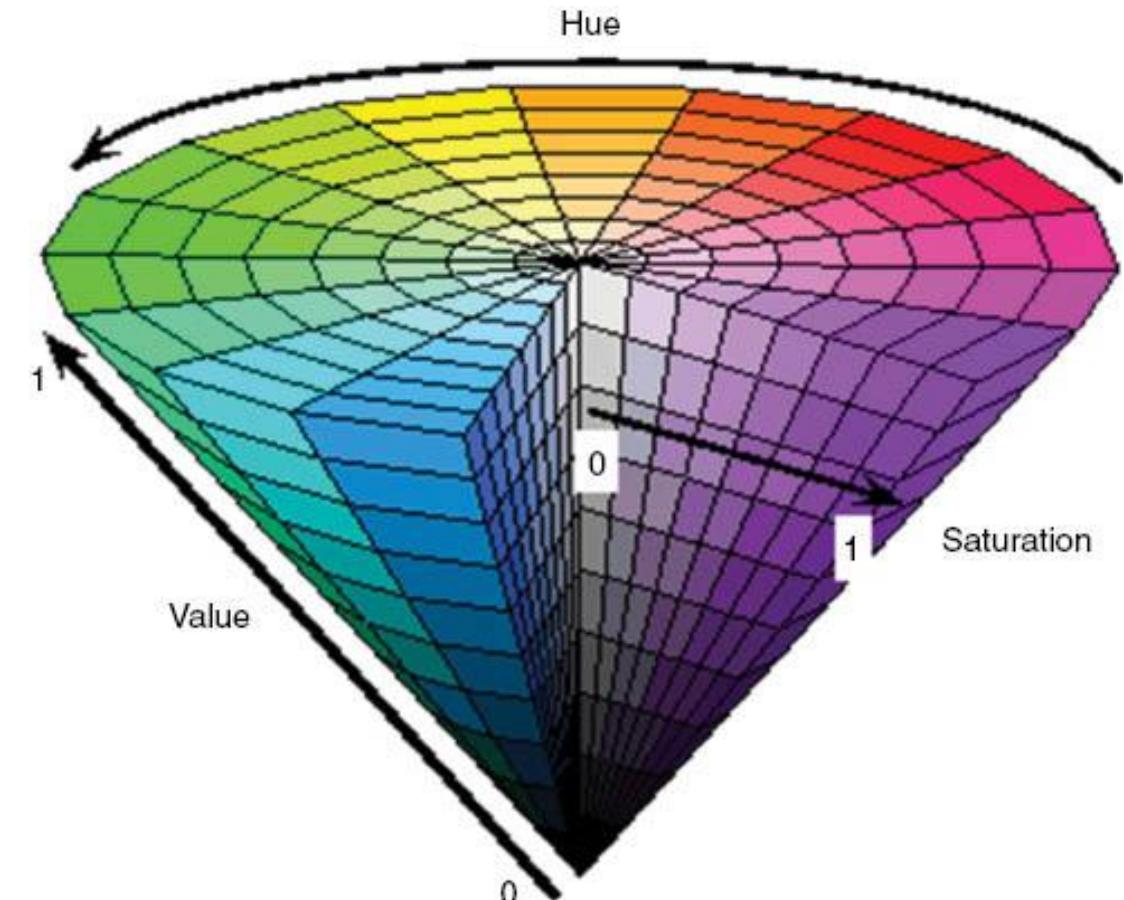
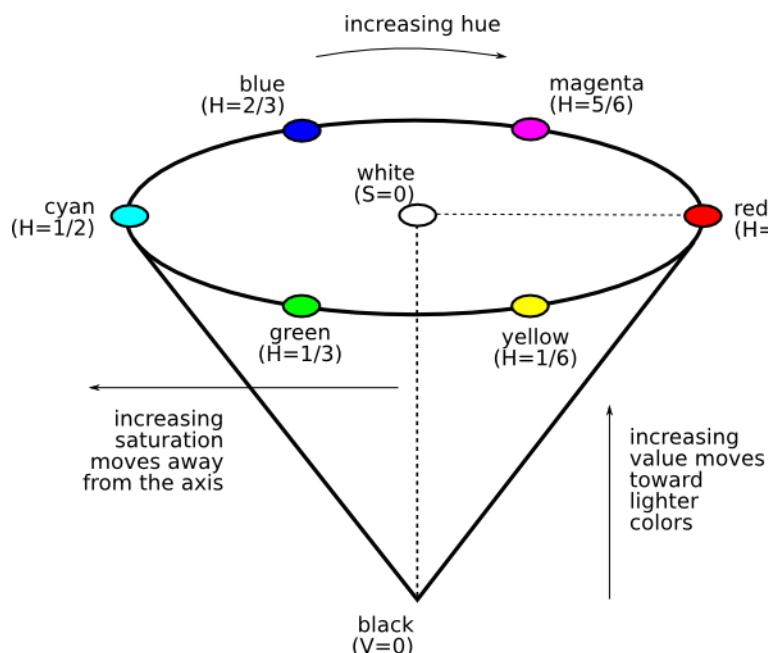
# Linear color space example: RGB

- Single wave-length primaries
- Appropriate for use in imaging devices (e.g., monitors), but not for human perception



# HSV colorspace

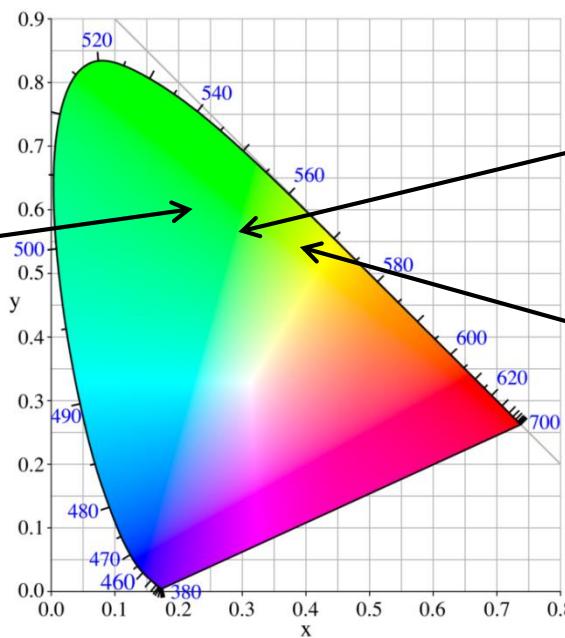
- Hue (barvnost), Saturation (nasičenje), Value (intenziteta)
- Nonlinear – hue coded by angle
- Matlab: `hsv2rgb`, `rgb2hsv`.



# Distances in colourspaces

---

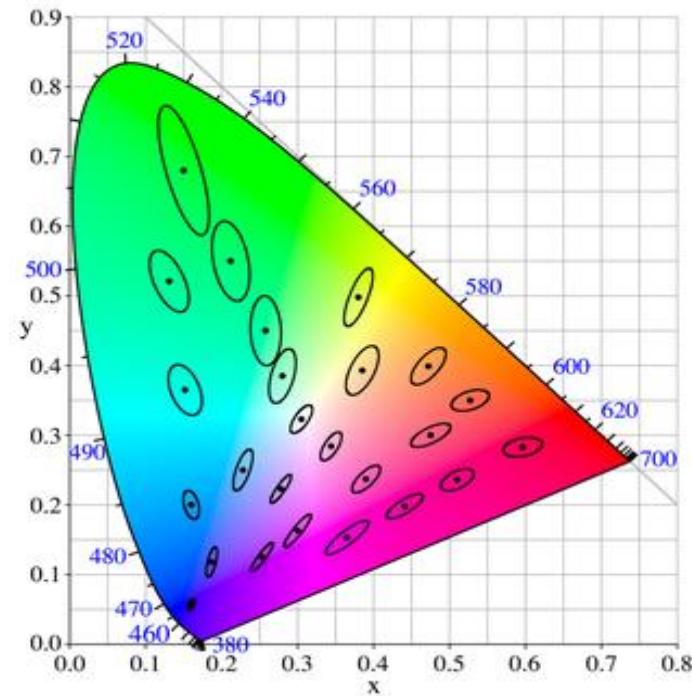
- Do distances between points in the colourspace make sense perceptually?



# Distances in colorspaces

---

- Not necessarily: CIE XYZ is **nonuniform colorspace** – *Euclidean distance between coordinates of colors in colorspace is not a good indicator of color similarity (in terms of human perception).*



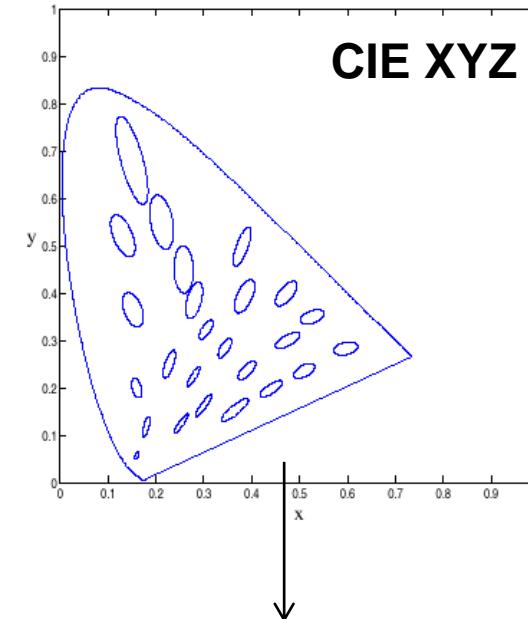
McAdam ellipses:  
Just (human) noticeable differences in color

# Uniform colour spaces

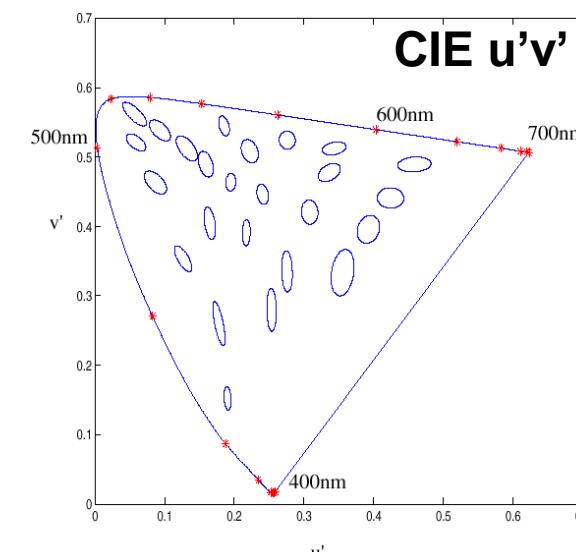
- Transforms such that ellipses are mapped into circles

→ *distances better replicate the human perception of color similarity.*

- Examples of uniform colour spaces:
  - CIE  $u'v'$
  - CIE Lab



Nonuniform colour space



Uniform colour space

Machine perception

# **COLOR SIMILARITY USING HISTOGRAMS**

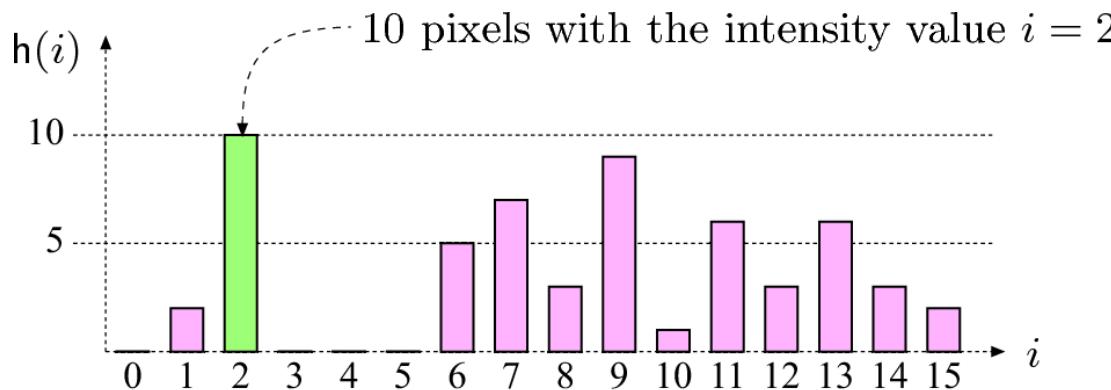
# What is a histogram?

- Image histogram records the frequency of intensity levels

$h(i)$  = the *number* of pixels in  $I$  with the intensity value  $i$

$$h(i) = \text{card}\{(u, v) \mid I(u, v) = i\}$$

- Example:

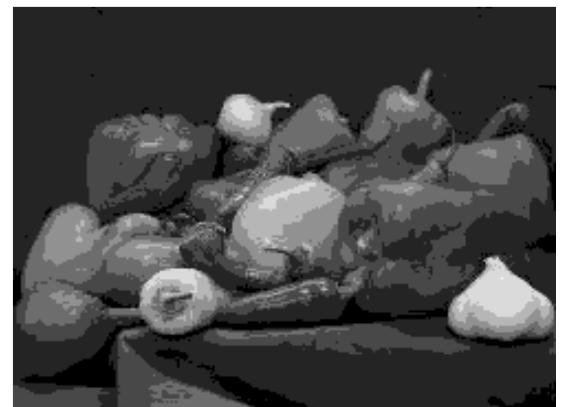


$h(i)$	0	2	10	0	0	0	5	7	3	9	1	6	3	6	3	2
$i$	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Intensity value																

256 intensity levels



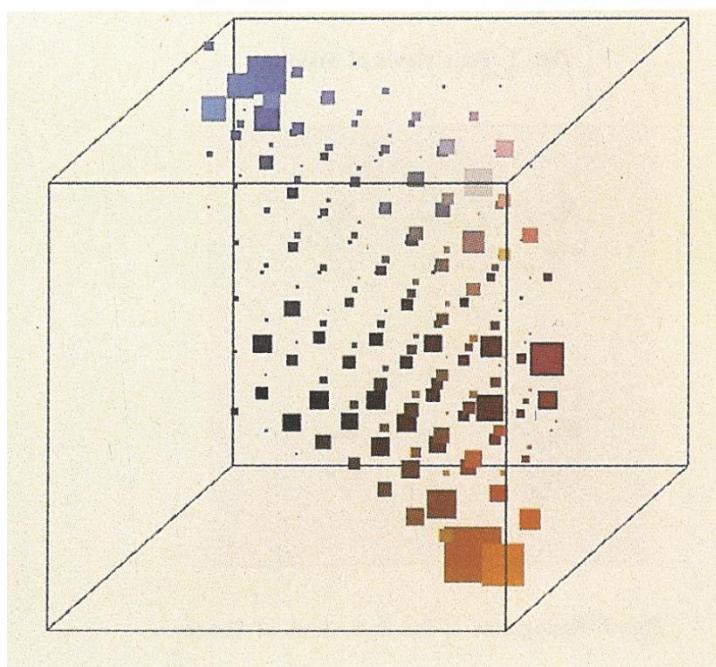
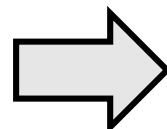
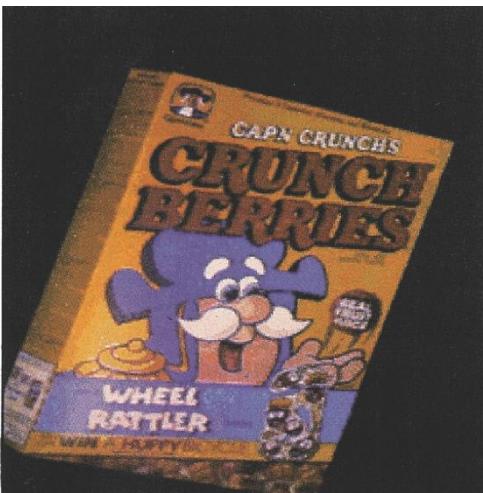
16 intensity levels



# Color histogram

---

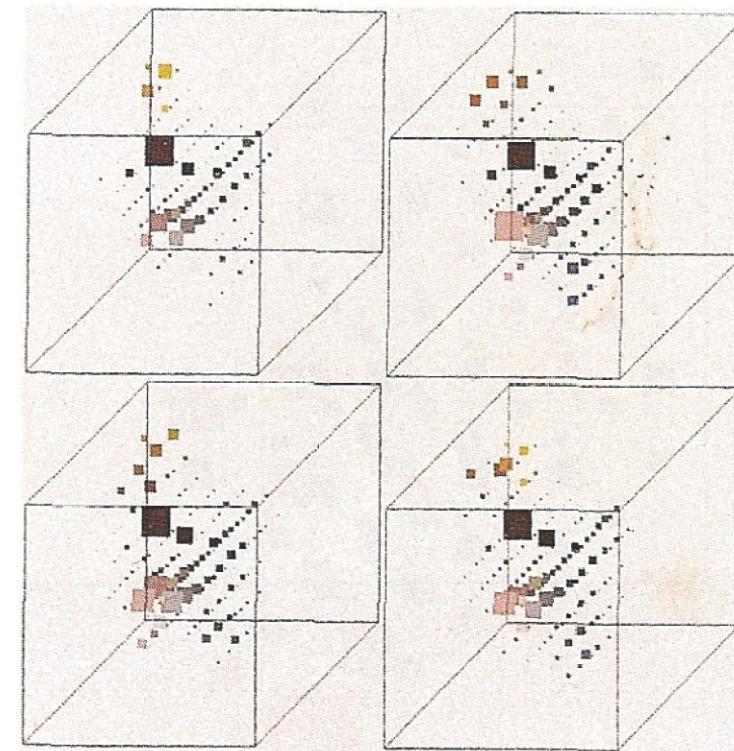
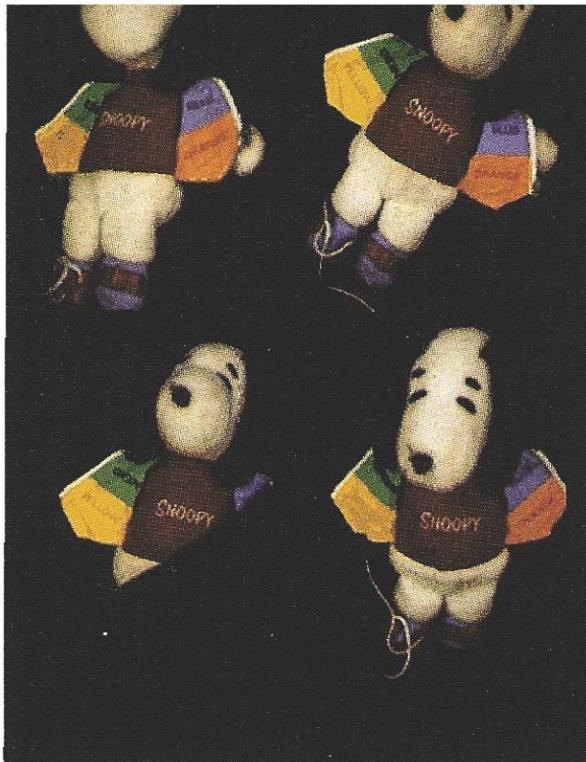
- Color statistic
  - Example of a 3D RGB histogram  $H(R, G, B)$  visualization
  - Each pixel color is a point in 3D space (RGB)
  - Calculate the 3D color histogram
    - $H(R, G, B) = \text{number of pixels with color } [R, G, B]$



# Color histogram

---

- Robust representation of images
  - Translation, scale, partial occlusion



# Intensity normalization

---

- Intensity is contained in each color channel
  - Multiplying a color by a scalar changes the intensity but not really the true „color“ (hue).
  - This means that we can normalize a color by its intensity.
    - Intensity is defined as:  $I = R + G + B$ :
    - Chromatic representation:

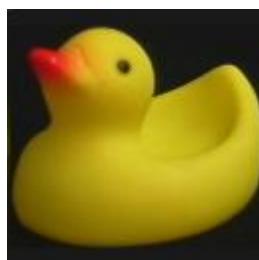
$$r = \frac{R}{R+G+B} \quad g = \frac{G}{R+G+B} \quad b = \frac{B}{R+G+B}$$

- We can now use only a 2D space (rg), since it holds that

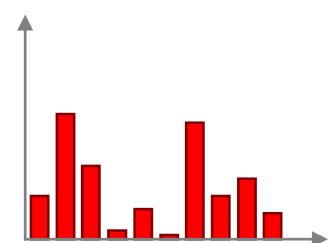
$$r + g + b = 1$$

# Color comparison via histograms

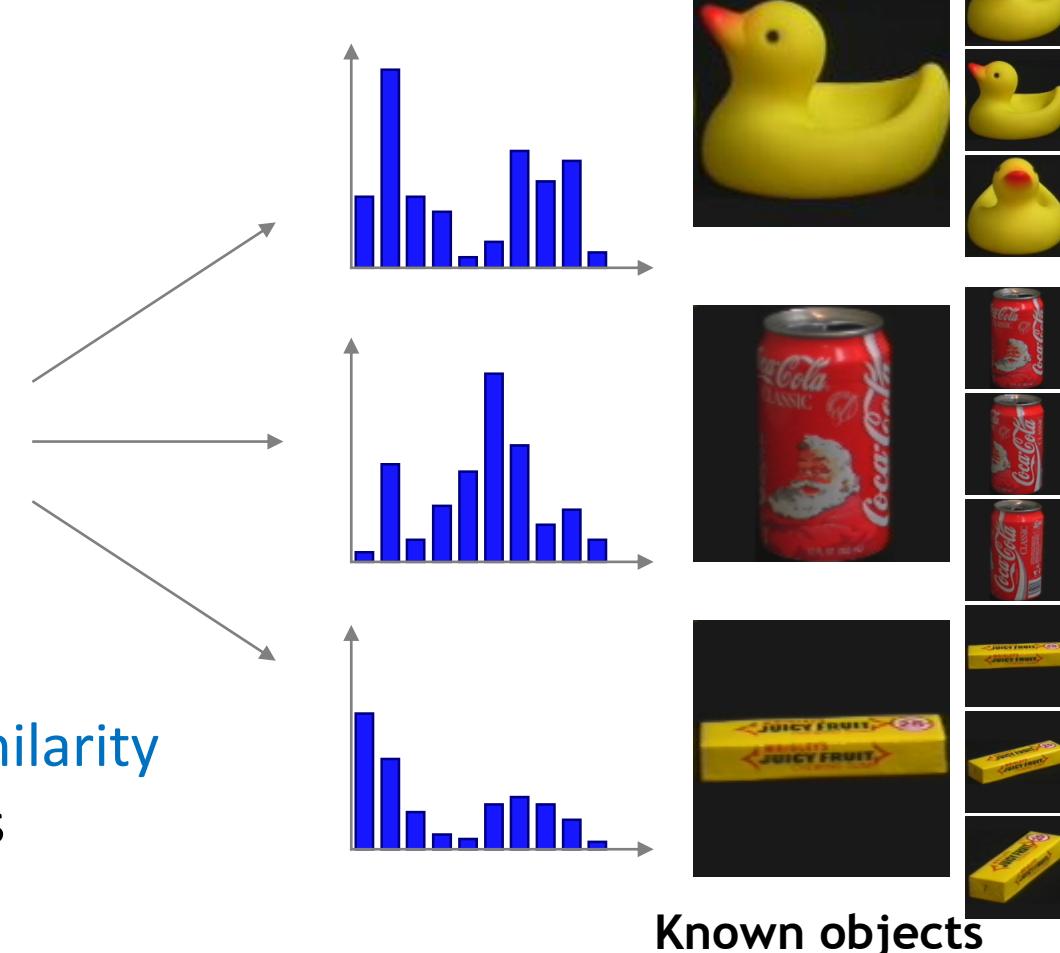
- Compare images indirectly – compare only their descriptors (histograms)



Test image



A measure of **distance/similarity** between the histograms is required!



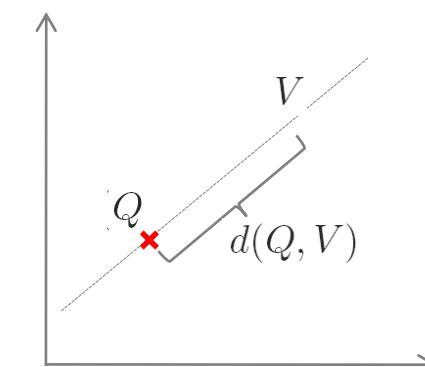
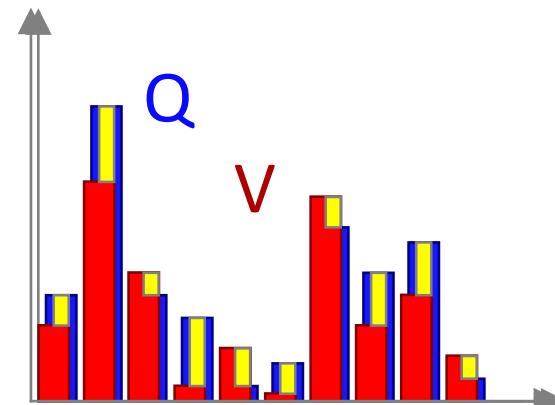
Known objects

# Popular distances: Euclidean distance

- Definition (=L<sub>2</sub> norm)

$$d(Q, V) = \sum_i (q_i - v_i)^2$$

- Explanation
  - Looks for differences in histogram cells.
  - Interpretation: Distance in feature space.
  - Range of output values: [0,1]
  - All cells receive equal weight.
  - Susceptible to noise!



# Popular distances: pdf similarity

---

- Similarity between two probability density functions
- Chi-squared (slo., hi-kvadrat):

$$\chi^2(Q, V) = \sum_i \frac{(q_i - v_i)^2}{q_i + v_i}$$

WATCH OUT FOR  $q_i=v_i=0!!$

- Kullback-Leibler divergence:

$$KL(Q, V) = \sum_i q_i \log \frac{q_i}{v_i}$$

Not a proper metric (not symmetric)

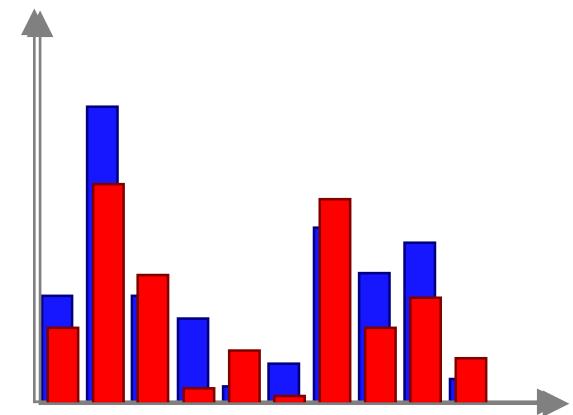
Symmetric version (Jeffreys' divergence):

$$JD(Q, V) = KL(Q, V) + KL(V, Q)$$

- Hellinger distance:

$$d_{\text{Hell}}(Q, V) = \sqrt{1 - \sum_i \sqrt{q_i v_i}}$$

Proper metric, constrained to interval [0,1]



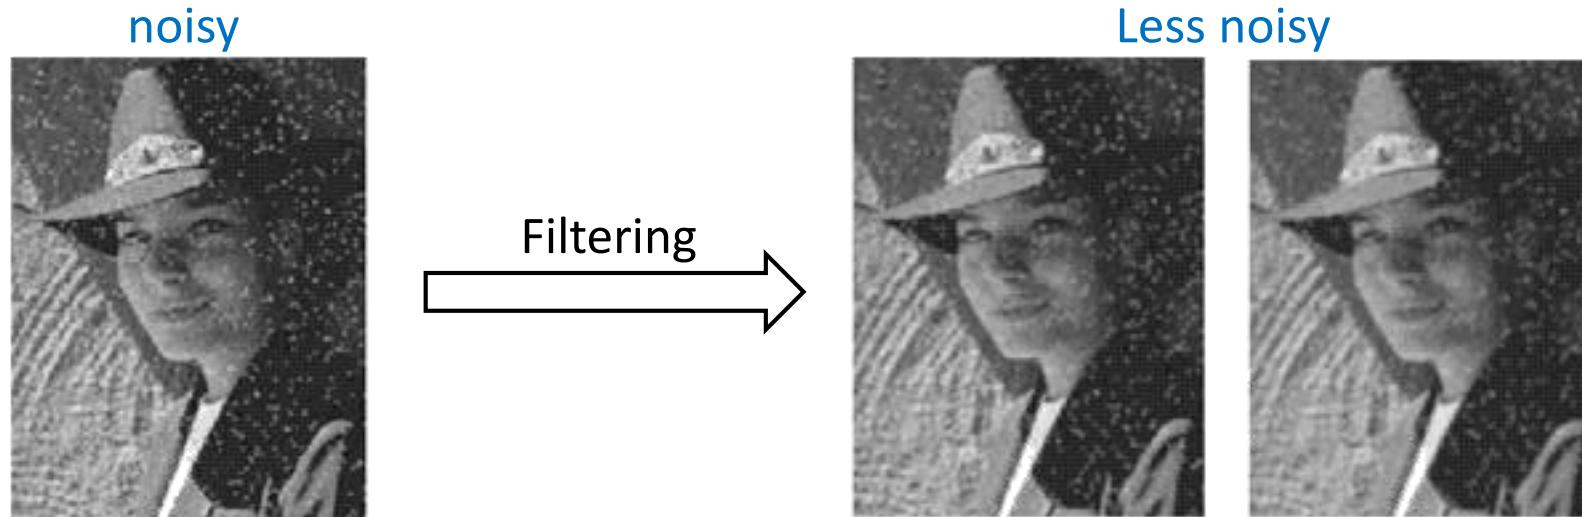
Machine perception

# **FILTERING**

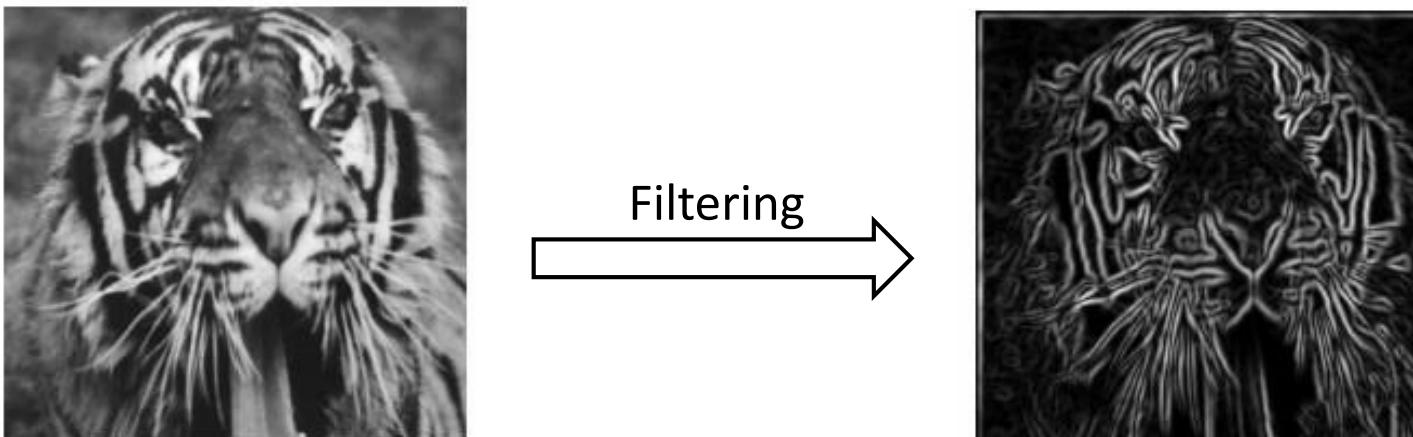
# Can be applied for...

---

- Noise reduction and image restoration



- Structure extraction/enhancement (later in the course)



# Types of image noise

---

- Salt and pepper (*sol in poper*)
  - Random black and white dots.
- Impulse noise (Impulzni šum)
  - Random occurrence of white dots.
- Gaussian noise (Gausov šum)
  - The intensity variation sampled from a Gaussian (Normal) distribution.



Original



Salt and pepper noise



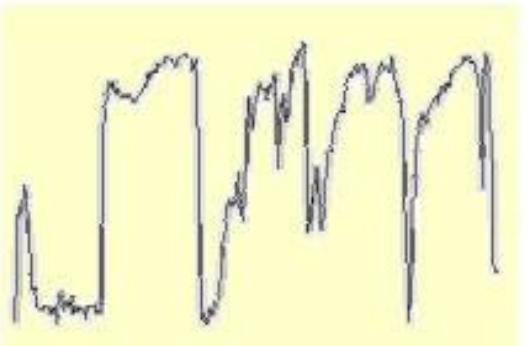
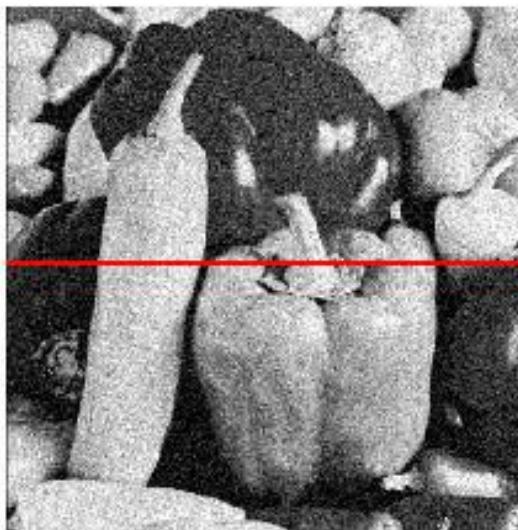
Impulse noise



Gaussian noise

# Gaussian noise

---



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

Matlab:

```
>> im = imread("peppers.jpg");  
>> noise = randn(size(im)).*5;  
>> output = im + noise;
```

# Let's try to remove the noise...

---

- Assumption:
  - Pixels are **similar** to their **neighboring pixels**
  - The **noise** is **independent** among pixels  
("i.i.d. = independent, identically distributed")
- So let's compute an **improved estimate** of pixel's intensity by **replacing** it with an **average of pixel intensities** in its immediate **neighborhood...**



# A moving average 2D

---

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

# A moving average 2D

---

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

			0	10	20					

# A moving average 2D

---

 $F[x, y]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $G[x, y]$ 


# A moving average 2D

---

 $F[x, y]$ 

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$ 

			0	10	20	30	30				

# A moving average 2D

---

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

# A moving average 2D

---

- Assume the **averaging window** size is  $2k+1 \times 2k+1$ :

$$G[i, j] = \frac{1}{(2k+1)^2} \underbrace{\sum_{u=-k}^k}_{\text{Equal weights for all pixels.}} \underbrace{\sum_{v=-k}^k}_{\text{A loop over all pixels within the neighbourhood of } F[i,j].} F[i + u, j + v]$$

- Now let's generalize this by making a **weight depend** on **relative position** from the central element.

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] \underbrace{F[i + u, j + v]}_{\text{Nonuniform weights}}$$

# Correlation filtering

---

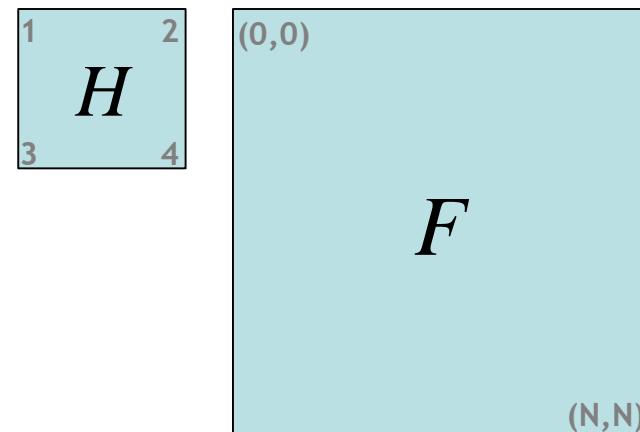
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

- This is called **cross-correlation** and abbreviate as:

$$G = H \otimes F$$

- Image filtering

- Replace image **intensity** with a **weighted sum** of a window centered at that pixel.
  - The weights in the **linear combination** are prescribed by the **filter's kernel**.



# Convolution as correlation

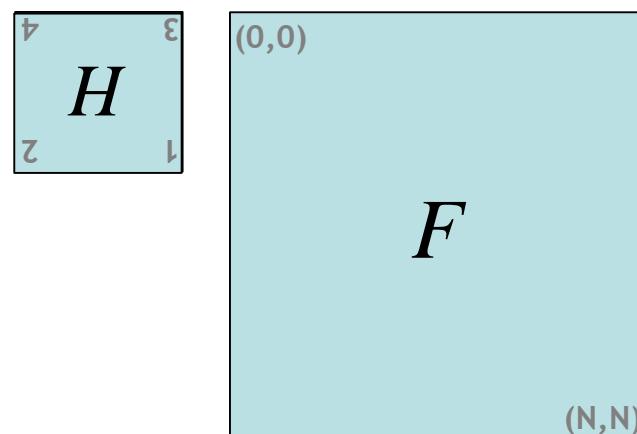
---

- Compute convolution by cross-correlation:
  - Flip the filter in both dimensions (horizontal + vertical)
  - Apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

$$G = H \star F$$

*↑*  
*convolution  
operator*



# Convolution vs. Correlation

---

- Correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$



Notice the difference?

Matlab:  
`filter2`  
`imfilter`

- Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Matlab:  
`conv2`

(we will also use “ $*$ ” to denote convolution,  
i.e.,  $G = H * F$ .)

- Comment:

- For a **symmetric** filter,  $H[-u, -v] = H[u, v]$ , correlation  $\equiv$  convolution.

# Linear shift-invariant system

---

- Shift-invariant:
  - Equal behavior **irrespective of position**, i.e., the output only **depends on the local pattern**, but not image coordinates.
- Linearity:
  - **Superposition:**  $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$
  - **Scaling:**  $h * (kf) = k(h * f)$  for a scalar  $k$

# Properties of convolution

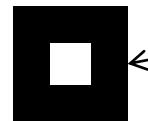
---

- Linear, shift-invariant
- Commutative:  $f * g = g * f$
- Associative:  
$$(f * g) * h = f * (g * h)$$
  - As result, application of **multiple filters** is **equal to** application of a **single filter** :  
$$((f * b_1) * b_2) * b_3 = f * (b_1 * b_2 * b_3)$$
- Identity:  
$$f * e = f$$
  - A **unit impulse**:  $e = [..., 0, 0, 1, 0, 0, ...]$ .
- Derivative:

$$\frac{\partial}{\partial x} (f * g) = \left( \frac{\partial}{\partial x} g \right) * f = \left( \frac{\partial}{\partial x} f \right) * g$$

# Recall smoothing by averaging

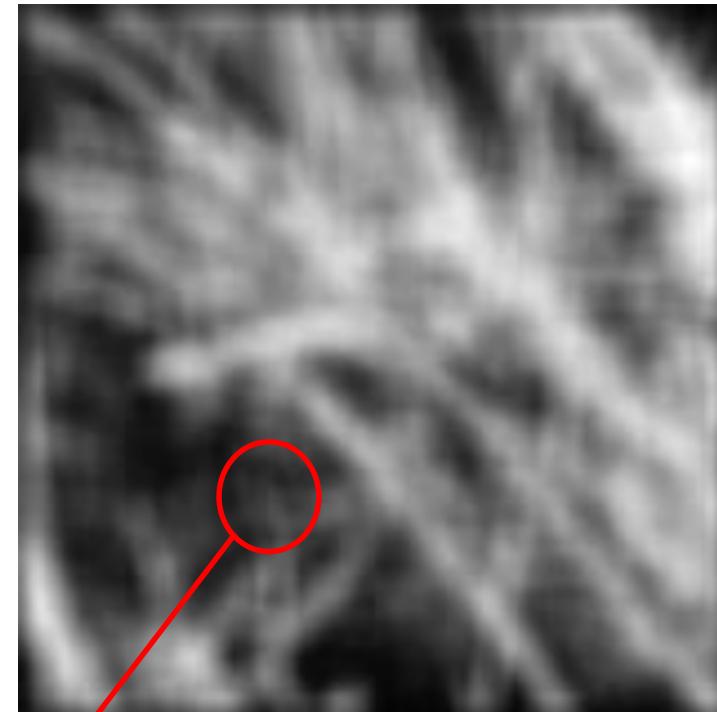
---



This is a box filter:  
white = large weight,  
black = low weight



*Original*

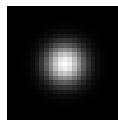


*Filtered*

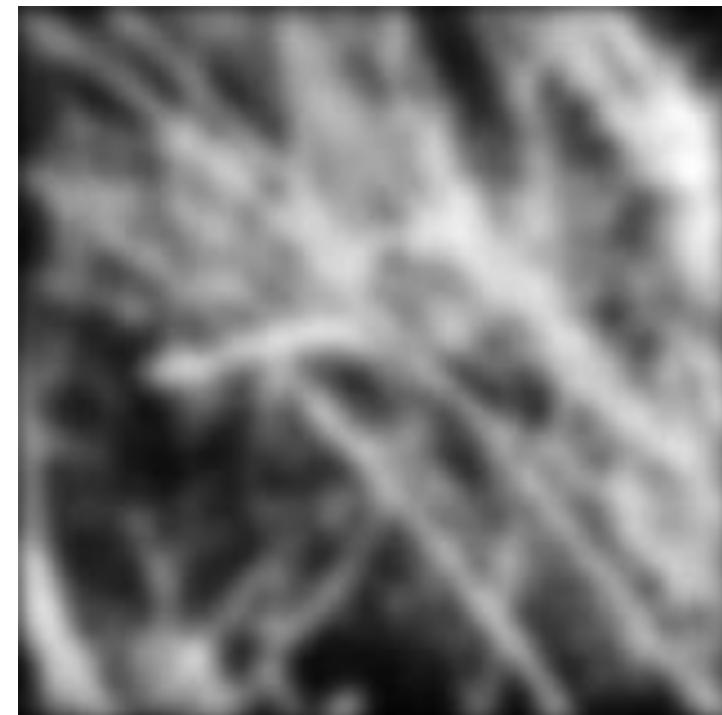
Side lobe artefacts!

# Smoothing by a Gaussian

---



*Original*



*Filtered*

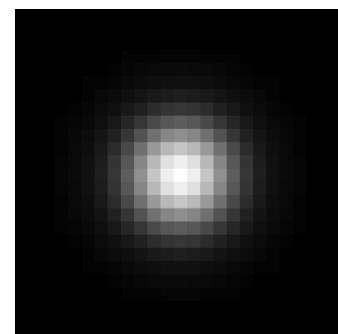
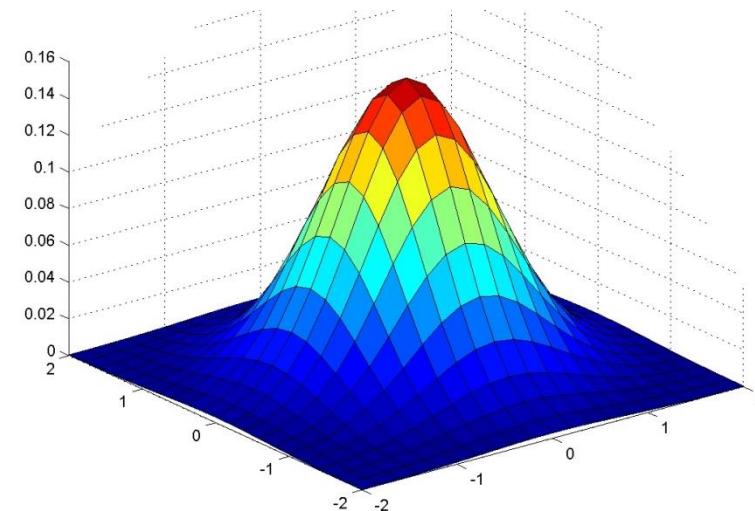
# Gaussian smoothing

---

- Gaussian kernel

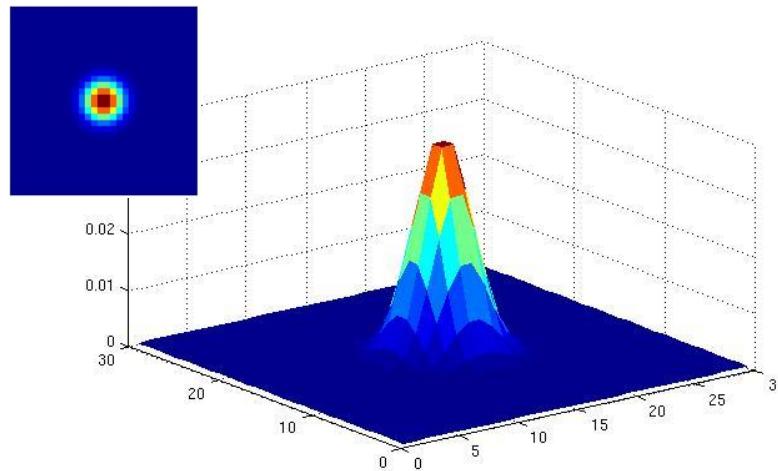
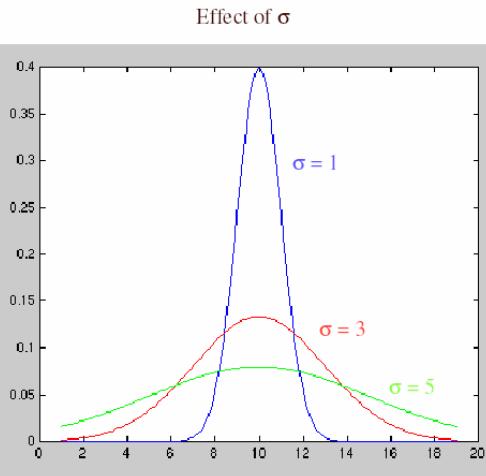
$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Rotation **symmetric**
- Pixels **closer** to center get **higher weight**
  - Makes sense for a probabilistic inference of a signal content.

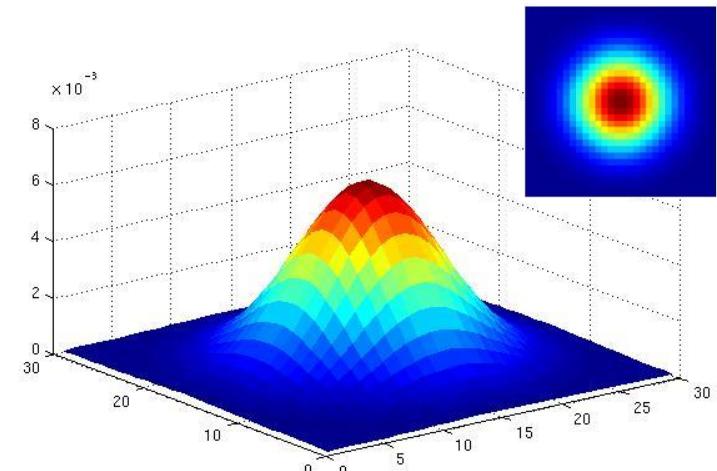


# Gaussian smoothing

- How about parameters?
- *Variance  $\sigma^2$*  determines the extent of smoothing...



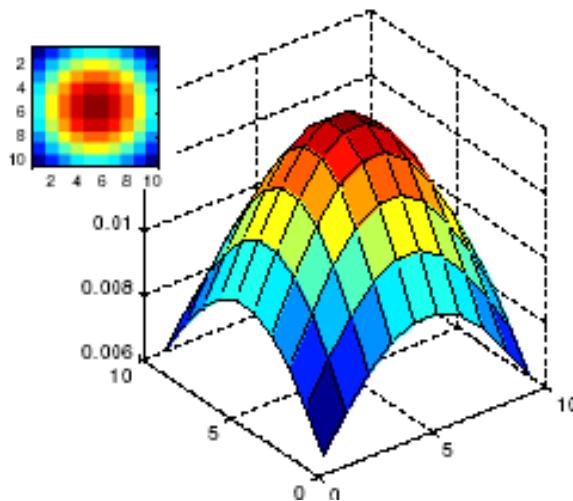
$\sigma = 2$  by kernel  
30x30



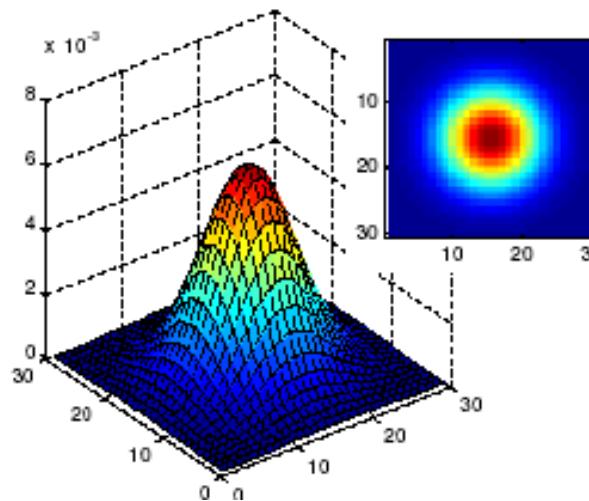
$\sigma = 5$  by kernel  
30x30

# Gaussian smoothing

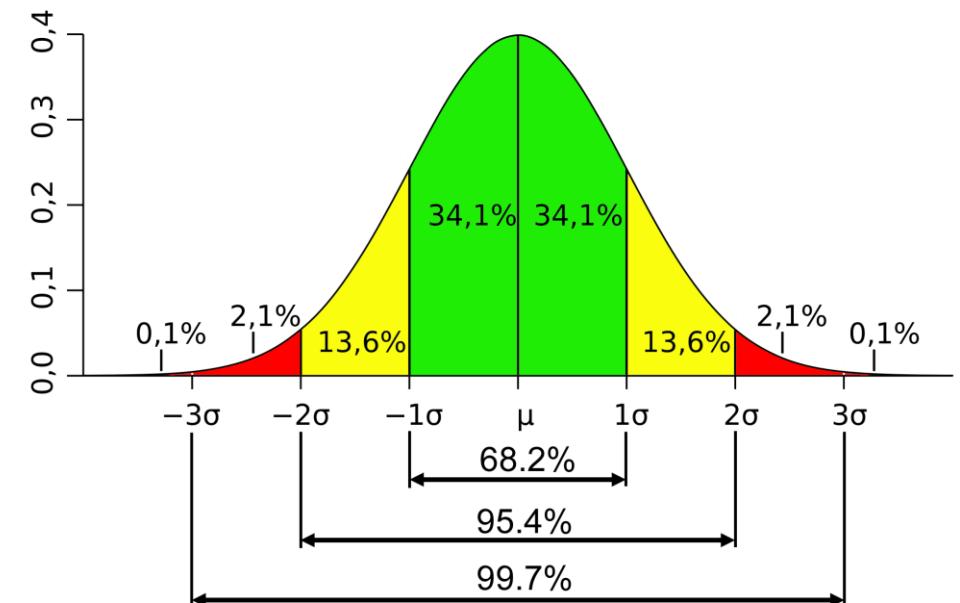
- How about **parameters**?
- *Kernel size!*
  - Infinite support, but discretization makes it **finite**.



$\sigma = 5$  with  $10 \times 10$  kernel



$\sigma = 5$  with  $30 \times 30$  kernel

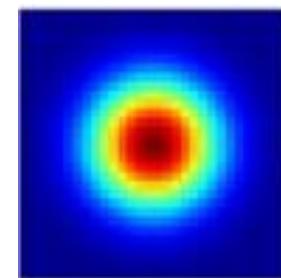
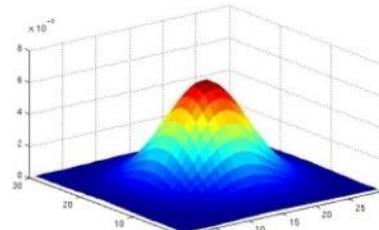


- Rule of thumb: set half size of the kernel to  $3\sigma$

# Gaussian filtering in Matlab

---

```
>> sigma = 5;  
>> hsize = 2*sigma*3+1;  
>> h = fspecial('gaussian', hsize, sigma);  
  
>> figure(1); surf(h);  
  
>> figure(2); imagesc(h); axis equal;  
  
>> outim = imfilter(im, h);  
>> figure(3); imshow(outim);
```



im

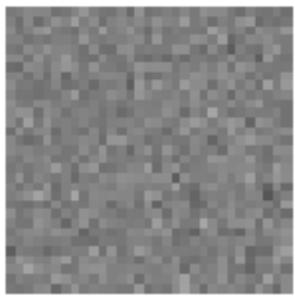


outim

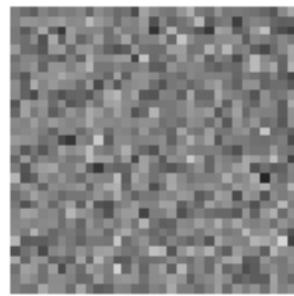
# Effects of smoothing

Increasing the noise extent →

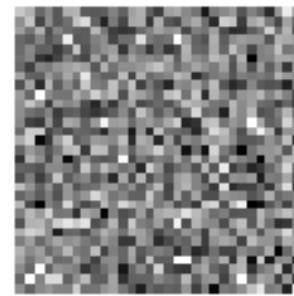
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no  
smoothing



Increasing the kernel size →



# Efficient implementation

- Both, Uniform as well as Gaussian **kernels are separable**:
  - Apply convolution at each row separately using a 1D kernel:

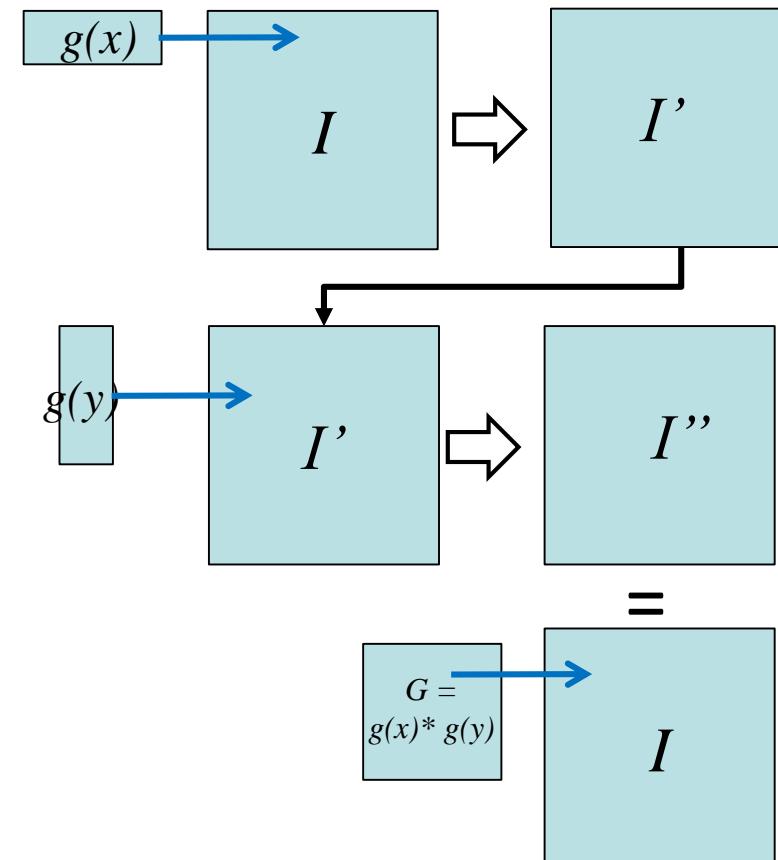
$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-x^2/(2\sigma^2))$$

- Next apply a 1D convolution at each column:

$$g(y) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-y^2/(2\sigma^2))$$

- **Why** is this separation possible?
  - Convolution is linear, associative + commutative

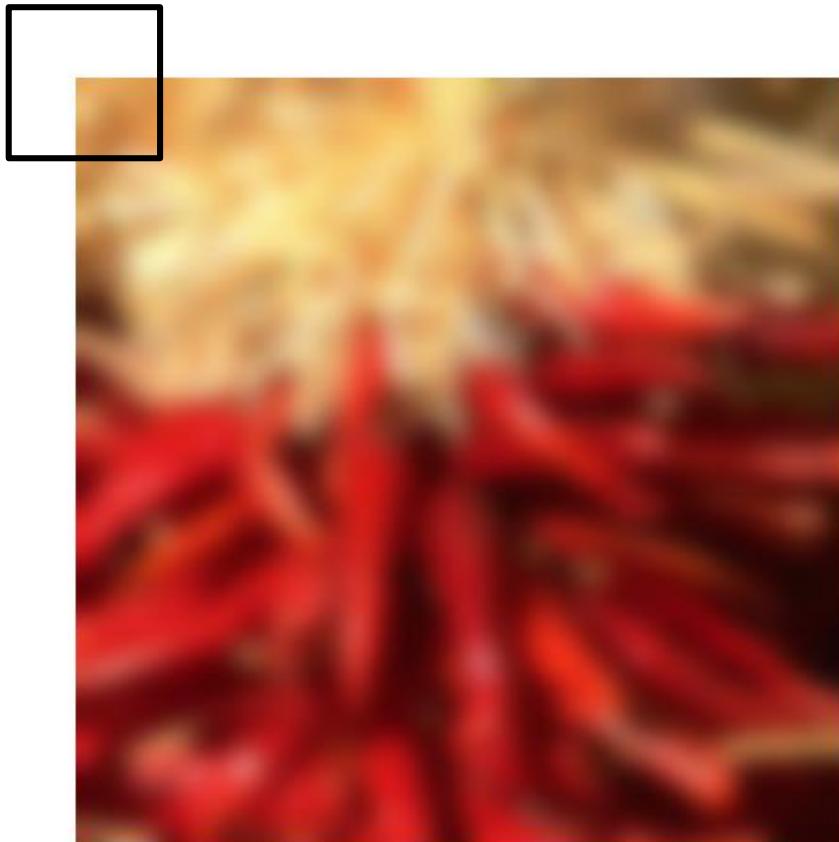
$$g_x * (g_y * I) = (g_x * g_y) * I$$



# Filtering: Boundary conditions

---

- What to do at the image **boundaries**?
  - The kernel **exceeds image boundaries** at the edge
  - Need for **extrapolation**
  - Methods:
    - Crop (black)
    - Bend image around
    - Replicate edges
    - Mirror image



# Filtering: Boundary conditions

---

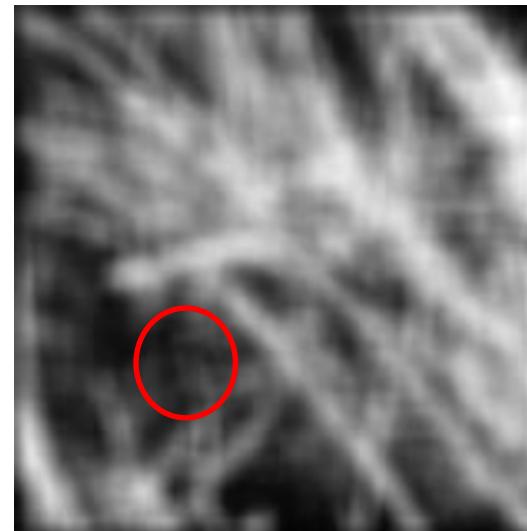
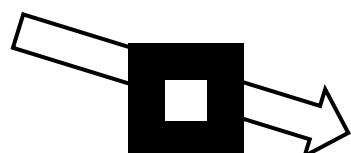
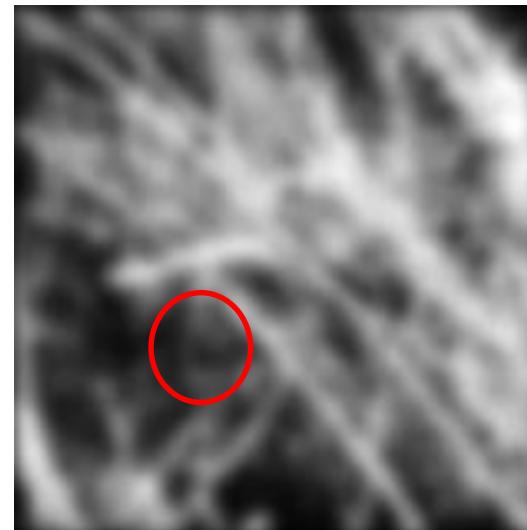
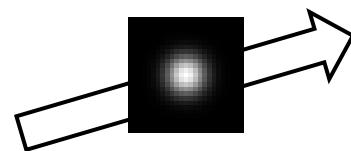
- What to do at the image boundaries?
  - The kernel exceeds image boundaries at the edge
  - Need for extrapolation
  - Methods (Matlab):
    - Crop (black):              `imfilter(f,g,0)`
    - Bend image:                `imfilter(f,g, 'circular')`
    - Replicate edges:          `imfilter(f,g, 'replicate')`
    - Mirror image:             `imfilter(f,g, 'symmetric')`

# Strange artefacts in convolution results...

---



*Original*



*Filtered*

# Convolution and spectrum

---

- Convolution of two functions in image space is equivalent to the product of their corresponding Fourier transforms (spectra).

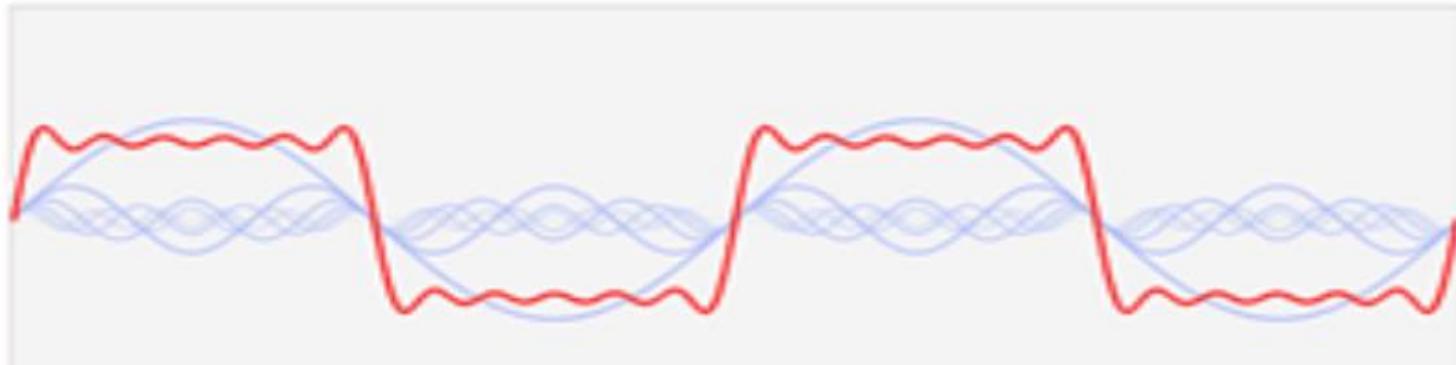
$$\mathcal{F}(f * g) = \mathcal{F}(f) \odot \mathcal{F}(g)$$

↑      ↓  
Image  $f$       Fourier transforms  
and filter  $g$       of  $f$  and  $g$ .

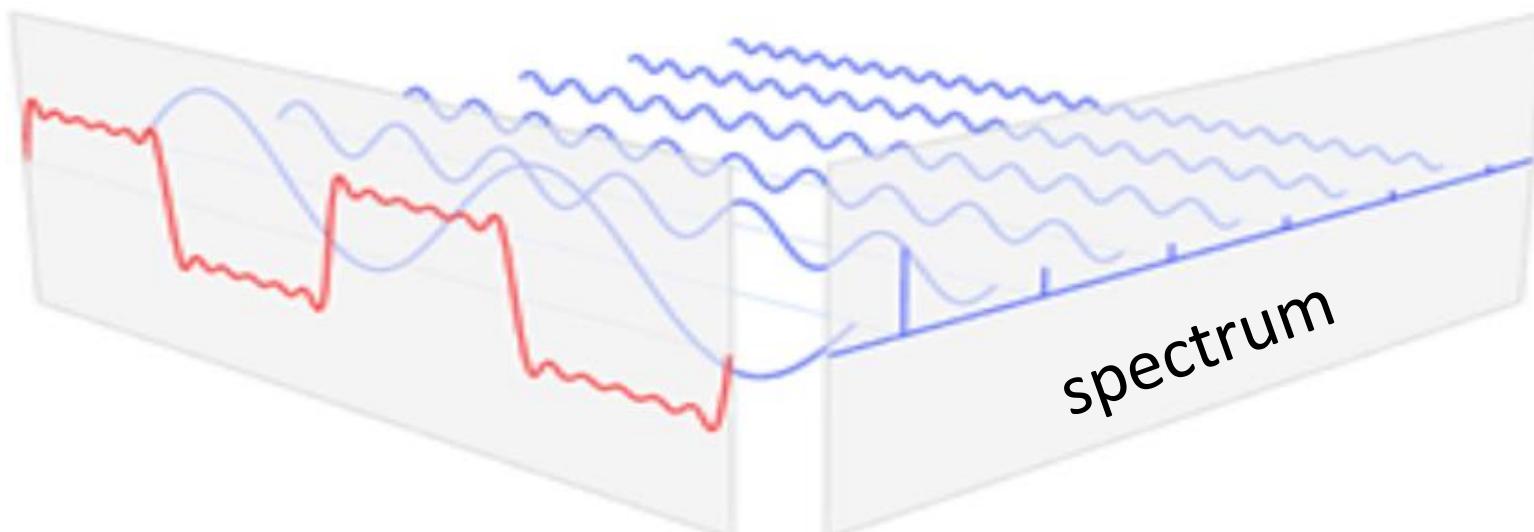
- Convolution manipulates the image spectrum
  - Enhancing/suppressing frequency bands in image.

# Recall the Fourier transform

A signal is represented as a sum of sines/cosines of various frequencies

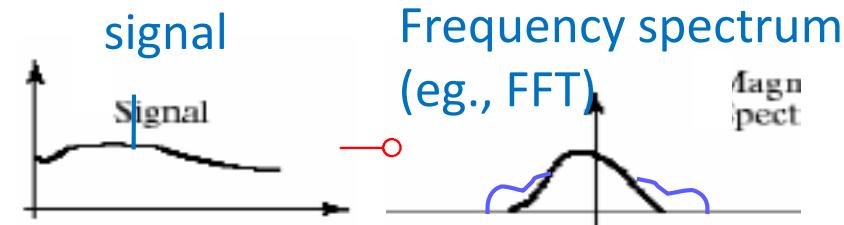


$$f(x) = \sum_n a_n \cos(nx) + b_n \sin(nx)$$

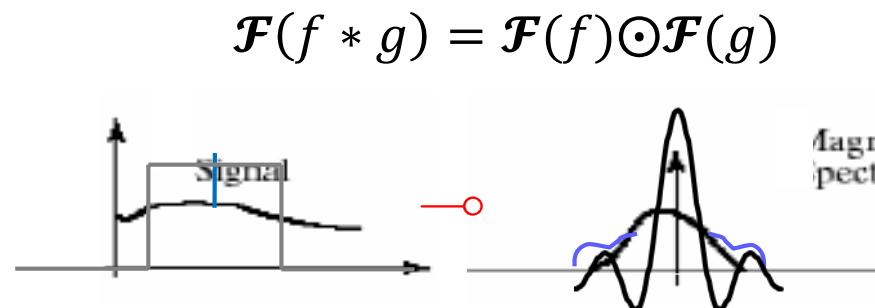


# Convolution: removing noise

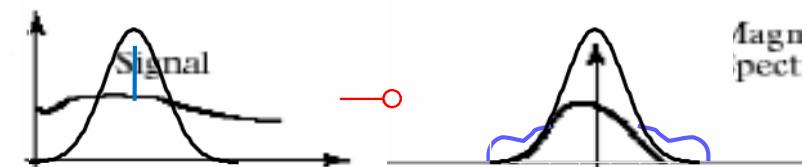
- Noise corresponds to adding **high frequencies**. To **remove** these, we apply a *low-band pass* filter.



- The **spatial box** filter also transforms to a **sinc** in frequency space, causing artefacts (side lobes).

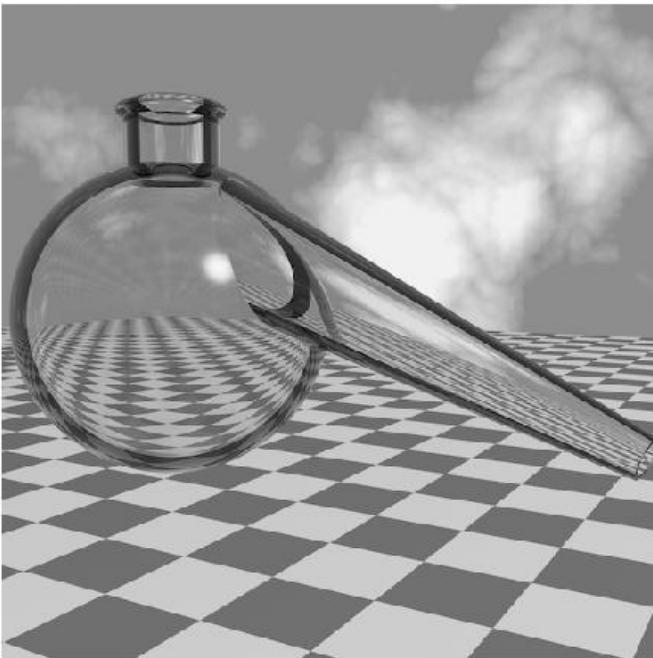


- A **Gaussian** maintains a **compact support** in both image and frequency space. Hence, it's **more appropriate** as a low-band-pass filter.

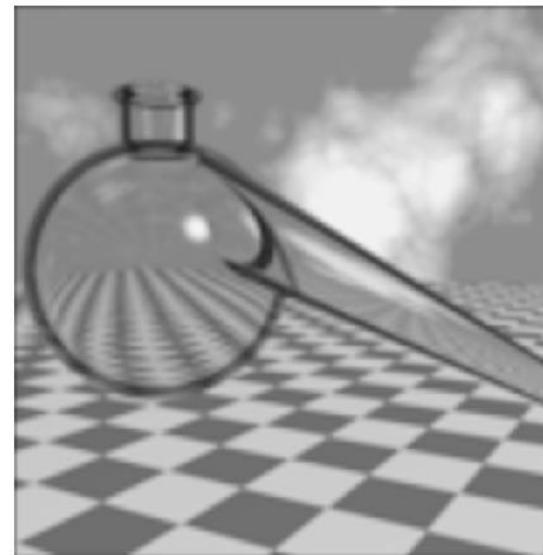


# Low-band vs High-band pass filters

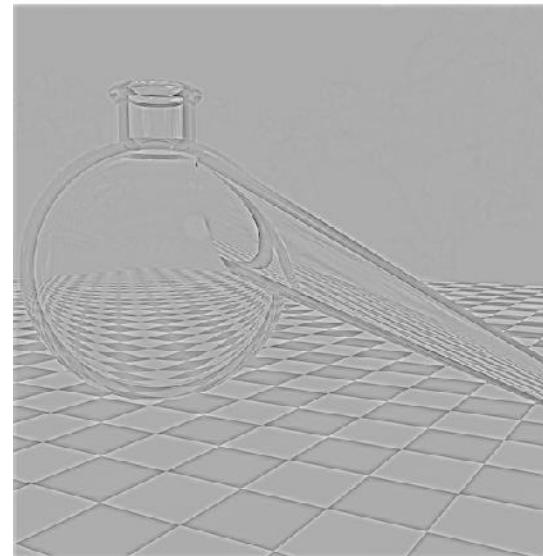
---



Original image



Filtered by a  
low-band-pass filter



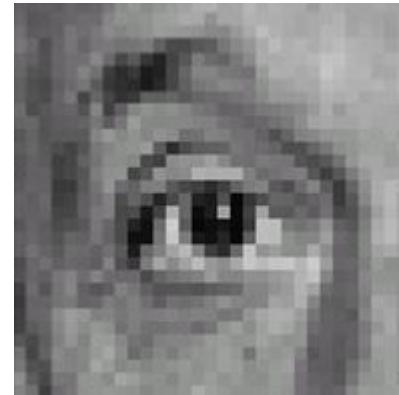
Filtered by a  
high-band-pass filter

# Linear filters in practice

---



$$\text{Original} * \left( \begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} \right) - \frac{1}{9} \left( \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) = \text{Sharpened Image}$$



Original

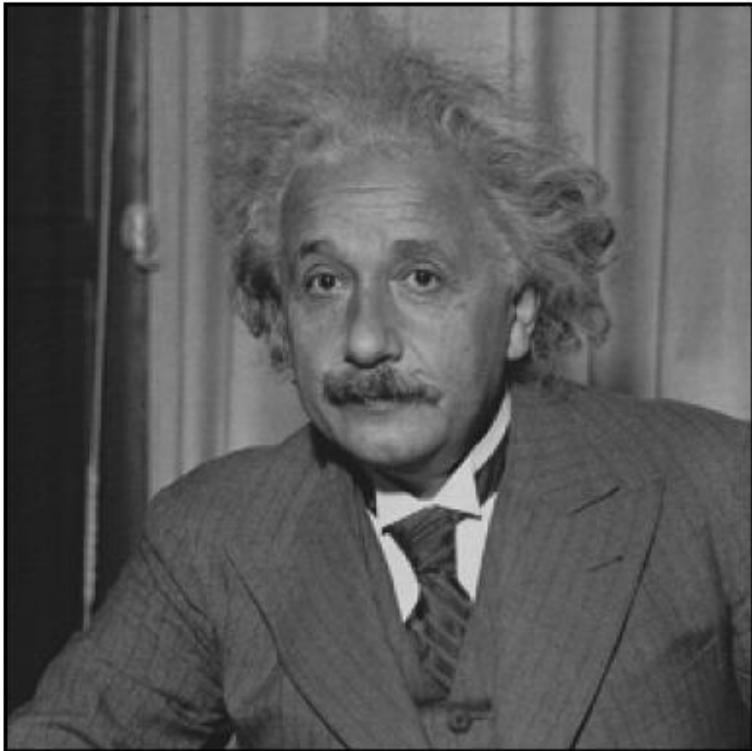
Sharpening filter:  
Enhances differences by local averaging.

To explain this, think about what happens in frequency domain.

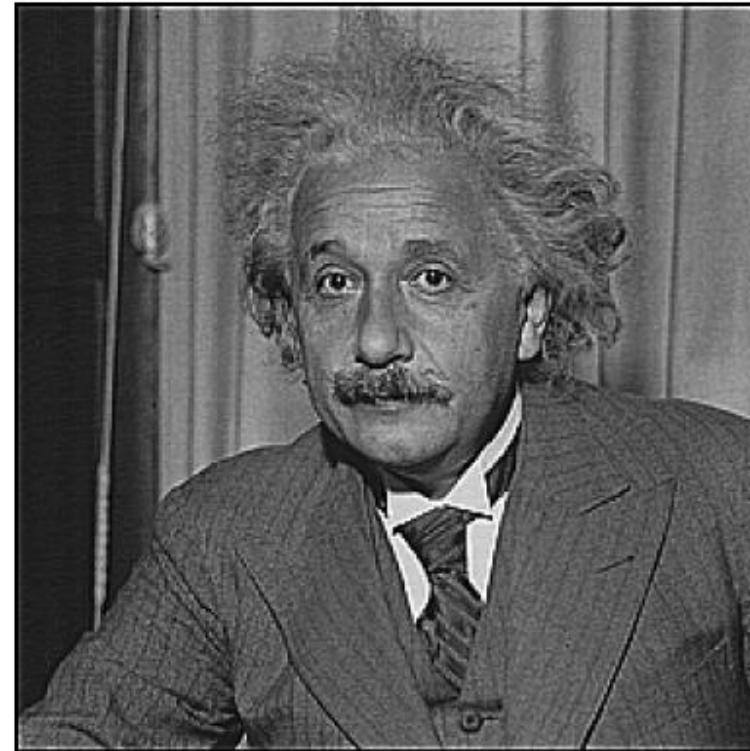
# Sharpening filter

---

before



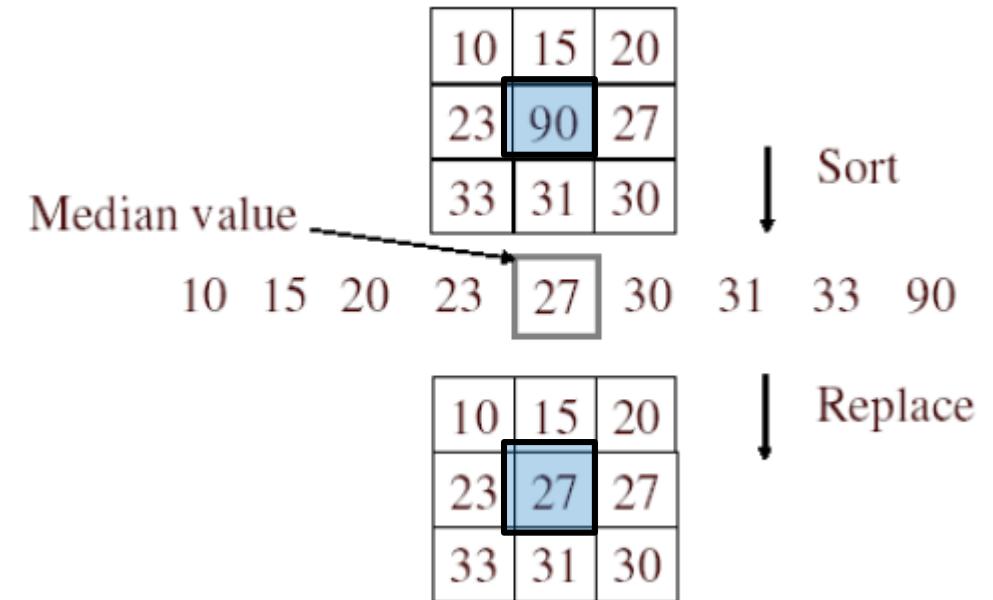
after



To explain this, think about what happens in frequency domain.

# Nonlinear filters: Median filter

- Basic idea
  - Replace the pixel intensity by a median of intensities within a small patch.

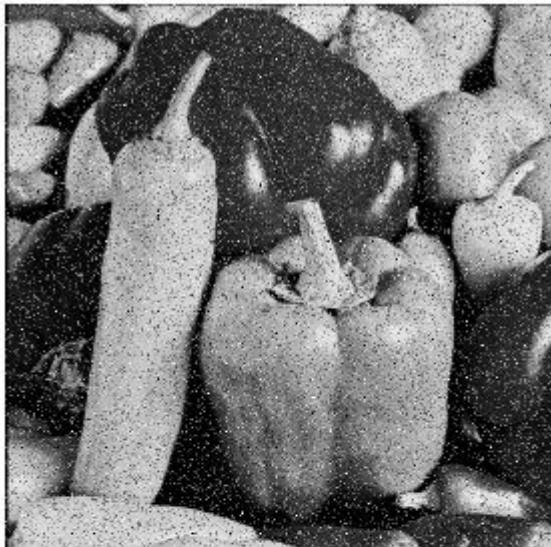


- Properties
  - Does not add new gray-levels into the image.
  - Removes impulse noise: appropriate for impulse noise and salt&pepper noise removal.

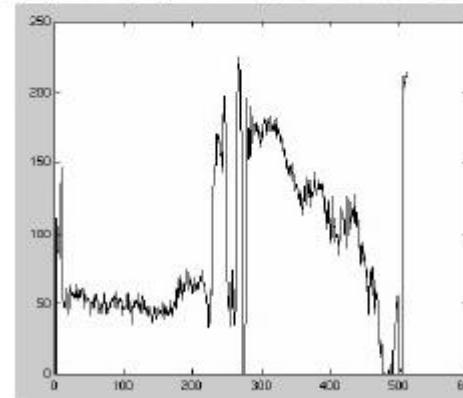
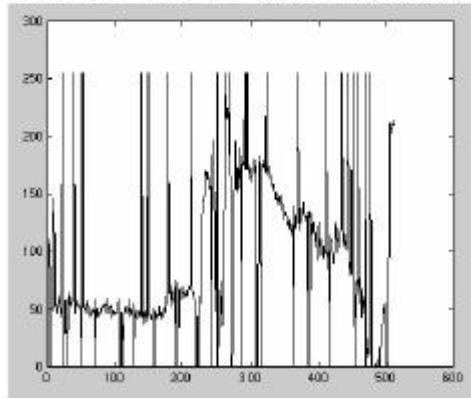
# The Median filter

---

Salt&pepper  
noise



After median  
filtering



Plot of a line in the image

# Median vs. Gaussian

---

Gaussian  
filter

3x3



5x5



7x7



Median  
filter

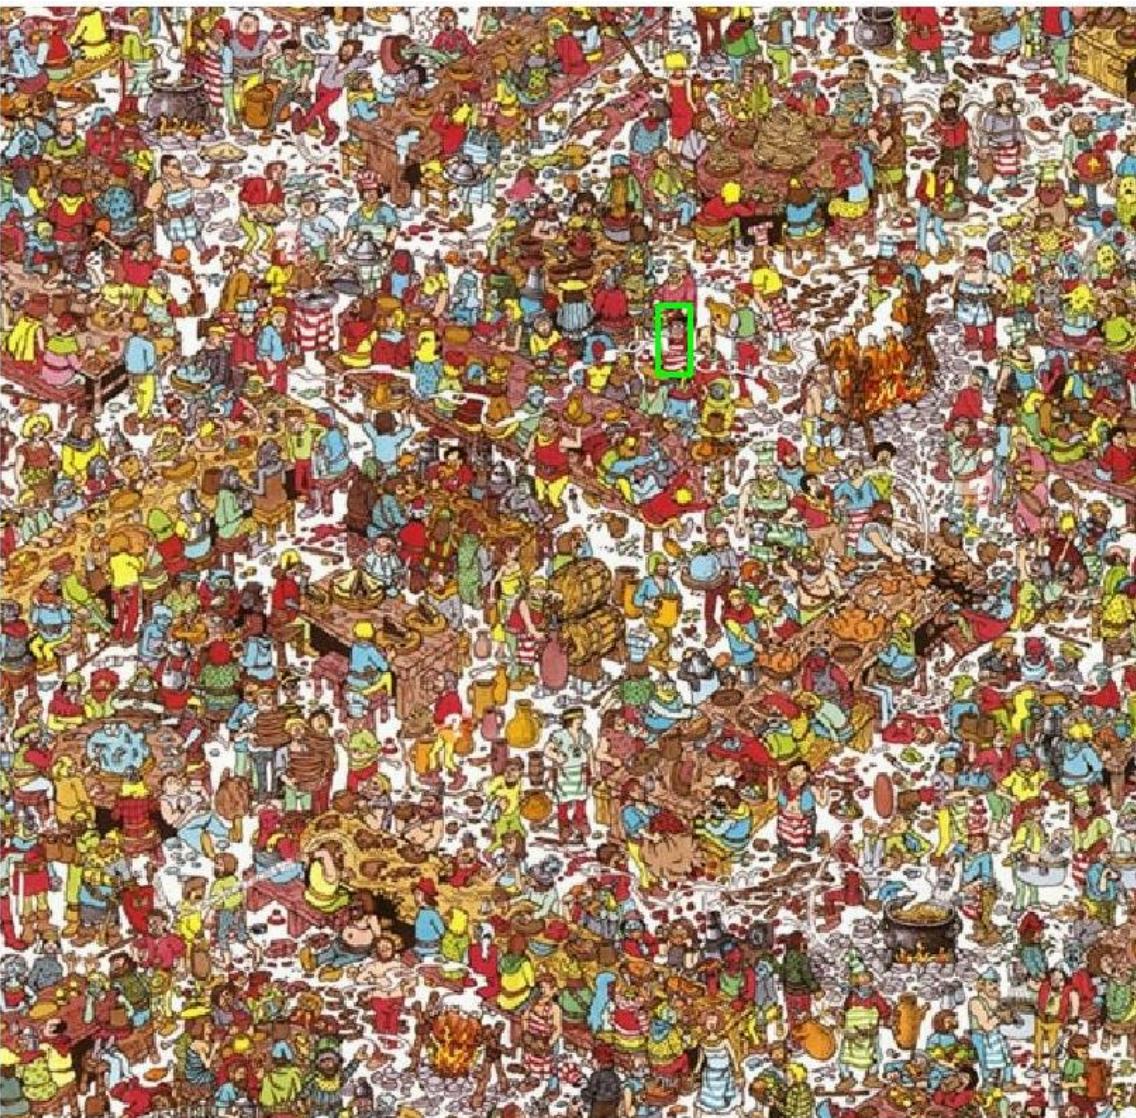


Machine perception

## **FILTERING AS TEMPLATE MATCHING**

# Filtering as template matching

---

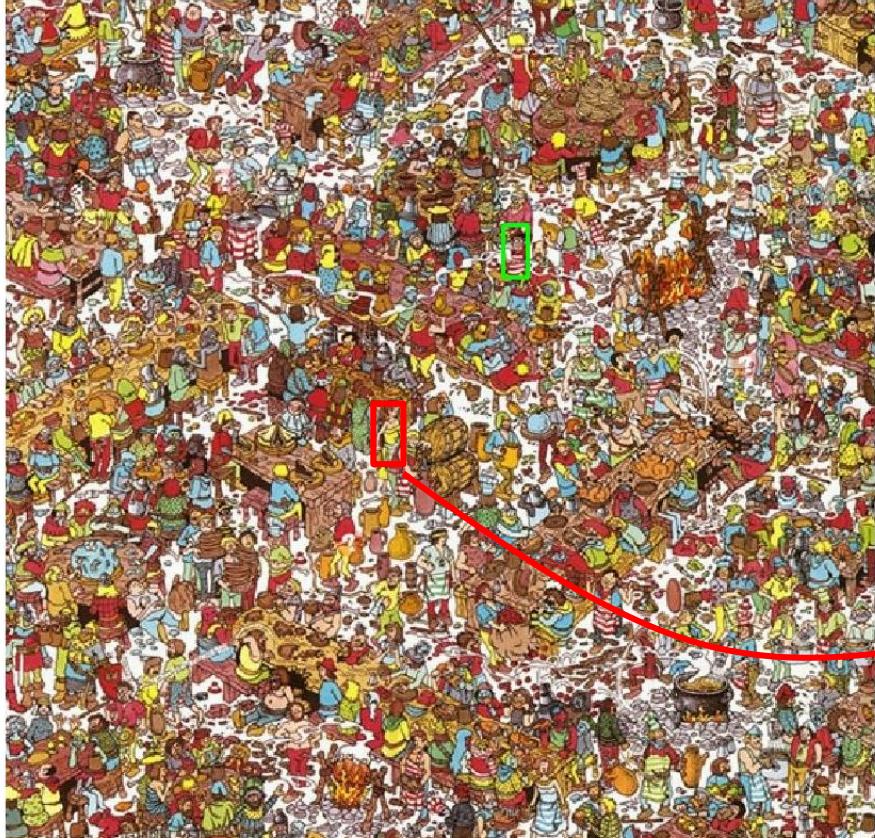


*Where's waldo?*



Template

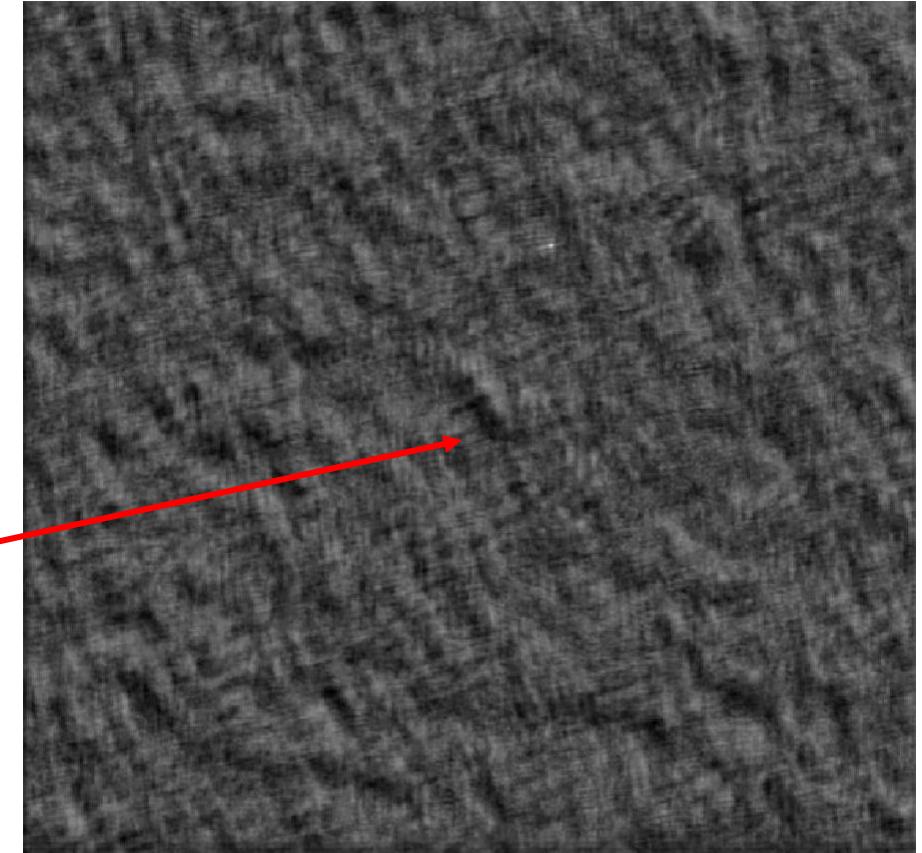
# Apply correlation with template



Input image

Template

$$\Sigma \quad \text{Template} \odot \text{Input} = \text{Correlation map}$$



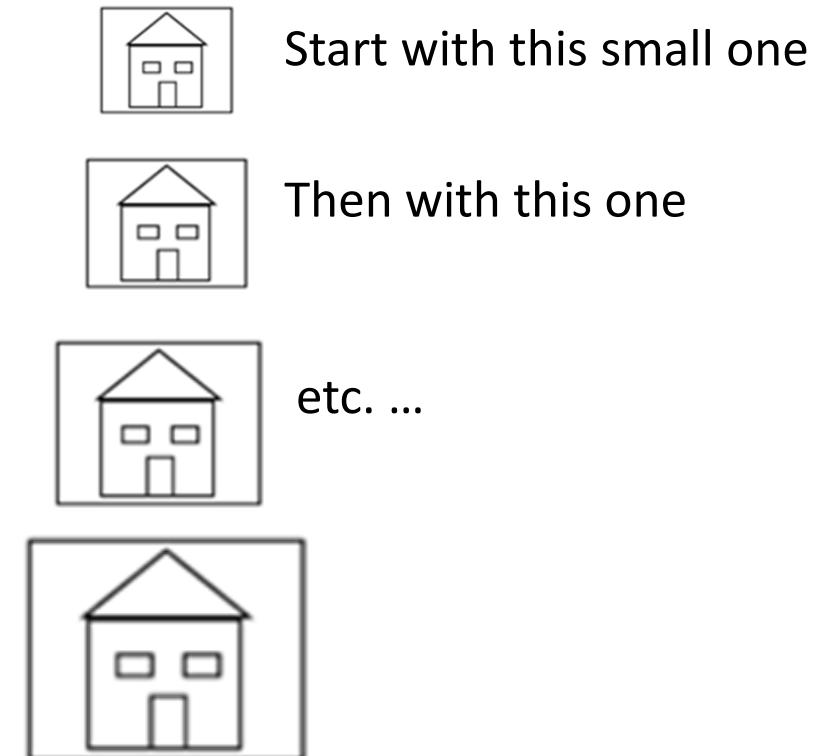
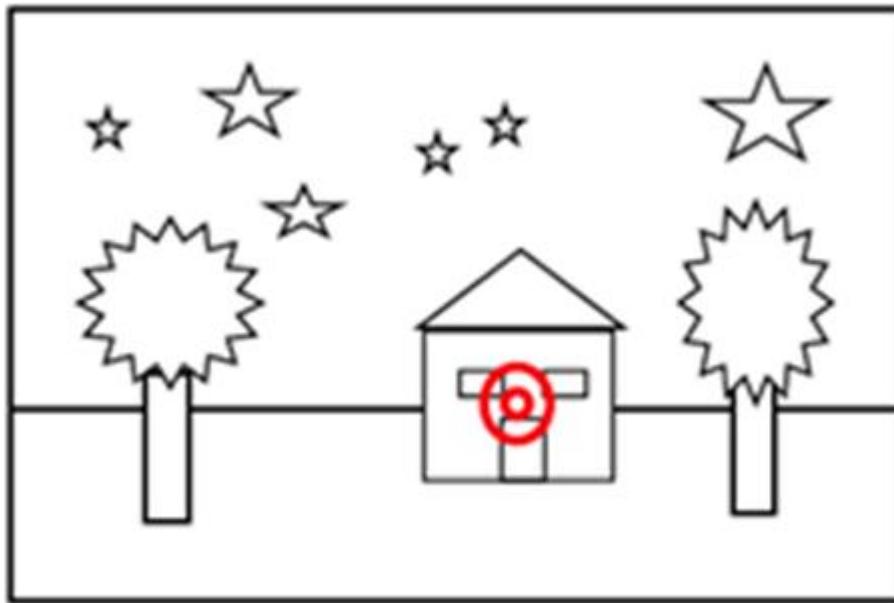
Correlation map

- By the way, this can be done *really* fast by Fast Fourier transform (FFT).

# Issues with template matching over scales

---

- But the object may be bigger/smaller in the image!
- Well, we could carry out correlation for different **scales** of the template...



# Template matching in scale space

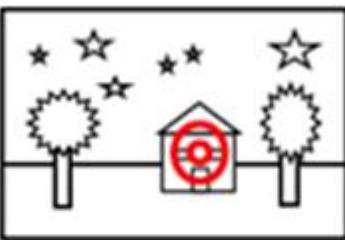
- But rather than template, we **scale the input image**

Reduce the image size

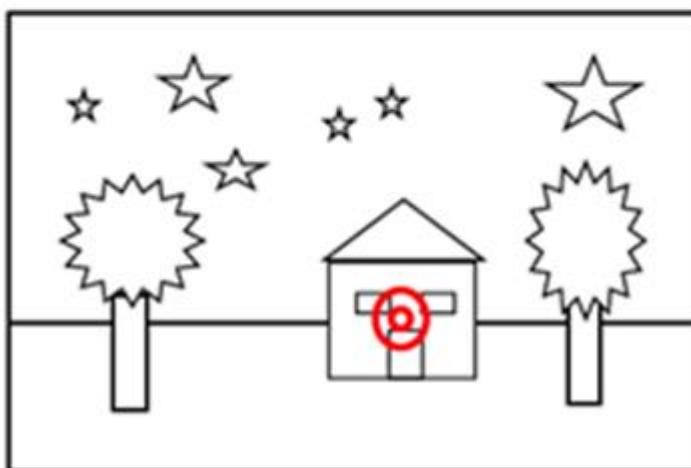


Keep the same size

Reduce the image size



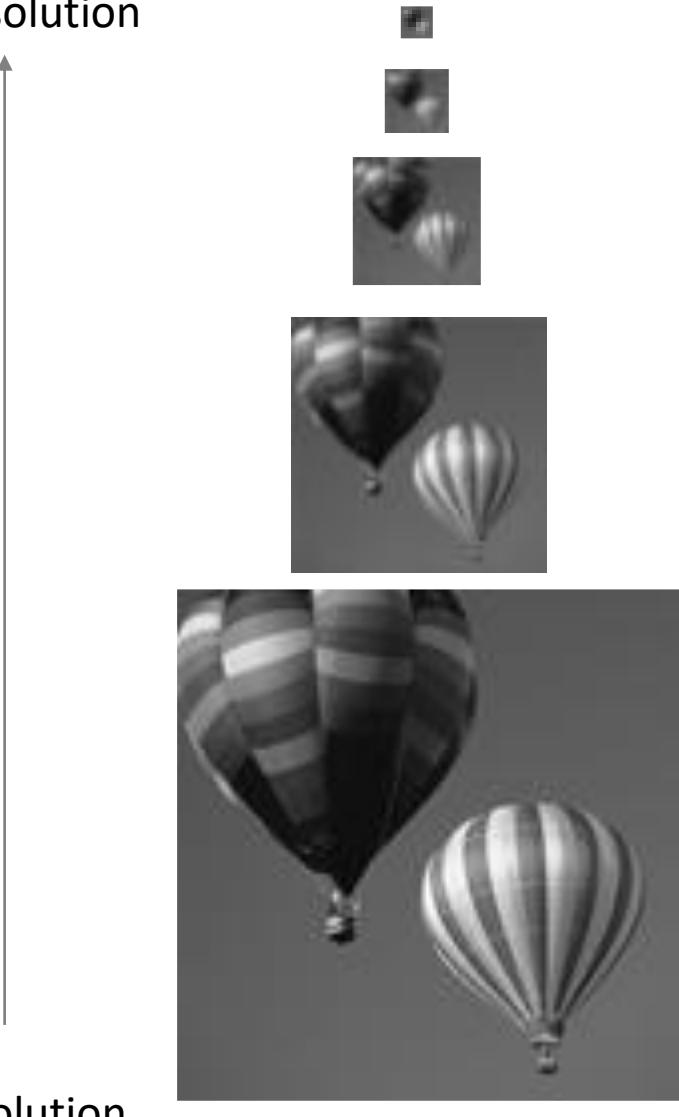
Keep the same size



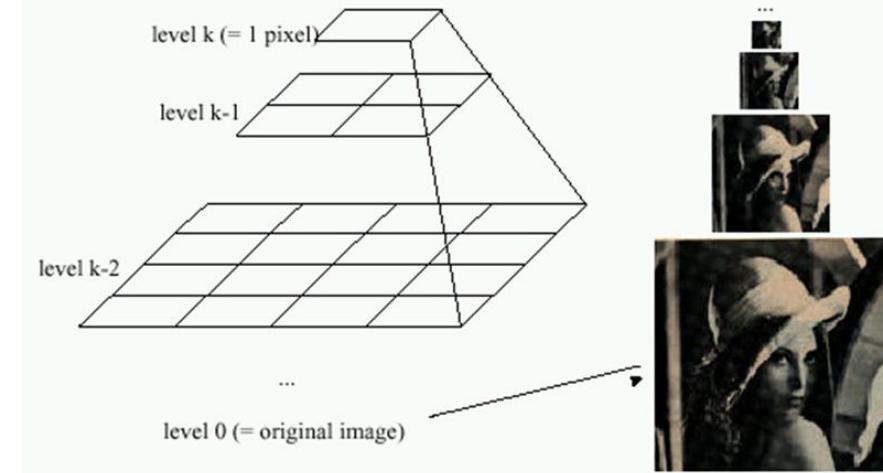
Start with this small one

# Efficient resizing: Image pyramids

Low resolution



High resolution

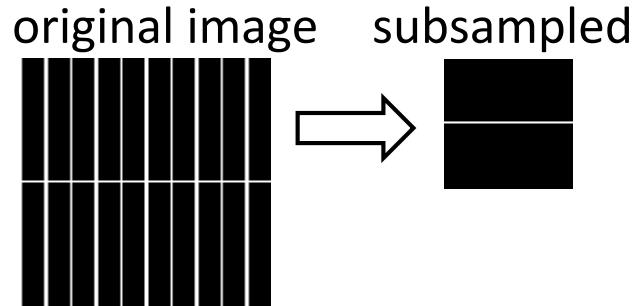


*Reduce (resample) the image!*

# How do we reduce an image?

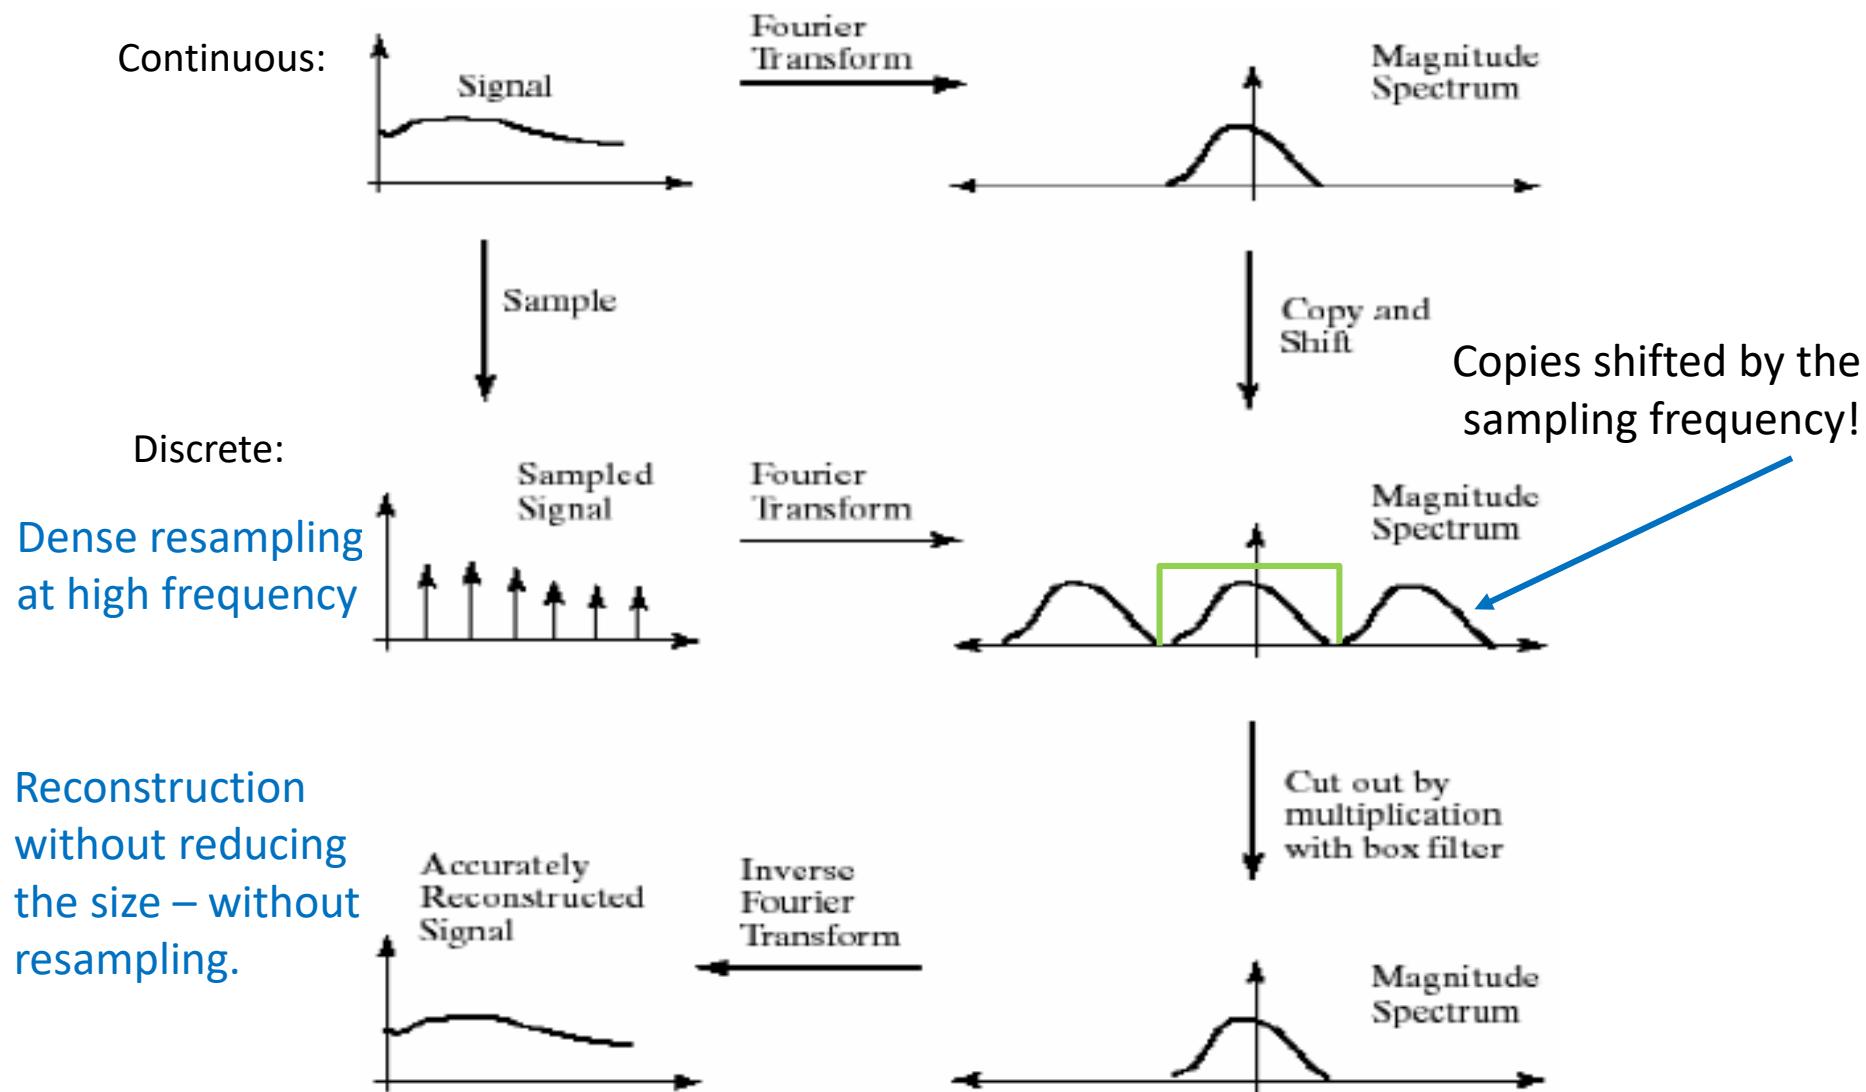
---

- Naive:
  - Remove **every second** pixel...

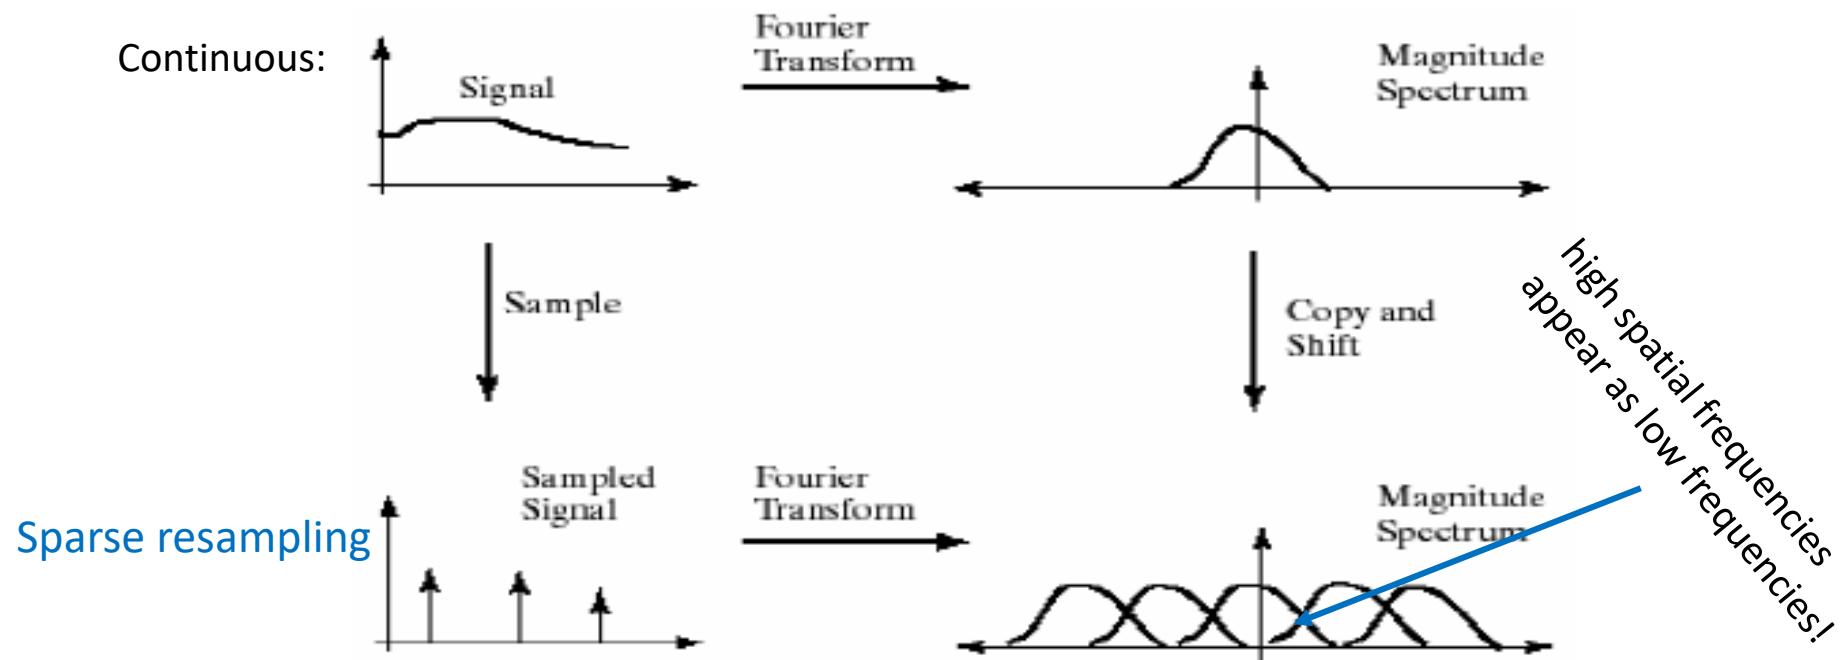


- Problem: *the structures in image change!*
- This effect is called *Aliasing*.
- Let's look into frequency domain to explain this...

# Sampling and aliasing (slo., prekrivanje)

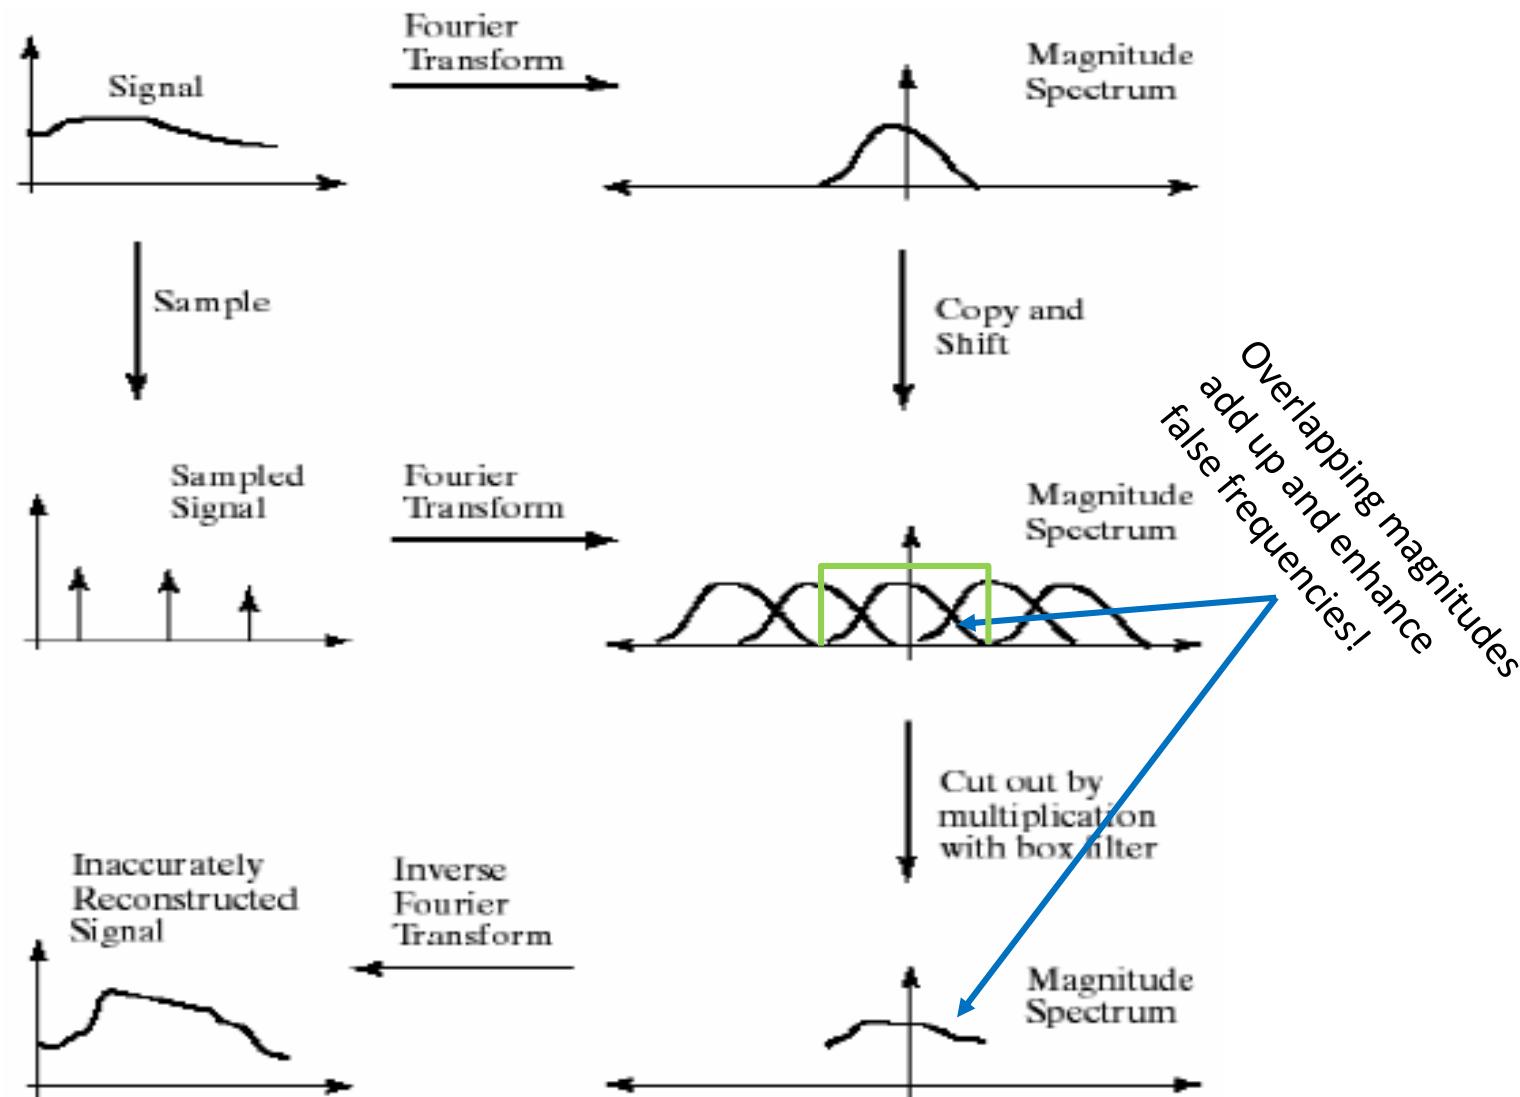


# Sampling and aliasing (slo., prekrivanje)



- Nyquist theorem:
  - If we want to reconstruct all frequencies up to  $f$ , we have to sample the signal by at least a frequency equal to  $2f$ .
  - This corresponds to a point after which repeated frequency spectra start to overlap (aliasing)!

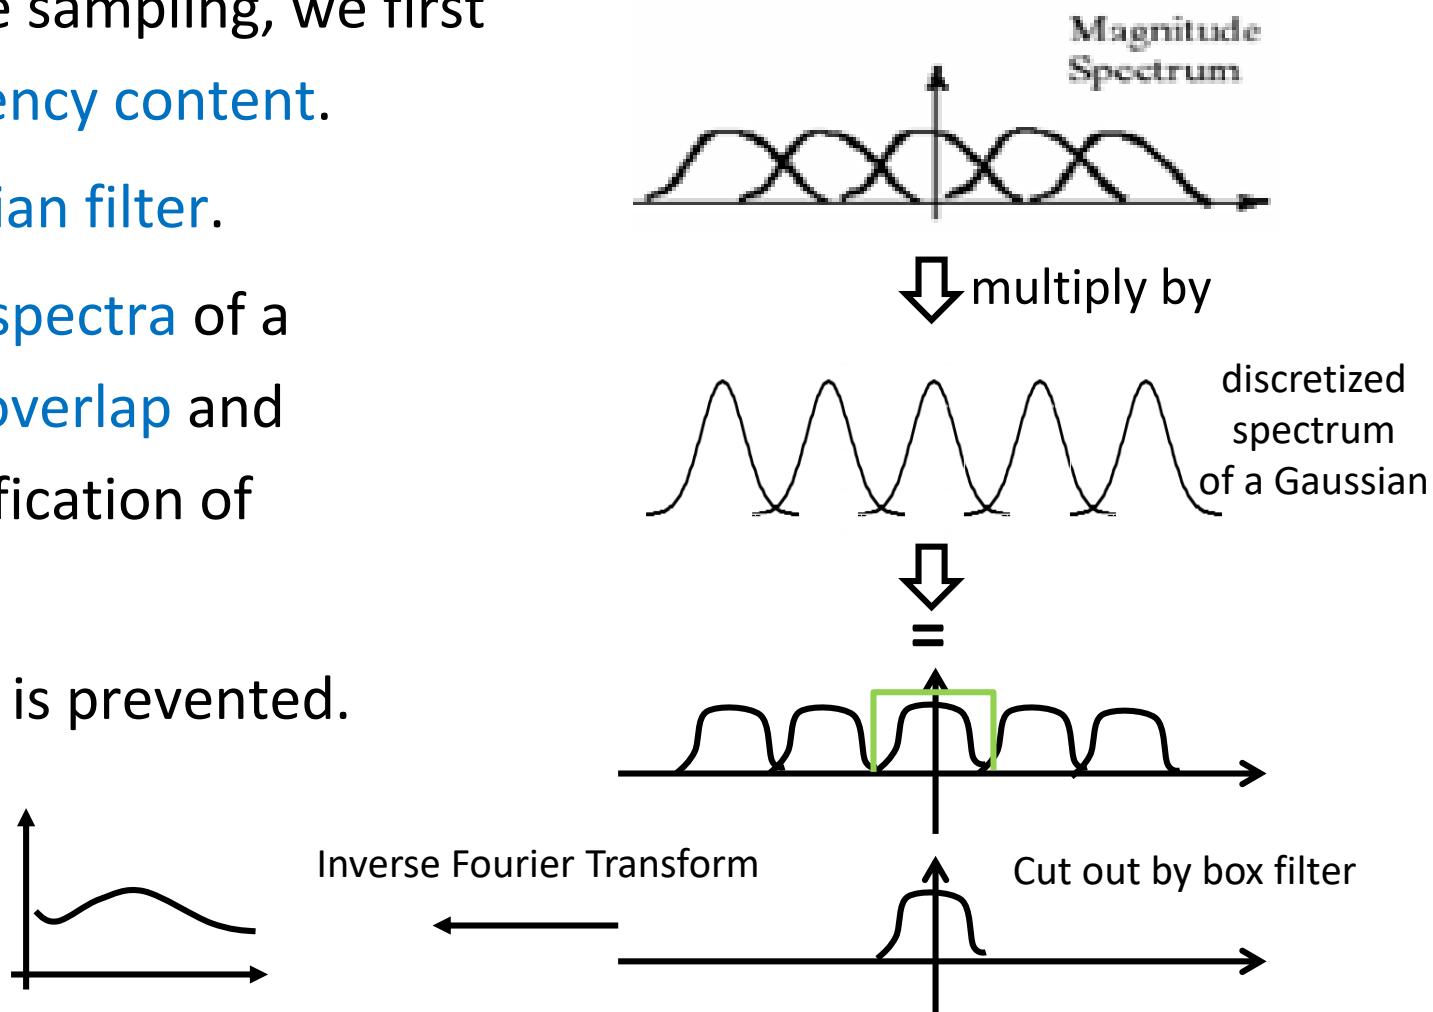
# Sampling and aliasing (slo., prekrivanje)



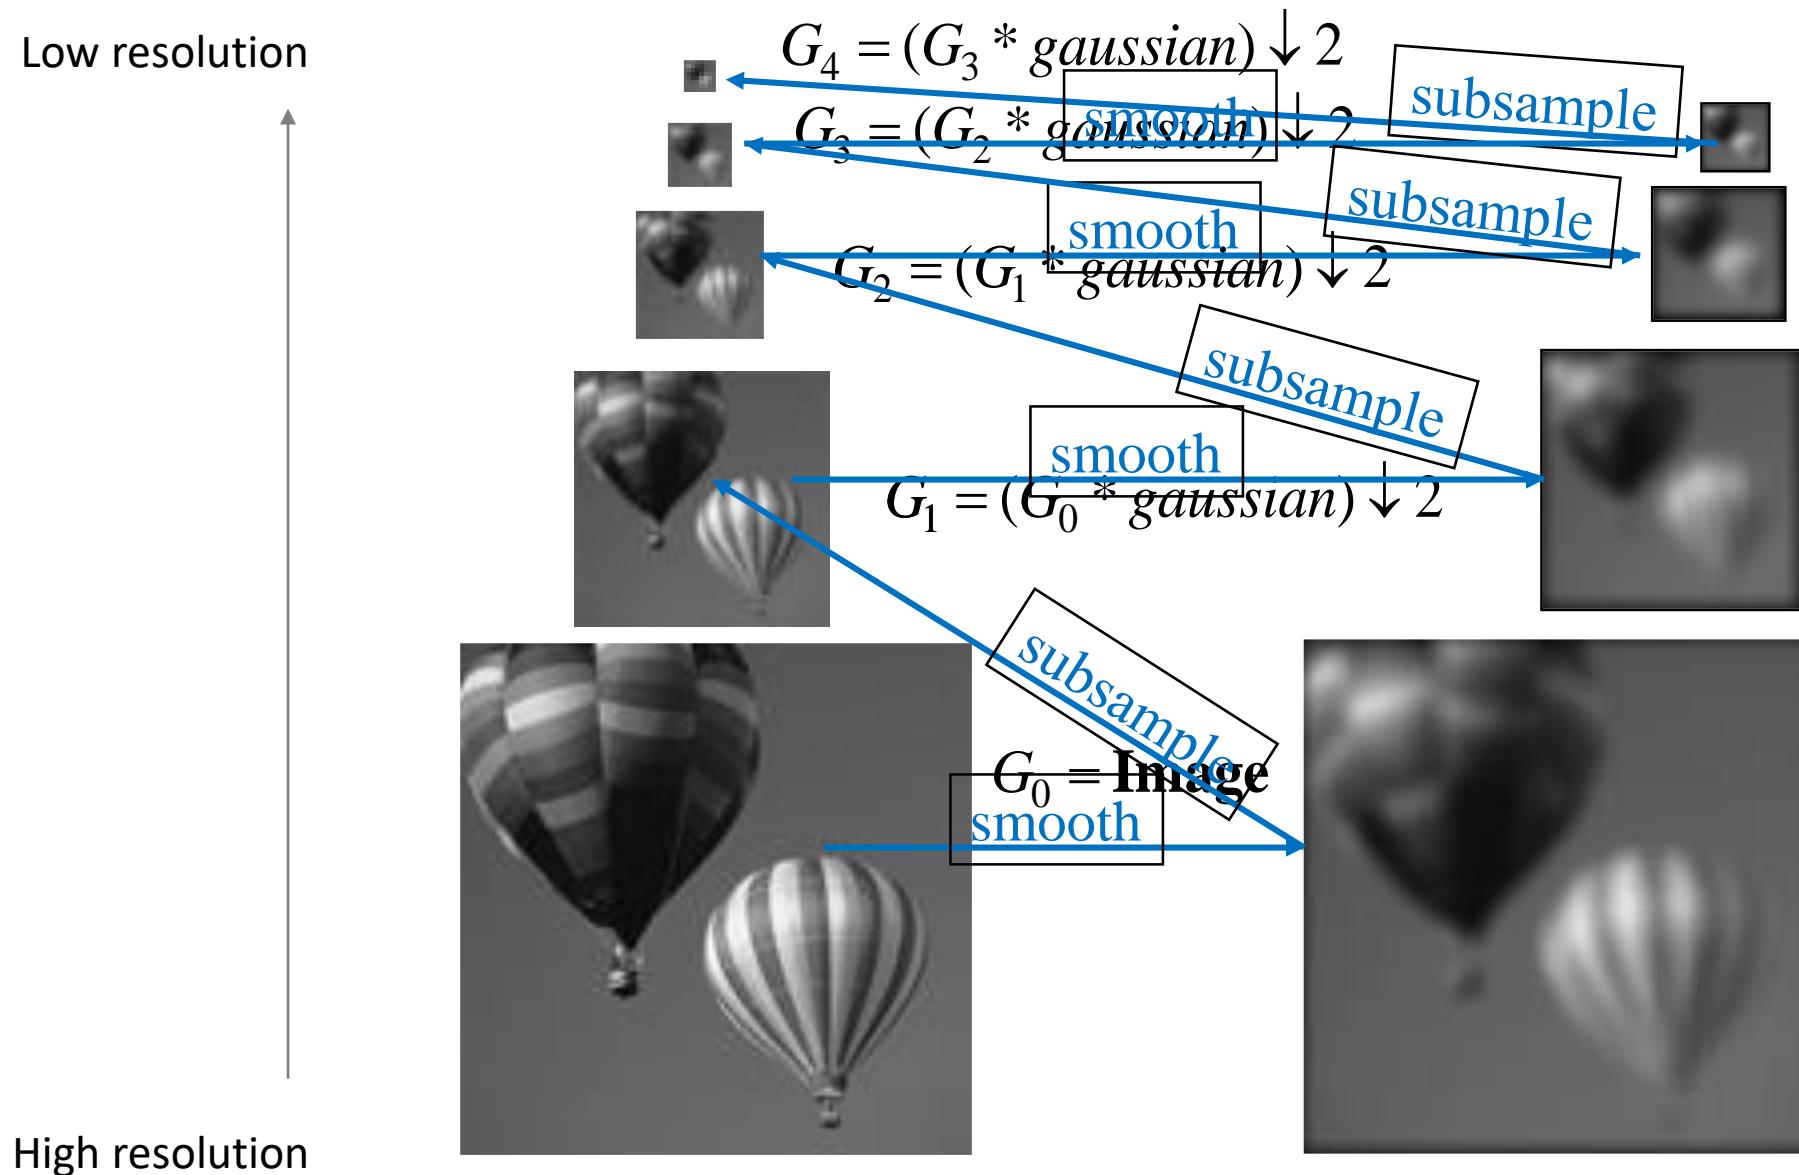
# Sampling and aliasing (slo., prekrivanje)

Proper resampling:

- When applying a sparse sampling, we first **remove the high frequency content**.
- This is done by a **Gaussian filter**.
- This way the **repeated spectra** of a sampled signal **do not overlap** and do not introduce amplification of false frequencies.
- Thus the *aliasing effect* is prevented.



# Gaussian pyramid



# Summary: Gaussian pyramid

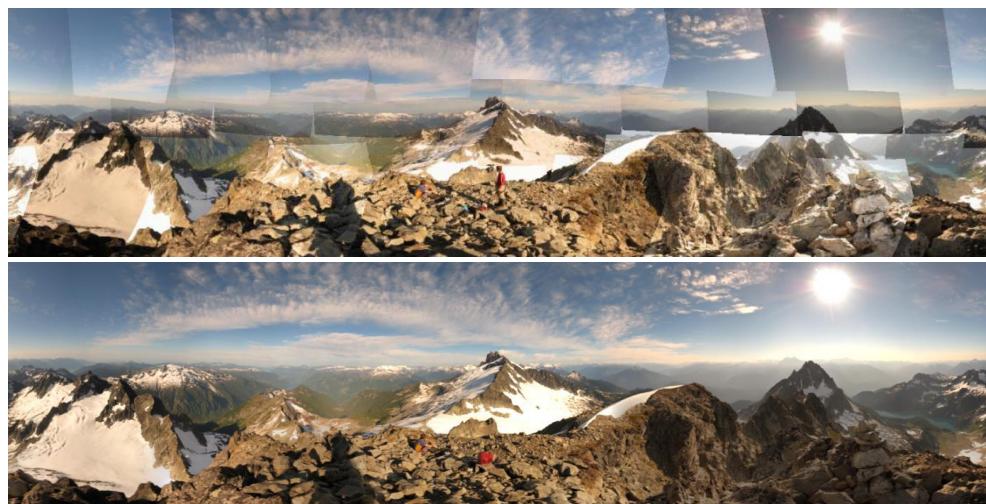
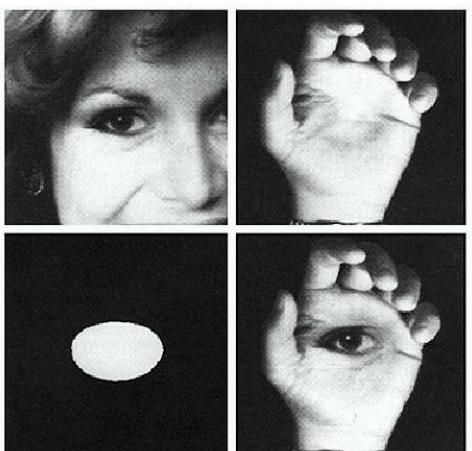
---

- Construction: get a **new level** directly from the **previous**
  - Smooth by a small filter and resample
- Reasons for Gaussian smoothing...
  - Convolution  $\text{Gauss} * \text{Gauss} = \text{new Gauss}$
  - $G(\sigma_1^2) * G(\sigma_2^2) = G(\sigma_1^2 + \sigma_2^2)$
- Reason for size reduction...
  - Gaussian is a **low-band-pass filter**, so we get a **redundant representation** of a smoothed image.  
⇒ No need to store a smoothed image in full resolution.

# Why a pyramid?

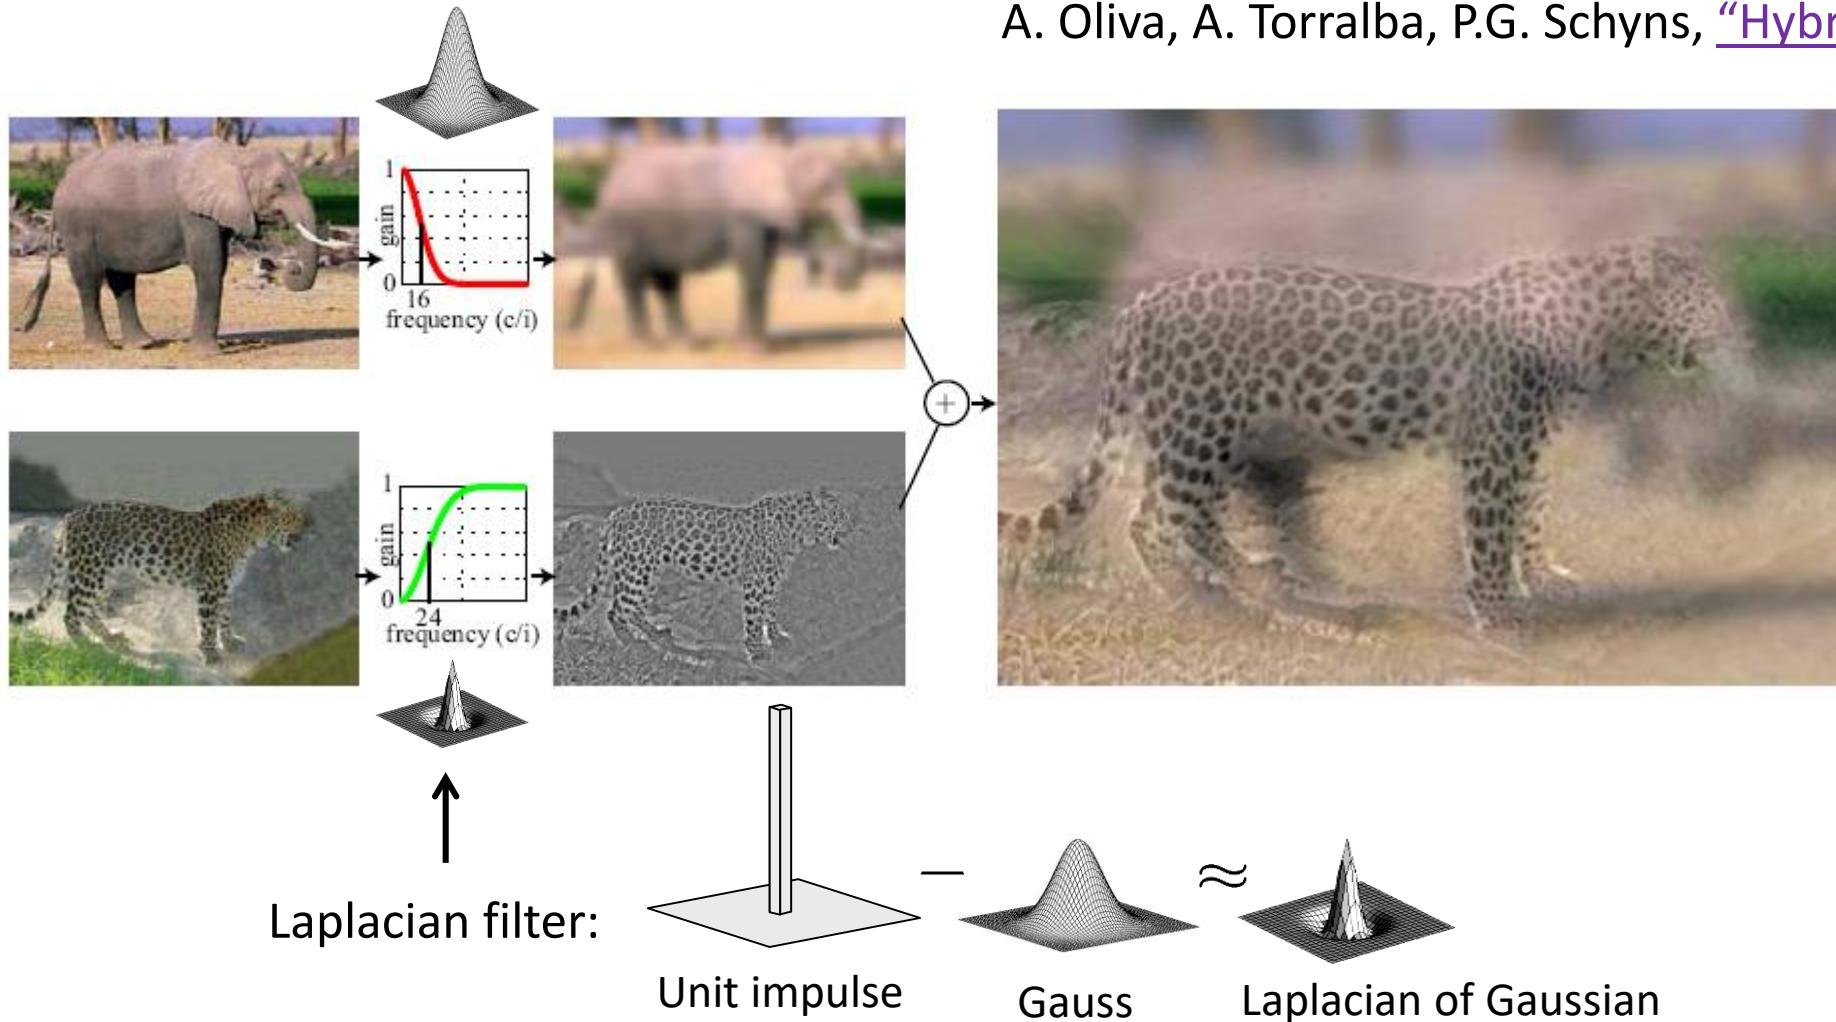
---

- Enables efficient implementation of many detection methods
- Multi-scale object detection ...
- Multi-scale edge detection ...
- Multi-scale feature point detection ...
- Manipulation of selected frequency bands ...
- Old stuff: Scale space



# Fun with hybrid images...

Gaussian filter



A. Oliva, A. Torralba, P.G. Schyns, ["Hybrid Images,"](#) SIGGRAPH 2006

# References

---

- David A. Forsyth, Jean Ponce, Computer Vision: A Modern Approach (2nd Edition), (prva izdaja dostopna na spletu)  
(Ozadje linearnih filtrov in povezavo s Fourierjevim transformom najdete v Poglavljih 7 in 8)
- R. Szeliski, Computer Vision: Algorithms and Applications, 2010
- Kristen Grauman, „Computer Vision“, lectures