




Machine perception

Fitting parametric models

Matej Kristan



Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani




Previously at MP...

- Edge detection (Canny)



- Fitting parametric models of shapes by voting (Hough transform)

- Lines
- Circles
- General shapes






Parametric models: Forward application

- Transformation parameterized by (many) parameters
 $\mathbf{x}'_i = f(\mathbf{x}_i; \mathbf{p})$
- Example: transform \mathbf{x}_i into \mathbf{x}'_i by a function $f(\mathbf{x}; \mathbf{p})$


$$\mathbf{x}_i = \begin{bmatrix} x_i \\ y_i \end{bmatrix} \in \mathbb{R}^2$$

$$\mathbf{p} = [\bar{p}_1, \bar{p}_2, \bar{p}_3, \dots, \bar{p}_M] \in \mathbb{R}^M$$



Parametric models: Use cases


- Inverse problem: “Given a set of correspondences, what are the parameters of the transformation?”



- Assuming the transformation can be well approximated by $f(x; p)$, what are the best parameter values for p ?

Parametric models: Use cases

- Best parameter values: *those that minimize the projection error*

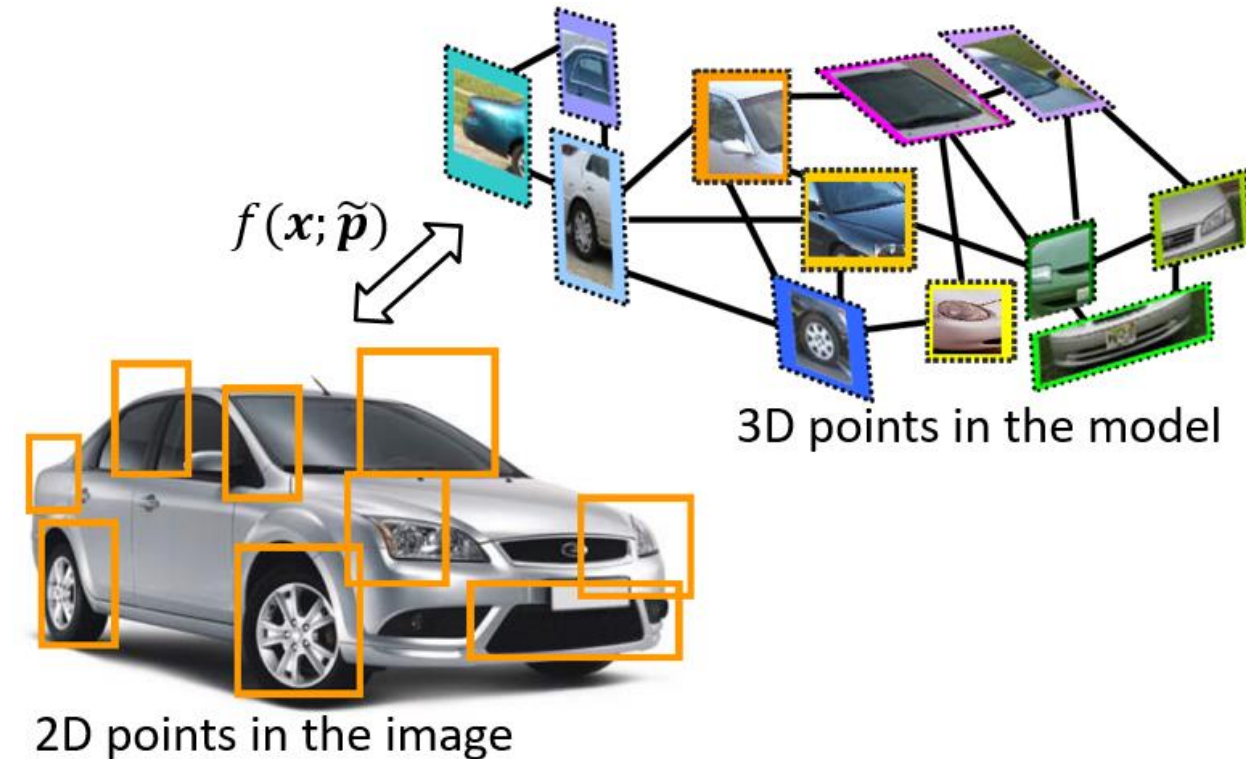


Stitched images:
Coordinates of all pixels in the left-hand
image transformed by $f(x; \tilde{p})$



Parametric models: Use cases

Example of a 3D pose estimation




TLD3.0 - 3D Tracking of Rigid Objects (ICCV 2017 demo)
<https://www.youtube.com/watch?v=i3cg8spZCrY>

Least squares: Line fitting

Problem formulation

- Data: $\{(x_1, y_1), \dots, (x_N, y_N)\}$
- Line equation:
 $y = f(x; \mathbf{p}) = xp_1 + p_2$
- Parameters:
 $\mathbf{p} = [p_1, p_2]^T$
- Projection error at i -th correspondence:
 $\varepsilon_i = f(x_i; \mathbf{p}) - y_i$
- The cost function (goodness of fit): $E(\mathbf{p}) = \sum_{i=1}^N \varepsilon_i^2$
- Best parameters: $\tilde{\mathbf{p}} = \arg \min_{\mathbf{p}} E(\mathbf{p})$



Least squares: Line fitting

Strategy:


1. Rewrite the cost function $E(\mathbf{p})$ into a vector-matrix form
2. Take derivative w.r.t. \mathbf{p} , set to zero, solve for \mathbf{p} .

$$f(x; \mathbf{p}) = xp_1 + p_2$$

$$\mathbf{p} = [p_1, p_2]^T$$

$$\varepsilon_i = f(x_i; \mathbf{p}) - y_i$$

$$E(\mathbf{p}) = \sum_{i=1}^N \varepsilon_i^2$$



Least squares: Line fitting

Strategy:

- Rewrite the cost function $E(\mathbf{p})$ into a vector-matrix form
- Take derivative w.r.t. \mathbf{p} , set to zero, solve for \mathbf{p} .

$$E(\mathbf{p}) = \sum_{i=1}^N \left(y_i - [x_i, 1] \begin{bmatrix} p_1 \\ p_2 \end{bmatrix} \right)^2 = \left\| - \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} + \begin{bmatrix} x_{1,1} \\ \vdots \\ x_{N,1} \end{bmatrix} \begin{bmatrix} p_1 \\ p_2 \end{bmatrix} \right\|^2 = \|\mathbf{-b} + \mathbf{A}\mathbf{p}\|^2$$

Normal equation:

$$\frac{dE(\mathbf{p})}{d\mathbf{p}} = 2\mathbf{A}^T \mathbf{A}\mathbf{p} - 2\mathbf{A}^T \mathbf{b} \equiv \mathbf{0}$$

Solution:

$$\mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b} = \mathbf{A}^\dagger \mathbf{b}$$

Pseudoinverse:

$$\mathbf{A}^\dagger = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T$$

$$\mathbf{A} \stackrel{\text{SVD}}{=} \mathbf{U} \mathbf{S} \mathbf{V}^T$$

$$\mathbf{A}^\dagger = \mathbf{V} \mathbf{S}^{-1} \mathbf{U}^T$$

A cookbook for normal equations:

1. Define the set of corresponding points

$$\{\mathbf{x}_i\}_{i=1:N}, \{\mathbf{x}'_i\}_{i=1:N}$$

2. Define the linear transformation

$$f(\mathbf{x}; \mathbf{p}): \mathbf{x} \rightarrow \mathbf{x}'$$

3. Define the per-point error and stack all errors into a single vector $\boldsymbol{\varepsilon}$:

$$E(\mathbf{p}) = \sum_{i=1}^N \varepsilon_i^2$$

$$\boldsymbol{\varepsilon}_i = f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i$$

$$\boldsymbol{\varepsilon} = [\varepsilon_1^T, \dots, \varepsilon_i^T, \dots, \varepsilon_N^T]^T$$

4. Rewrite the error into a form $\boldsymbol{\varepsilon} = \mathbf{A}\mathbf{p} - \mathbf{b}$


5. Solve by pseudoinverse: $\mathbf{p} = \mathbf{A}^\dagger \mathbf{b}$

Matlab: `p = A \ b`

Note: point errors $\boldsymbol{\varepsilon}_i$ are of same dimensionality as the points \mathbf{x}' .

Least squares: A simple image alignment

- Task: Align two images based on correspondences



- Assume a similarity transform (scale, rotation, translation)

$$\mathbf{x}' = f(\mathbf{x}; \mathbf{p})$$


- The similarity transform is parameterized by (See Szeliski, Section 2.1.2):

$$\underline{x}'_i = \begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} p_1 x_i - p_2 y_i + p_3 \\ p_2 x_i + p_1 y_i + p_4 \end{bmatrix}, \quad \mathbf{p} = [p_1 \ p_2 \ p_3 \ p_4]^T$$

Weighted least squares: Line fitting

Problem formulation

- Data: $\{(x_1, y_1), \dots, (x_N, y_N)\}$
- All points are *not* equally accurately measured!
- Weight at each point: w_i
- Projection error at i -th correspondence:
$$\varepsilon_i = f(x_i; \mathbf{p}) - y_i$$
- A **weighted cost**:
$$E(\mathbf{p}) = \sum_{i=1}^N w_i \varepsilon_i^2$$
- Best parameters:
$$\tilde{\mathbf{p}} = \arg \min_{\mathbf{p}} E(\mathbf{p})$$



Weighted least squares: Line fitting

Strategy:

- Rewrite the cost function $E(\mathbf{p})$ into a vector-matrix form
- Take derivative w.r.t. \mathbf{p} , set to zero, solve for \mathbf{p} .


$$f(x; \mathbf{p}) = xp_1 + p_2$$

$$\mathbf{p} = [p_1, p_2]^T$$

$$\varepsilon_i = f(x_i; \mathbf{p}) - y_i$$

$$E(\mathbf{p}) = \sum_{i=1}^N w_i \varepsilon_i^2$$

$$\tilde{\mathbf{p}} = \arg \min_{\mathbf{p}} E(\mathbf{p})$$



Weighted least squares: Line fitting

Strategy:

- Rewrite the cost function $E(\mathbf{p})$ into a vector-matrix form
- Take derivative w.r.t. \mathbf{p} , set to zero, solve for \mathbf{p} .

$$E(\mathbf{p}) = \sum_{i=1}^N w_i \left(y_i - [x_i, 1] \begin{bmatrix} p_1 \\ p_2 \end{bmatrix} \right)^2$$

$$E(\mathbf{p}) = [\varepsilon_1, \dots, \varepsilon_N] \begin{bmatrix} w_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & w_N \end{bmatrix} \begin{bmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_N \end{bmatrix} = \varepsilon^T \mathbf{W} \varepsilon$$

$$\frac{dE(\mathbf{p})}{d\mathbf{p}} = 2\mathbf{A}^T \mathbf{W} \mathbf{A} \mathbf{p} - 2\mathbf{A}^T \mathbf{W} \mathbf{b} \equiv 0 \quad \longleftarrow \text{Normal equation}$$

$$\mathbf{p} = (\mathbf{A}^T \mathbf{W} \mathbf{A})^{-1} \mathbf{A}^T \mathbf{W} \mathbf{b}$$

A cookbook for weighted least squares:

1. Define a weighted set of corresponding points

$$\{\mathbf{x}_i\}_{i=1:N}, \{\mathbf{x}'_i\}_{i=1:N}, \{w_i\}_{i=1:N}$$

Note: $\mathbf{x}' \in \mathbb{R}^d, w \in \mathbb{R}^1$

2. Define the linear transformation

$$f(\mathbf{x}; \mathbf{p}): \mathbf{x} \rightarrow \mathbf{x}'$$

3. Rewrite the error into a form $\boldsymbol{\varepsilon} = \mathbf{A}\mathbf{p} - \mathbf{b}$

4. Create a weight matrix \mathbf{W} as

$$\mathbf{W} = \text{diag}([\mathbf{w}_1^T, \dots, \mathbf{w}_N^T])$$

$$\text{with } \mathbf{w}_i^T = w_i[1, \dots, 1]_{1 \times d}$$


Note: think about why are \mathbf{w}_i^T vectors of same dimensionality as the points \mathbf{x}' .

5. Solve by : $\mathbf{p} = (\mathbf{A}^T \mathbf{W} \mathbf{A})^{-1} \mathbf{A}^T \mathbf{W} \mathbf{b}$

To practice: solve the
“sailboat” example

Robust least squares

- Quadratic cost function behaves poorly with outliers:




- To see where the problem lies, we will have to **rewrite our cost** function into a general form.
- The cost can be generally written as: $E(\mathbf{p}) = \sum_{i=1}^N h(\varepsilon_i)$
- For ordinary least squares we had: $h(\varepsilon_i) = \|\varepsilon_i\|^2$

Robust least squares

$$E(\mathbf{p}) = \sum_{i=1}^N h(\varepsilon_i)$$


Not a robust function $h(\epsilon)$:

Squared-
error




Robust cost functions $h(\epsilon)$:


Blake-
Zisserman




corrupted
Gaussian




Cauchy




L_1



Huber



pseudo-
Huber



Robust least squares

- For a cost function with robust error function $h(\varepsilon_i)$

$$E(\mathbf{p}) = \sum_{i=1}^N h(\varepsilon_i)$$

- It is possible to find an equivalent weighted L_2 cost

$$E_W(\mathbf{p}) = \sum_{i=1}^N w(\varepsilon_i) ||\varepsilon_i||^2$$

$$\text{with } w = \frac{h'(\varepsilon)}{\varepsilon} \text{ and } h'(\varepsilon) = \frac{\partial h(\varepsilon)}{\partial \varepsilon} .$$

- **Problems:**
 1. Weights depend on the errors incurred by the optimal parameters of our model.
 2. But the *parameters are unknown* and so are the weights.
- **Solution:** Can apply an iterative approach that will converge as long as $h(\sqrt{|\epsilon|})$ is concave¹.

¹Aftab, K. and Hartley, R., Convergence of Iteratively Re-weighted Least Squares to Robust M-estimators, WACV 2015

R. Hartley, Robust Optimization Techniques in Computer Vision, [Session 3](#), ECCV2014 tutorials

Iterative reweighted least squares

1. Set all the weights to $w_i^{t-1} = 1$.
2. Solve for \mathbf{p}^t by the weighted least squares problem.
3. Using the estimated parameters \mathbf{p}^t re-calculate per-point projection errors $\boldsymbol{\varepsilon}_i^t$.
4. Using the projection errors re-calculate new weights w_i^t from:
$$w = \frac{h'(\varepsilon)}{\varepsilon} \quad h'(\varepsilon) = \frac{\partial h(\varepsilon)}{\partial \varepsilon}$$
5. Go back to step 2 and continue until the change in parameters is negligibly small (convergence).

Note: $(\cdot)^t$ indicates a step of iteration in the iterative reweighted least squares.

For an instructive discussion on parameters of the Huber cost function from data, please see:

J. Fox, [Robust Regression--Appendix to An R and S-PLUS Companion to Applied Regression](#), 2002, "1.1 Objective Functions".

Constrained least squares

- Often we will seek parameters \mathbf{p} that satisfy constraints.
- Reconsider line-fitting example, but this time we'll minimize *perpendicular* distances!

$$E(\mathbf{p}) = \sum_{i=1}^N ||\varepsilon_i||^2$$

- Re-parameterize:


$$\mathbf{p} = [p_1, p_2, p_3]^T$$

- Distance of a point to line:

$$||\varepsilon_i||^2 = (x_i p_1 + y_i p_2 - p_3)^2$$

- Let's minimize:

$$E(\mathbf{p}) = \sum_{i=1}^N ||\varepsilon_i||^2$$




Constrained least squares

- The solution: $\frac{dE(\mathbf{p})}{d\mathbf{p}} = 2\mathbf{A}^T \mathbf{A} \mathbf{p} \equiv \mathbf{0}$
- Trivial solution: $\mathbf{p} = \mathbf{0}$
- A nontrivial solution is obtained by constraint $||\mathbf{p}||^2 = 1$

$$\mathbf{p} = [p_1, p_2, p_3]^T$$

$$||\boldsymbol{\varepsilon}_i||^2 = (x_i p_1 + y_i p_2 - p_3)^2$$

$$E(\mathbf{p}) = \sum_{i=1}^N ||\boldsymbol{\varepsilon}_i||^2$$



Constrained least squares

- The solution: $\frac{dE(\mathbf{p})}{d\mathbf{p}} = 2\mathbf{A}^T \mathbf{A} \mathbf{p} \equiv \mathbf{0}$
- Trivial solution: $\mathbf{p} = \mathbf{0}$
- A nontrivial solution is obtained by constraint $||\mathbf{p}||^2 = 1$
- Taking the derivative of a Lagrangian and setting to 0:
$$\mathbf{A}^T \mathbf{A} \mathbf{p} = \lambda \mathbf{p} \quad \longleftarrow \text{Homogenous equation!}$$
- The solution is the eigenvector of $(\mathbf{A}^T \mathbf{A})$ corresponding to the smallest eigenvalue.
- Actually, it can be shown that this is also the eigenvector corresponding to the smallest eigenvalue of \mathbf{A} . (see notes on “Avoid computing $\mathbf{A}^T \mathbf{A}$ ”)

Recognizing the hammer for your nail!

- Problems that can be written as systems of equations (*normal equations*):

$$A\mathbf{p} = \mathbf{b}$$

(if you have weights on equations, then $\mathbf{W}A\mathbf{p} = \mathbf{W}\mathbf{b}$)

can be solved by ordinary LS or IRWLS

Matlab: $\mathbf{p} = A \setminus \mathbf{b}$;

- Problems that result in a homogenous system:

$$A\mathbf{p} = \mathbf{0}$$



can be solved by putting the constraint $\|\mathbf{p}\|^2 = 1$, the solution is the eigenvector corresponding to the smallest eigenvalue.

(If required, rescale the solution for \mathbf{p})

Matlab: $[U,S,V] = \text{svd}(A)$; $\mathbf{p} = V(:,\text{end})$;

Generally for nonlinear cost functions

- Often **nonlinear error** functions are used, which **cannot** be minimized analytically in a **closed form**.
- **Popular** approaches:
 - Gradient descend
 - [Newton's method](#)
 - Gauss-Newton method
 - Levenberg-Marquardt
 - Alternate direction method of multipliers (ADMM) [!wery powerful & simple]
- More about these:
 - [Fua and Lepetit: Computer Vision Fundamentals: Robust Non-Linear Least-Squares and their Applications](#)
 - Griva et al., [Linear and Nonlinear Optimization](#) (See appendix on Matrix Algebra)
 - [The Matrix Cookbook](#) (List of common vector/matrix solutions)
 - Forsyth, Ponce, „Computer Vision – A modern approach“, (Appendix in *2nd ed.*)



Need to deal even better with outliers

- Large **disagreements** in only a few points (outliers) cause failure of the least-squares-based methods.
- The **detection, localization** and **recognition** in CV have to operate in significantly **noisy data**.
- In some cases **>½ data** is expected to be outliers.
- **Standard** methods for robust estimation can **rarely deal** with such a **large** proportion of **outliers**.

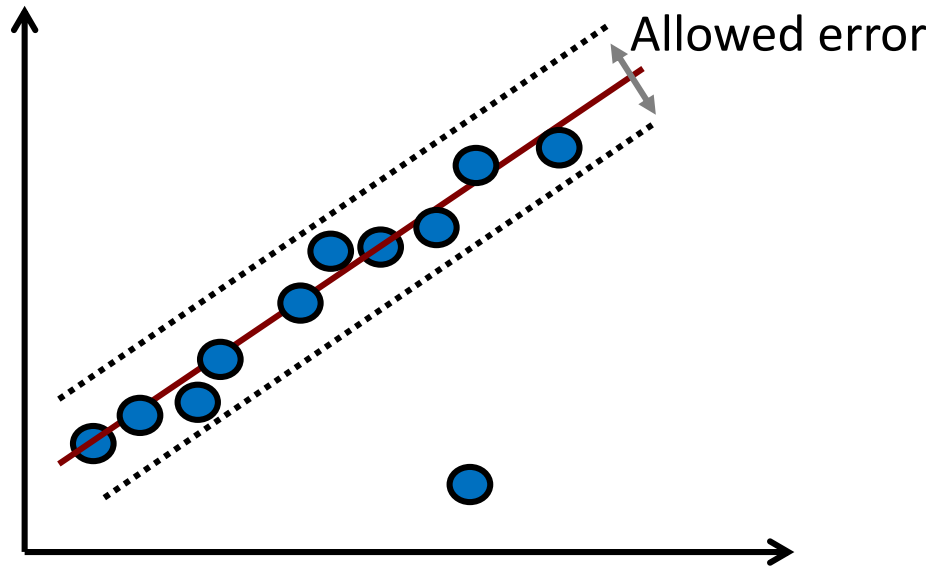
RANSAC

- The **RANSAC** algorithm (random sample consensus).
- Very **popular** due to its generality and simplicity.
- Can deal with **large** portions of outliers.
- Published in 1981 (Fischler in Bolles)
- One of the **most cited** papers in Computer Vision
- Many improvements proposed since!


M. A. Fischler, R. C. Bolles. [Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography](#). Comm. of the ACM, Vol 24, pp. 381-395, 1981.

RANSAC: Intuition by line fitting

- A good estimate of our model should have a strong support in data:
“recognize a good model when you see it”



10 point support this line!




4 point support this line!

- How to find a model with a strong support?
- By randomly sampling potential models.

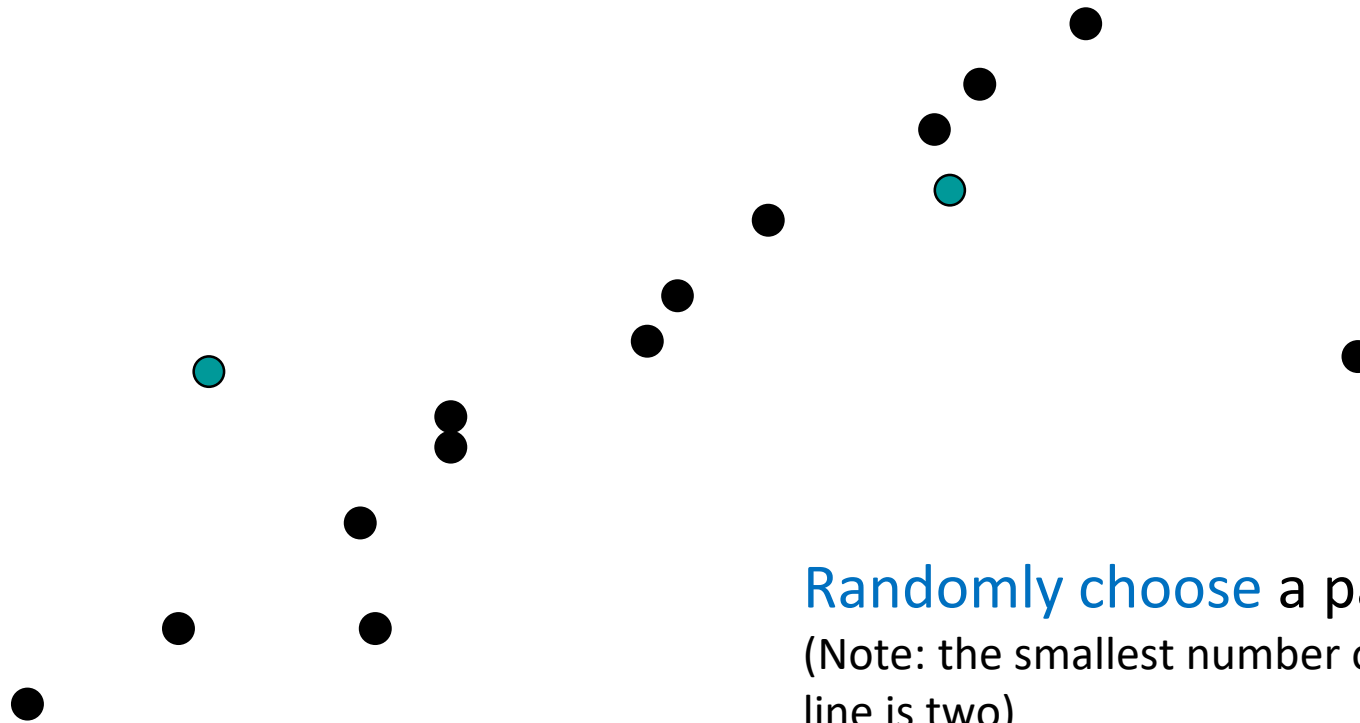
RANSAC: Intuition by line fitting

- Task: Robustly estimate **the most likely** line



RANSAC: Intuition by line fitting


- Task: Robustly estimate **the most likely** line



Randomly choose a pair of points
(Note: the smallest number of points to fit a line is two)

RANSAC: Intuition by line fitting

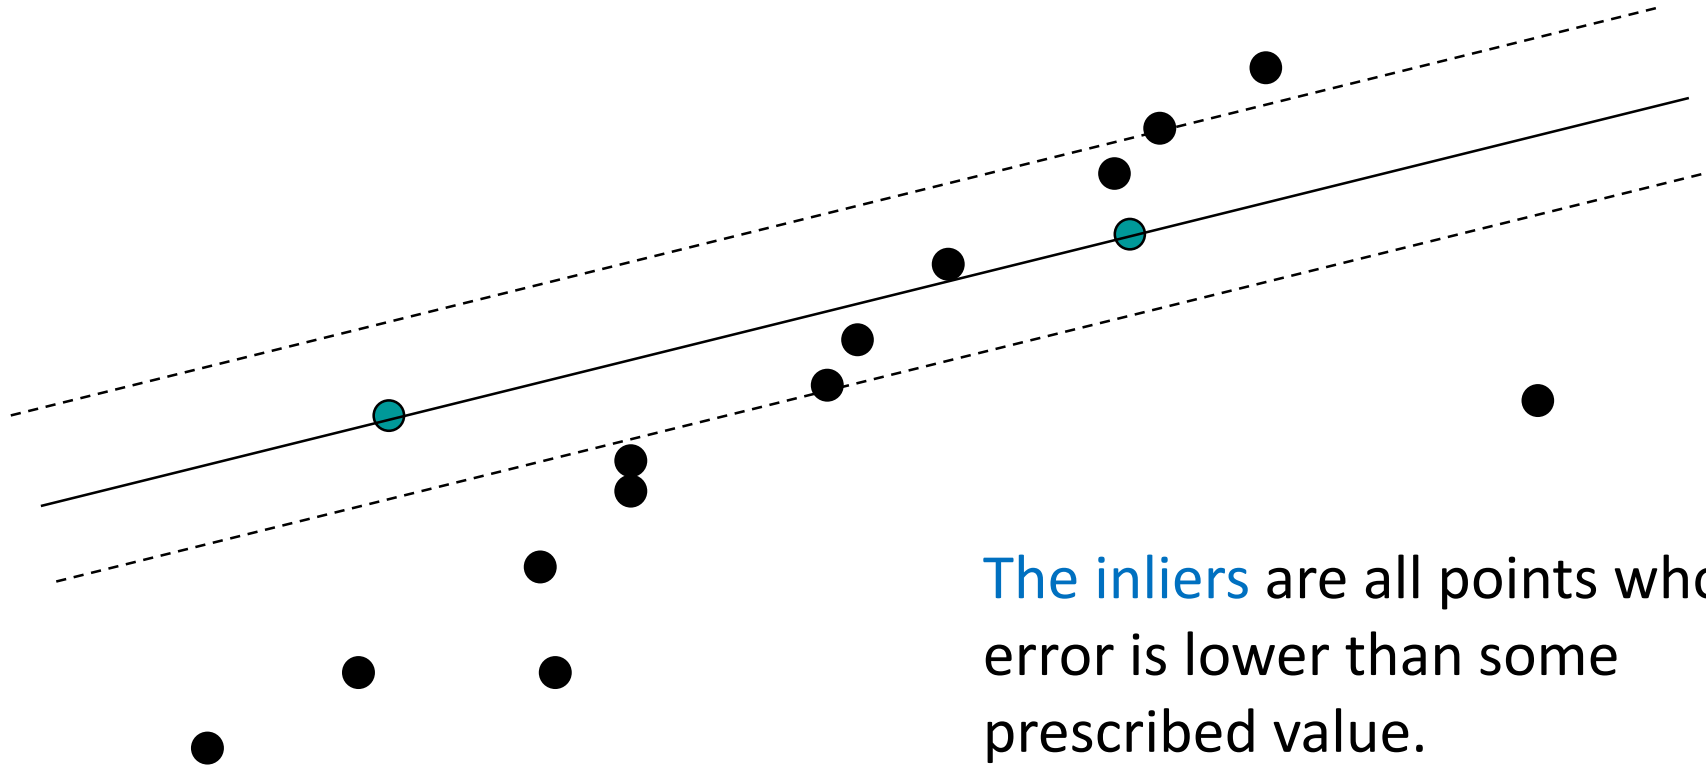
- Task: Robustly estimate **the most likely** line



Fit the line to the selected points.

RANSAC: Intuition by line fitting

- Task: Robustly estimate **the most likely** line




The **inliers** are all points whose error is lower than some prescribed value.

$$\varepsilon_i = |f(x_i; \mathbf{p}) - y_i|$$

RANSAC: Intuition by line fitting


- Task: Robustly estimate **the most likely** line



Repeat N-iterations, or, until the support becomes strong enough
(actually this is an oversimplification).


RANSAC: line fitting

- Another example




A general setting

1. Define the set of “potentially” corresponding points:
 $\{\mathbf{x}_i\}_{i=1:N}$, $\{\mathbf{x}'_i\}_{i=1:N}$
2. Define the transformation model: $f(\mathbf{x}; \mathbf{p}): \mathbf{x} \rightarrow \mathbf{x}'$



A simple RANSAC loop




$$\{\mathbf{x}_i\}_{i=1:N}, \{\mathbf{x}'_i\}_{i=1:N}$$

$$f(\mathbf{x}; \mathbf{p}): \mathbf{x} \rightarrow \mathbf{x}'$$

1. Randomly select the **smallest** group of correspondences, from which we can estimate the parameters of our model.
2. Fit the parametric model **to the selected** correspondences: $\tilde{\mathbf{p}}$

A simple RANSAC loop




$$\{\mathbf{x}_i\}_{i=1:N}, \{\mathbf{x}'_i\}_{i=1:N}$$

$$f(\mathbf{x}; \mathbf{p}): \mathbf{x} \rightarrow \mathbf{x}'$$

1. Randomly select the smallest group of correspondences, from which we can estimate the parameters of our model.
 2. Fit the parametric model to the selected correspondences: $\tilde{\mathbf{p}}$
 3. Count how many of all correspondences are in agreement with the fitted model – number of inliers.
- Remember the model parameters that maximize the number of inliers.

The choice of parameters

- **How many** correspondences " s " are required?
 - Typically the smallest number that allows estimating the model parameters, i.e., as many as the model parameters.
- **Threshold** distance t for proclaiming the inliers
 - Choose t , such, that the probability that an inlier falls below the threshold is equal to p_w . For example ($p_w=0.95$)
 - Assuming a Gaussian noise on the measurements.
The noise standard dev. σ : $t=2\sigma$
- **Number** of sampling iterations N
 - Chose N such, that the **probability** p of drawing a sample with all inliers at least once **is high enough**.



The choice of parameters: N

- Setting the number of sampling iterations N :
 - Assume we know the proportion e of outliers (probability of selecting an outlier at random).
 - Choose N such, that the probability of drawing a sample set with all inliers at least once in N draws is p , (e.g., $p=0.99$).
 - Derive the probability of drawing a bad sample in N trials, $1 - p = p_{bad}^N$, and expose N
 - Probability of choosing a single inlier: $1 - e$
 - Probability of an all-inlier sample:
→ s -times sample an inlier: $(1 - e)^s$
 - Probability, of a bad sample:
→ at least one of s not an inlier: $[1 - (1 - e)^s]$
 - Probability of always drawing a bad sample in N trials: $(1 - (1 - e)^s)^N$


$$1 - p = (1 - (1 - e)^s)^N \quad \Rightarrow \quad N = \frac{\log(1-p)}{\log(1-(1-e)^s)}$$

The choice of parameters: N

Number of iterations N required to sample an inlying model with s parameters at least once with probability p if the proportion of outliers is e :


s	portion of outliers: e						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Tabulated values of N for $p = 0.99$



After RANSAC: Refit by LS

- RANSAC **splits** the data into **inliers** and **outliers**, and calculates the model parameters using a **minimal number** of correspondences.
- **Improve** the model parameters by applying least squares to the inliers.



Beyond the simple RANSAC

- A great deal of research was invested by many researchers into improving RANSAC
- Particularly in [finding](#) the right [solution faster](#)
- Or improving [resiliency to outliers](#)
- Please see online resources:
 - [MLESAC](#) (uses maximum likelihood on hypothesis verification)
 - [PROSAC](#) (better chooses the order of samples)
 - Or an entire presentation dedicated to recent developments on RANSAC:
J. Matas [RANSAC in 2011 – 30 years after](#), CVPR, 2011

RANSAC: Summary

- Pros
 - Very **simple** and **general**
 - Applicable to **many real-life problems**
 - Often **used in practice**
- Cons
 - Requires setting some **parameters**.
 - Potentially **many iterations required** to find the optimum.
 - Fails at **very small number** of inliers.
 - In some cases **more accurate procedures**, that do not require brute-force sampling, can be found.

Fitting: Challenges

- If we know the inliers how to estimate the parameters?
 - Least squares
- What if our data includes outliers?
 - Robust least squares, RANSAC
- What if we have multiple instances of our model (e.g., multiple lines)?
 - Apply voting: sequential RANSAC, Hough transform
- What if we have multiple models (e.g., unknown degree of a polynomial)?
 - Apply model selection (e.g., MDL, BIC, AIC)
- Complicated nonparametric models
 - Generalized Hough (GHT)
 - Iterative Closest Point, (ICP) == iterative local least squares

Further reading

- Another simple and interesting way to iteratively fit a complicated model to data:
Iterative Closest Point method
Matlab implementation: ICP
- A very nice and accessible tutorial on nonlinear optimization in computer vision: <http://cvlabwww.epfl.ch/~fua/courses/lsg/Intro.htm>

References

- R. Szeliski, [Computer Vision: Algorithms and Applications](#), 2010
- [David A. Forsyth](#), [Jean Ponce](#), Computer Vision: A Modern Approach (2nd Edition), (*second edition!*)
 - See appendix on Normal equations and Homogeneous systems
- Igor Griva, Stephen G. Nash, Ariela Sofer, [Linear and Nonlinear Optimization](#)
 - See appendix on Matrix Algebra
- [The Matrix Cookbook](#)
 - List of common vector/matrix solutions